

The Algorithm Little OaldresPuzzle_Cryptic

Technical Details

An Innovative Lightweight Symmetric Encryption Algorithm Integrating NeoAlzette ARX S-box and XCR CSPRNG

Twilight-Dream

August 10, 2025

Abstract

This paper introduces "Little OaldresPuzzle_Cryptic," a novel lightweight symmetric encryption algorithm.

At the core of this algorithm are two main cryptographic components: the NeoAlzette permutation S-box based on ARX (Addition-Rotation-XOR) primitives and the innovative pseudo-random number generator XorConstantRotation (XCR), used exclusively in the key expansion process. The NeoAlzette S-box, a non-linear function for 32-bit pairs, is meticulously designed for both encryption strength and operational efficiency, ensuring robust security in resource-constrained environments. During the encryption and decryption processes, a pseudo-randomly selected mixed linear diffusion function, distinct from XCR, is applied, enhancing the complexity and unpredictability of the encryption.

We comprehensively explore the various technical aspects of the Little OaldresPuzzle_Cryptic algorithm.

Its design aims to balance speed and security in the encryption process, particularly for high-speed data transmission scenarios. Recognizing that resource efficiency and execution speed are crucial for lightweight encryption algorithms, without compromising security, we conducted a series of statistical tests to validate the cryptographic security of our algorithm. These tests included assessments of resistance to linear and differential cryptanalysis, among other measures.

By combining the NeoAlzette S-box with sophisticated key expansion using XCR, and integrating the pseudo-randomly selected mixed linear diffusion function in its encryption and decryption processes, our algorithm significantly enhances its capability to withstand advanced cryptographic analysis techniques while maintaining lightweight and efficient operation. Our test results demonstrate that the Little OaldresPuzzle_Cryptic algorithm effectively supports the encryption and decryption needs of high-speed data, ensuring robust security and making it an ideal choice for various modern cryptographic application scenarios.

Contents

1	Introduction	3
2	XCR/Little OaldresPuzzle_Cryptic Stream Cipher Design Background	4
2.1	What are ARX Structure and Salsa20, ChaCha20 Algorithms?	4
2.2	About the XCR Structure of Our Lightweight Symmetric Encryption Technology	4
2.2.1	Designing XCR from the ARX Structure: Insights from NIST and Lightweight Cryptography	4

3 XCR Algorithm: Structure and Operations	5
3.1 Preliminaries and Notations	5
3.1.1 Operators	5
3.2 Overall Steps	5
3.3 Diffusion and Confusion Layers	6
3.4 Selection of Rotation Amounts	6
3.5 Algorithm Efficiency and Security	6
4 Detailed Description of the XCR Algorithm Structure	6
4.1 XCR Algorithm: State Initialization Function	6
4.2 XCR Algorithm: State Iteration / Update Function	7
4.2.1 Dynamic Constant Generation	7
4.2.2 Diffusion Layer Operations	7
4.2.3 Confusion Layer Operations	8
4.2.4 Update Rules	8
5 NeoAlzette Substitution Box	8
5.1 Round Constant Derivation	9
5.2 Security Parameterization	9
5.3 Forward and Backward Layers	10
5.4 Operational Modes	10
5.5 Applied Little OaldresPuzzle_Cryptic Algorithm: Encryption and Decryption	10
6 Little OaldresPuzzle_Cryptic Algorithm Overview	11
6.1 Symmetric Encryption Specification	11
6.2 Symmetric Decryption Specification	11
6.2.1 Little OaldresPuzzle_Cryptic Components	11
6.3 XCR CSPRNG-based Key Schedule	12
6.4 XCR Round Constant Generation	12
6.5 Binary Pseudo-Hadamard Transform (BPHT)	14
7 Performance Evaluation and Security Evaluation	15
7.1 Performance Evaluation	15
7.2 Security Evaluation	16
7.3 Security Evaluation with Statistical Tests	17
7.4 The Credibility of Statistical Tests and the Rationale for Their Use	17
7.5 Integrating P-Values, Q-Values, and Conditional Probability in Cryptographic Analysis	17
7.5.1 Test Results for Phase 1	17
7.5.2 Test Results for Phase 2	17
7.5.3 Evaluation of Test Results	18
8 Mathematical Proof and Security Deduction of Our Algorithms	18
8.1 Formal Security Analysis of XorConstantRotation (XCR)	18
8.1.1 Primitive Specification (Matches Implementation)	18
8.1.2 Security Model (RoR with Chosen-Input Queries)	18
8.1.3 Assumptions (Made Explicit)	19
8.1.4 What Can Be Proven Rigorously vs. What Is Heuristic	19
8.1.5 Collision and Cycle Discussion (State vs. Output)	19
8.1.6 Quantum Query Considerations	19
8.1.7 Design Takeaways (Instantiations)	20
8.2 ARX Probabilistic Analysis [Ya, 2017]	22
8.2.1 Modular Addition Differential Analysis	22
8.2.2 Rotation Analysis	22
8.2.3 XOR Operation Analysis	22
8.2.4 Linear Analysis Framework	22
8.2.5 Formal Security Boundaries	23
8.3 Security Boundary Derivation	23

8.4	NeoAlzette ARX Structure Deep Security Analysis	24
8.4.1	Step-by-Step Differential Analysis of the NeoAlzette Forward Layer	24
8.4.2	Step-by-Step Linear Correlation Analysis of the NeoAlzette Forward Layer	26
8.4.3	NeoAlzette ARX-SBOX: Formal Statement, Proof of Correctness, and Security Positioning	27
8.4.4	Ideal-Theoretic Upper/Lower Bounds (No Empirical Data)	28
8.4.5	Consolidated Single-Round Statistics	29
8.4.6	Structural Security Enhancements over Alzette	29
8.4.7	Rotation Offset Analysis	29
8.5	XCR/ZUC-Based Key Schedule: Security Review and Analysis	30
8.5.1	XCR/ZUC-Driven Key Schedule for LOPC Wrapper: Clean Two-PRG Hybrid Reduction	31
8.6	Keyed Switching Layer as $T \circ L \circ S$: Minimal Facts Needed for the Reduction	31
8.7	Precise Key Mixing(Add / Subtract Round Key) Analysis	33
8.7.1	Algebraic Analysis	33
8.7.2	Explicit Matrix Analysis for 8-bit Model (Matrix-Exact, Piecewise- and Jacobian-Linear with Random-Walk Probabilistic Extension)	34
8.7.3	Computing the complexity of 8-bit ARX via matrix-analytic key mixing models	38
8.7.4	Generalizing the matrix picture to 64-bit ARX	38
8.8	Cryptographic Game Analysis of Encryption and Decryption Functions	39
9	Conclusion	41
A	NeoAlzette ARX S-box Analysis Python Code:	41
B	NeoAlzette MITL Differential characteristic search	48
C	NeoAlzette SAT SMT search	50
D	Statistical Tests for Randomness Assessment	51
E	About Git repository and run test	52
E.1	Contributions	52
F	Additional Quantum-Resistant OaldresPuzzle_Cryptic Block Cipher Algorithm	52
G	Documents referenced	52
H	Data Images	53
H.1	ASCON vs Little OaldresPuzzle Cryptic (Phases 1)	53
H.2	ASCON vs Little OaldresPuzzle Cryptic (Phases 2)	53
H.3	Chacha20 Statistical Test Phases	53
I	Data Tables	53

1 Introduction

This paper presents "Little OaldresPuzzle_Cryptic," a symmetric sequence cryptographic algorithm designed to meet the demands of our digital age. As data generation and consumption rates increase, so do the threats to data integrity and security. This necessitates cryptographic algorithms that can guarantee secure data exchange and high-speed performance.

Our algorithm uses the same cryptographic key for both encryption and decryption, providing significant speed advantages over asymmetric alternatives. The sequence cryptographic approach adds complexity to the encryption and decryption process, making it more challenging for unauthorized entities to interpret encrypted data.

The algorithm combines speed and security by utilizing a cryptographically secure pseudo-random number generator (CSPRNG) called "XorConstantRotation" (XCR), which produces highly unpredictable number sequences.

The remainder of this paper will examine the algorithm's mechanics, demonstrating how it utilizes mathematical constants, bitwise operations, and sequence-based cryptography to ensure both speed and security in data transmission.

2 XCR/Little OaldresPuzzle_Cryptic Stream Cipher Design Background

2.1 What are ARX Structure and Salsa20, ChaCha20 Algorithms?

ARX, or Addition/Bit-Rotation/Exclusive-OR, is a class of symmetric key algorithms constructed using the following simple operations: modulo addition, bit rotation, and Exclusive-OR. Unlike S-box-based designs, where the only nonlinear element is the substitution box (S-box), ARX designs rely on nonlinear hybrid functions such as addition and rotation [Ranea et al., 2022] [Liu et al., 2021]. These functions are easy to implement in both software and hardware and offer good diffusion and resistance to differential and linear cryptanalysis [Fei, 2012] [Aumasson et al., 2007]. There are some ECRYPT PowerPoint presentations that summarize this design structure [ARX-based Cryptography](#).

Salsa20 is a stream cipher algorithm proposed by Daniel J. Bernstein in 2005, which utilizes the ARX structure. Salsa20/8 and Salsa20/12 are two variants of Salsa20 that run 8 and 12 rounds of encryption, respectively [Bernstein, 2005] [Bernstein, 2008]. These algorithms were evaluated in the eSTREAM project and accepted as finalists in 2008 [Tsunoo et al., 2007]. Salsa20 was designed with a focus on simplicity and efficiency, and it is favored in the cryptography community for its speed and security [Maitra et al., 2015] [Ghafoori and Miyaji, 2022].

ChaCha20 [Bernstein et al., 2008a] [Bernstein et al., 2008b] is an ARX-based high-speed stream cipher proposed by Daniel J. Bernstein in 2008 as an improved version of Salsa20. ChaCha20 uses a 512-bit permutation function to convert a 512-bit input vector into a 512-bit output vector. The output vector is then added to the input vector to obtain a 512-bit keystream block. The input vector consists of four parts: a constant, a key, a counter, and a random number. The security of ChaCha20 relies on the complexity and irreversibility of the permutation function, as well as the randomness and uniqueness of the input vector. ChaCha20 has been utilized in a wide range of applications, including TLS v1.3, SSH, IPsec, and WireGuard [Serrano et al., 2022] [Cai et al., 2022].

2.2 About the XCR Structure of Our Lightweight Symmetric Encryption Technology

2.2.1 Designing XCR from the ARX Structure: Insights from NIST and Lightweight Cryptography

The National Institute of Standards and Technology (NIST) has underscored the critical need for lightweight encryption algorithms to secure the growing number of resource-constrained devices, particularly within the Internet of Things (IoT) ecosystem [NIST-LCS Website](#). These devices, such as sensors and RFID tags, necessitate encryption solutions that are both efficient and secure, a balance that traditional cryptographic methods struggle to achieve due to their computational intensity.

Our approach to designing the XCR structure is informed by insights from the seminal work "State of the Art in Lightweight Symmetric Cryptography" by Biryukov and Perrin [Biryukov and Perrin, 2017]. This paper outlines the design constraints and trends in lightweight symmetric cryptography, emphasizing the necessity for algorithms tailored to specific hardware and use cases. It highlights the crucial trade-offs among performance, security, and resource consumption, which are fundamental in developing lightweight encryption algorithms.

The paper specifies criteria for lightweight cryptographic algorithms in resource-constrained environments, advocating for a small block size, ideally 64 bits or less, and a minimum key size of 80 bits to balance security and efficiency. The round function should be simple, leveraging straightforward operations that are easily implemented on low-power devices. Additionally, a simple key scheduling mechanism is essential to prevent vulnerabilities and minimize complexity.

The risks of poorly implemented lightweight cryptography are underscored by the study "Speck-R: an ultra-lightweight encryption scheme for the Internet of Things" [Steem and Couturier, 2021]. This research and its references illustrate the severe consequences of inadequate lightweight encryption, including security breaches and denial-of-service attacks on small devices, serving as a cautionary tale about the importance of meticulous design and implementation.

Our "Little OaldresPuzzle_Cryptic" algorithm is developed with these considerations in mind, adhering to the standards set by contemporary lightweight cryptography research. It aims to deliver a robust and efficient solution for secure data transmission in IoT and other resource-constrained environments.

By leveraging the ARX structure, known for its simplicity and efficiency, we have designed the XCR to generate a sequence of pseudo-random numbers that are both unpredictable and computationally intensive. This structure ensures a consistent architecture adaptable across various device types, enabling faster encryption and decryption processes.

3 XCR Algorithm: Structure and Operations

Overview. XCR is an ARX (Addition, Rotation, Exclusive-OR) construction tailored for CSPRNG applications. It combines: (i) \oplus for rapid state mixing; (ii) constant injection via fixed, non-secret values derived from well-known irrational constants; and (iii) bit rotations for diffusion. The design emphasizes simplicity (word-oriented, branch-free operations) while aiming for statistical unpredictability suitable for keystream generation and authentication tags.

3.1 Preliminaries and Notations

All words are 64-bit unless stated otherwise, i.e., elements of $\{0, 1\}^{64}$ interpreted as unsigned integers. We write $\langle a, i \rangle$ for the i -th least significant bit of a (LSB indexed by $i = 1$).

- $x, y, \text{state} \in \{0, 1\}^{64}$: internal registers; unless stated, initialized to 0.
- $\text{number_once} \in \{0, 1\}^{64}$: per-round selector derived deterministically from key/nonce/**counter**.
- **ROUND_CONSTANTS**: a fixed set of 64-bit constants ("nothing-up-my-sleeve" values).
- $I, O \in \{0, 1\}^{64}$: input and output words, with $O = I \oplus y$ after the round transformation.
- **counter**: iteration index controlling selection within **ROUND_CONSTANTS**.

3.1.1 Operators

- $a \wedge b, a \vee b$: bitwise AND/OR.
- $\neg a$: bitwise NOT.
- $a \oplus b$: bitwise XOR.
- $a \ominus b := \neg a \oplus b$.
- $a \boxplus_n b := (a + b) \bmod 2^n, a \boxminus_n b := (a - b) \bmod 2^n$.
- $a \lll_n r, a \ggg_n r$: left/right rotation by r bits with word size n ($0 \leq r < n$).
- $a \ll_n s, a \gg_n s$: logical left/right shift by s bits with word size n ($0 \leq s < n$).

3.2 Overall Steps

1. **State Initialization Function**: set x, y, state , and **counter**.
2. **Round Constant Selection**: deterministically select $RC0, RC1, RC2 \in \text{ROUND_CONSTANTS}$ from $(\text{number_once}, \text{counter}, \text{state})$.
3. **State Iteration / Update**: apply the diffusion and confusion layers to (x, y) and update **state**.

4. **Counter Increment:** $\text{counter} \leftarrow \text{counter} + 1 \bmod 2^{64}$.
5. **Output:** $O \leftarrow I \oplus y$.

3.3 Diffusion and Confusion Layers

XCR separates mixing into two complementary phases:

- **Diffusion Layer.** Word-wise rotations and modular additions propagate single-bit changes across many bit positions in both x and y within a few rounds.
- **Confusion Layer.** XOR with round constants and cross-word feeds (x into y and vice versa) obscure linear structure and bias.

3.4 Selection of Rotation Amounts

Rotation pairs (r_0, r_1) are chosen to avoid short linear dependencies and to ensure rapid bit spreading. In particular, pairs such as $(1, 63)$ are complementary on 64-bit words. More generally, for a schedule (r_i) we require co-primeness across adjacent (and preferably small sliding windows):

$$\gcd(r_i, r_{i+1}) = 1, \quad \gcd(r_i, r_{i+1}, r_{i+2}) = 1, \dots$$

so that no subset of bit positions remains invariant under the induced rotation group.

3.5 Algorithm Efficiency and Security

ARX operations map efficiently to commodity CPUs (single-cycle XOR, addition; constant-time rotations). The absence of table lookups reduces side-channel surface. Security intuition follows from: (i) constant injection (nothing-up-my-sleeve RCj) preventing slide/structural symmetries; (ii) co-prime rotations yielding fast avalanche; (iii) cross-feeding $x \leftrightarrow y$ suppressing linear trails. Empirical tests (e.g., avalanche and bias checks) should complement formal analysis of differential/linear resistance.

4 Detailed Description of the XCR Algorithm Structure

A single XCR update (“round”) acts on (x, y) with selected $RC0, RC1, RC2$ and rotation amounts (r_0, r_1) :

$$\begin{aligned} x &\leftarrow x \oplus RC0, \\ x &\leftarrow x \boxplus_{64} (y \oplus RC1), \\ x &\leftarrow x \lll_{64} r_0, \\ y &\leftarrow y \oplus x, \\ y &\leftarrow y \boxplus_{64} RC2, \\ y &\leftarrow y \ggg_{64} r_1. \end{aligned}$$

After a prescribed number of iterations (as determined by `number_once` and `counter`), the output word is

$$O = I \oplus y.$$

Equivalently, writing **XCR**(`number_once`) for the above state transformation,

$$O = \mathbf{PRG}(I, \text{number_once}) = I \oplus \mathbf{XCR}(\text{number_once}) = I \oplus y.$$

4.1 XCR Algorithm: State Initialization Function

The procedure `StateInitialize(seed → random)` follows a GGM-style expansion with whitening. Let κ denote the security parameter (words remain 64-bit).

Definition 4.1 (State Initialization). *Given `seed` ∈ {0, 1} $^\kappa$, perform:*

1. *Register Initialization:*

$$\text{counter} \leftarrow 0, \quad \text{state} \leftarrow \begin{cases} 1, & \text{if } \text{state} = 0, \\ \text{state}, & \text{otherwise.} \end{cases}$$

2. *State Anchoring:* $(\text{state}_0, \text{random}) \leftarrow (\text{state}, \text{state})$.

3. *GGM-Style Expansion (4 rounds):*

For $r = 0, 1, 2, 3$, update $\text{random} \leftarrow \mathcal{G}_r(\text{random})$, where each $\mathcal{G}_r : \{0, 1\}^\kappa \rightarrow \{0, 1\}^\kappa$ iterates the XCR core to extract bits:

(a) Set $\text{next_random} \leftarrow 0^\kappa$.

(b) For $i = 1$ to κ :

$\text{random} \leftarrow \text{StateIterate}(\text{random}), \quad b \leftarrow \langle \text{random}, 1 \rangle, \quad \text{next_random} \leftarrow (\text{next_random} \ll \kappa 1) \vee (1 - b)$.

(c) Set $\text{random} \leftarrow \text{next_random}$.

4. *Whitening:* $\text{state} \leftarrow \text{state} \oplus (\text{state}_0 \boxplus_{64} \text{random})$.

GGM Tree View. Let $G : \{0, 1\}^\kappa \rightarrow \{0, 1\}^{2\kappa}$ be a PRG with κ -bit branches $G_0(x), G_1(x)$. For depth $d = 4$ and $x \in \{0, 1\}^\kappa$,

$$\mathcal{G}(x) := \bigoplus_{r=0}^{d-1} G_{b_r}(x_r), \quad b_r = \langle x_r, 1 \rangle, \quad x_{r+1} = G_{b_r}(x_r) \bmod 2^\kappa.$$

This realizes a forward-iterable expansion in which disclosure of a later node does not reveal prior nodes, while the final whitening step ties the internal state to both its anchored and expanded images.

4.2 XCR Algorithm: State Iteration / Update Function

Definition 4.2 (State Iteration / Update Function — Diffusion–Confusion Layer Specification). *Let $\text{StateIteration}(\text{nonce}) : \{0, 1\}^{64} \circ \{0, 1\}^{64} \circ \{0, 1\}^{64} \circ \{0, 1\}^{64} \rightarrow \{0, 1\}^{64}$ denote the core transformation operating on a 64-bit state s and nonce n . The function comprises:*

- Rotational constants $\mathbf{r} = (r_1, r_2, r_3, r_4) = (7, 19, 32, 47)$ satisfying

$$\gcd(r_i, r_{i+1}) = 1 \quad \text{and} \quad \gcd(r_i, r_{i+1}, r_{i+2}) = 1 \quad \forall i \in \{1, 2\}.$$

- Round constants $\text{RC}[\cdot]$ with $|\text{RC}| = 300$ (pseudo-random numbers computed by nonlinear functions).

4.2.1 Dynamic Constant Generation

Define three nonce-dependent round constants:

$$\begin{aligned} \text{RC0} &\leftarrow \text{RC}[n \bmod |\text{RC}|], \\ \text{RC1} &\leftarrow \text{RC}[(c + n) \bmod |\text{RC}|], \\ \text{RC2} &\leftarrow \text{RC}[s \bmod |\text{RC}|], \end{aligned}$$

where c denotes the counter register and s the current state.

4.2.2 Diffusion Layer Operations

- 1: Initialize $x, y \leftarrow 0^{64}$
- 2: **if** $x = 0$ **then**
- 3: $x \leftarrow \text{RC0}$
- 4: **else**
- 5: $y \leftarrow y \oplus (x \lll r_2) \oplus (x \lll r_3)$
- 6: $s \leftarrow s \oplus [(y \lll r_3) \oplus (y \lll r_4) \oplus (y \lll 63)] \oplus c$
- 7: $x \leftarrow x \oplus [(s \lll r_1) \oplus (s \lll r_2) \oplus \text{RC0} \oplus n]$
- 8: **end if**

4.2.3 Confusion Layer Operations

$$\begin{aligned} s &\leftarrow s \boxplus_{64} (y \oplus (y \ggg 1) \oplus \text{RC0}), \\ x &\leftarrow x \oplus [s \boxplus_{64} (s \ggg 1) \boxplus_{64} \text{RC1}], \\ y &\leftarrow y \boxplus_{64} (x \oplus (x \ggg 1) \oplus \text{RC2}). \end{aligned}$$

4.2.4 Update Rules

- Counter increment: $c \leftarrow c + 1 \bmod 2^{64}$.
- Keystream output: $z \leftarrow y$.

Lemma 4.1 (Diffusion Properties). *With $\mathbf{r} = (7, 19, 32, 47)$, the iteration achieves full diffusion within 3 rounds under the avalanche criterion:*

$$\forall \Delta \in \{0, 1\}^{64} \setminus \{0\} : \text{HW}(\Delta) \geq 1 \Rightarrow \mathbb{E}[\text{HW}(\text{Stateliterate}(s \oplus \Delta) \oplus \text{Stateliterate}(s))] \geq 32,$$

where HW denotes Hamming weight.

Proof sketch.

1. Pairwise coprime rotations yield a large effective branch number and prevent short rotational cycles.
2. Inclusion of a 63-bit rotation (complementary to 1) forces cross-word bit migration.
3. Nonlinear modular additions (\boxplus_{64}) disrupt linear XOR trails and attenuate bias accumulation.

□

Remark 4.1. The structure combines rotational ARX principles [Khovratovich et al., 2015] with the diffusion-confusion philosophy of substitution-permutation networks [Shannon, 1949], while the counter-dependent constant generation enforces round asymmetry over a 2^{64} -state space.

5 NeoAlzette Substitution Box

The *NeoAlzette* substitution box (S-box), a central primitive within the **Little OaldresPuzzle_Cryptic** construction, refines the ARX-based design lineage of Alzette [Beierle et al., 2019] while preserving its word-oriented efficiency. Engineered for paired 32-bit words, the S-box layers modular additions, rotations, and XORs to attain stronger statistical behavior (nonlinearity, bias suppression, and rapid avalanche) without sacrificing constant-time implementability. In comparative testing under identical settings, the original Alzette trails *NeoAlzette* on several indicators of resistance to linear and differential patterns, underscoring the benefits of the added diffusion/mixing phases.

The design incorporates:

- multi-phase round-constant injection with Fibonacci/irrational blending,
- rotational asymmetry via prime-derived offsets, and
- bidirectional diffusion through cross-coupled modular additions.

Definition 5.1 (NeoAlzette S-box Operation). *Let $(a, b) \in (\mathbb{F}_2^{32})^2$ be the input word pair. The forward map $(a', b') = \text{NeoAlzette_ForwardLayer}(a, b, rc)$ consists of four diffusion phases:*

1. **Vertical Mixing.** XOR then rotated-sum:

$$b^{(1)} = b \oplus a, \quad a^{(1)} = (a \boxplus_{64} b^{(1)}) \ggg 31 \oplus rc_1.$$

2. **Cross Feedback.** Staggered constant addition after left rotation:

$$a^{(2)} = (a^{(1)} \oplus b^{(1)}) \lll 24 \boxplus_{64} rc_2.$$

3. **Diagonal Branching.** Parallel rotation and constant injection:

$$b^{(2)} = (b^{(1)} \lll 8) \oplus rc_3, \quad a^{(3)} = a^{(2)} \boxplus_{64} b^{(2)}.$$

4. **Convergence Layer.** Asymmetric rotation with final mixing:

$$b' = (a^{(3)} \boxplus_{64} b^{(2)}) \ggg 17 \oplus rc_4, \quad a' = a^{(3)} \oplus b'.$$

5.1 Round Constant Derivation

The 16-entry round-constant array is assembled from four sources as follows.

1) Fibonacci Concatenation. Let F_k be the k -th Fibonacci number. Concatenate F_1, \dots, F_{20} in decimal, then parse into 32-bit words:

$$\begin{aligned} F_1 \| F_2 \| \cdots \| F_{20} &= 123581321345589144233377610987159725844181 \\ &\Downarrow \text{Hex grouping} \\ rc_{0:3} &= [0x16B2C40B, 0xC117176A, 0x0F9A2598, 0xA1563ACA]. \end{aligned}$$

2) Irrational Constants. Let $\pi^{(n)}, \phi^{(n)}, e^{(n)}$ denote the first 128 fractional bits of π, ϕ, e respectively; extract 4 consecutive 32-bit words:

$$\begin{aligned} \pi_{3:0} &= \lfloor 2^{128}\pi \rfloor \triangleright 32 = [0x243F6A88, 0x85A308D3, 0x13198102, 0xE0370734], \\ \phi_{3:0} &= \lfloor 2^{128}\phi \rfloor \triangleright 32 = [0x9E3779B9, 0x7F4A7C15, 0xF39CC060, 0x5CEDC834], \\ e_{3:0} &= \lfloor 2^{128}e \rfloor \triangleright 32 = [0xB7E15162, 0x8AED2A6A, 0xBF715880, 0x9CF4F3C7]. \end{aligned}$$

3) Composite Array. Interleave the sequences to form:

$$rc[0:15] = [\underbrace{rc_0, \dots, rc_3}_{\text{Fibonacci}}, \underbrace{\pi_0, \dots, \pi_3}_{\pi}, \underbrace{\phi_0, \dots, \phi_3}_{\phi}, \underbrace{e_0, \dots, e_3}_e].$$

5.2 Security Parameterization

The following indicative metrics summarize the attained properties:

Metric	NeoAlzette	Alzette [Beierle et al., 2019]
Differential Uniformity	2^{-64}	2^{-32}
Linear Bias Bound	2^{-32}	2^{-16}
Full Diffusion Rounds	2	4
Minimum Nonlinearity	28	18

Contributors to the improvements include:

- 4× constant-injection points per round (vs. Alzette),
- 56-bit effective rotation diversity (vs. Alzette’s 24-bit span),
- dual modular-addition branches per phase.

Constants are injected at designated rotation phases to:

1. disrupt slide symmetries via aperiodic (Fibonacci) offsets,
2. hinder rotational cryptanalysis using irrational bit patterns,
3. mitigate constant-correlation through multi-source blending.

```

1 constexpr std::array<std::uint32_t, 16> ROUND_CONSTANT
2 {
3     // 1,2,3,5,8,13,21,34,55,89,144,233,377,610,987,1597,2584,4181 (Fibonacci)
4     // Concatenation: 123581321345589144233377610987159725844181
5     // Hex: 16b2c40bc117176a0f9a2598a1563aca6d5
6     0x16B2C40B, 0xC117176A, 0x0F9A2598, 0xA1563ACA,
7
8     /*
9      Mathematical Constants - Millions of Digits
10     http://www.numberworld.org/constants.html
11 */
12
13     // (3.243f6a8885a308d313198a2e0370734)
14     0x243F6A88, 0x85A308D3, 0x13198102, 0xE0370734,
15     // (1.9e3779b97f4a7c15f39cc0605cedc834)
16     0x9E3779B9, 0x7F4A7C15, 0xF39CC060, 0x5CEDC834,
17     // e (2.b7e151628aed2a6abf7158809cf4f3c7)
18     0xB7E15162, 0x8AED2A6A, 0xBF715880, 0x9CF4F3C7
19 };

```

5.3 Forward and Backward Layers

The S-box exposes forward and inverse layers to support encryption and decryption.

Algorithm 1 NeoAlzette ARX S-box Layers

```

1: function NEOALZETTE_FORWARDLAYER(a, b, rc)
2:   b ← b ⊕ a
3:   a ← (a ⊕32 b) ≫ 31
4:   a ← a ⊕ rc
5:   b ← b ⊕32 a
6:   a ← (a ⊕ b) ≪ 24
7:   a ← a ⊕32 rc
8:   b ← (b ≪ 8) ⊕ rc
9:   a ← a ⊕32 b
10:  a ← a ⊕ b
11:  b ← (a ⊕32 b) ≫ 17
12:  b ← b ⊕ rc
13:  a ← a ⊕32 b
14:  b ← (a ⊕ b) ≪ 16
15:  b ← b ⊕32 rc
16:  return a, b
17: end function

18: function NEOALZETTE_BACKWARDLAYER(a, b, rc)
19:   b ← b ⊖32 rc
20:   b ← (b ≫ 16) ⊕ a
21:   a ← a ⊖32 b
22:   b ← b ⊕ rc
23:   b ← (b ≪ 17) ⊖32 a
24:   a ← a ⊕ b
25:   a ← a ⊖32 b
26:   b ← (b ⊕ rc) ≫ 8
27:   a ← a ⊖32 rc
28:   a ← (a ≫ 24) ⊕ b
29:   b ← b ⊖32 a
30:   a ← a ⊕ rc
31:   a ← (a ≪ 31) ⊖32 b
32:   b ← b ⊕ a
33:   return a, b
34: end function

```

5.4 Operational Modes

Forward Mode. Realizes S_{rc} using 12 ARX primitives:

- 4 modular additions with round constants,
- 5 interleaved XORs,
- 5 rotations (right by 31, 17; left by 24, 8, 16).

Backward Mode. Computes S_{rc}^{-1} by precisely inverting operation order, rotation directions, and constant injections via \ominus_{32} .

5.5 Applied Little OaldresPuzzle_Cryptic Algorithm: Encryption and Decryption

Within **Little OaldresPuzzle_Cryptic**, multiple rounds of the forward (resp. backward) layer are invoked per block for encryption (resp. decryption). Each round consumes a distinct entry of **ROUND_CONSTANT**, with the total number of rounds fixed by the global parameter to balance security and throughput.

6 Little OaldresPuzzle_Cryptic Algorithm Overview

Little OaldresPuzzle_Cryptic is a symmetric, word-oriented block cipher operating on 128-bit blocks (two 64-bit lanes). It follows the ARX paradigm and applies r rounds of bijective transformations arranged as $\Pi \rightarrow \Theta \rightarrow \Gamma$. In each round, Π provides non-linear diffusion via the NeoAlzette S-box over four 32-bit words, Θ applies key-dependent linear/mixing steps (including a per-lane binary pseudo-Hadamard transform and a bit-tweak), and Γ performs key mixing with modular addition and rotations. Decryption mirrors encryption by applying the round inverses in reverse order, with the key schedule reproducing the exact per-round state (subkeys, selectors, rotation amounts, and constant index). Round constants are selected from `ROUND_CONSTANTS` (Section 5.1) in a deterministic cyclic fashion.

6.1 Symmetric Encryption Specification

Let \mathbb{B}^{128} be the 128-bit data space and \mathcal{K} the key space. For $P \in \mathbb{B}^{128}$ and $K \in \mathcal{K}$, the encryption is the composition

$$C = E(P, K) = (\Gamma_{r-1} \circ \Theta_{r-1} \circ \Pi_{r-1}) \circ \dots \circ (\Gamma_0 \circ \Theta_0 \circ \Pi_0)(P).$$

```

1: KeyState  $\triangleq \{ sk = (sk^{(0)}, sk^{(1)}), cf, \alpha \ll\gg, \beta \ll\gg, index \}$ 
2: Initialize  $DataState^{(0)} \leftarrow P$ ;  $KeySchedule(K) \rightarrow \{KeyState_i\}_{i=0}^{r-1}$ 
3: for  $i \leftarrow 0$  to  $r-1$  do
4:    $DataState^{(i+1)} \leftarrow \Gamma_i(\Theta_i(\Pi_i(DataState^{(i)}; index_i)))$ 
5: end for
6: return  $C \leftarrow DataState^{(r)}$ 
```

Note: The key schedule deterministically reproduces $KeyState_i$ for all i , including $index_i \equiv i \bmod 16$ for the `ROUND_CONSTANTS` lookup.

6.2 Symmetric Decryption Specification

For $C \in \mathbb{B}^{128}$ and $K \in \mathcal{K}$,

$$P = D(C, K) = (\Pi_0^{-1} \circ \Theta_0^{-1} \circ \Gamma_0^{-1}) \circ \dots \circ (\Pi_{r-1}^{-1} \circ \Theta_{r-1}^{-1} \circ \Gamma_{r-1}^{-1})(C).$$

```

1: KeyState  $\triangleq \{ sk, cf, \alpha \ll\gg, \beta \ll\gg, index \}$ 
2: Initialize  $DataState^{(r)} \leftarrow C$ ;  $KeySchedule(K) \rightarrow \{KeyState_i\}_{i=0}^{r-1}$ 
3: for  $i \leftarrow r-1$  down to 0 do ▷ Reverse round order
4:    $DataState^{(i)} \leftarrow \Pi_i^{-1}(\Theta_i^{-1}(\Gamma_i^{-1}(DataState^{(i+1)})); index_i)$ 
5: end for
6: return  $P \leftarrow DataState^{(0)}$ 
```

6.2.1 Little OaldresPuzzle_Cryptic Components

Let the 128-bit state be stored as two 64-bit lanes `lane0`, `lane1`. Define the packing

$$\text{lane0} = (w_0 \parallel w_1), \quad \text{lane1} = (w_2 \parallel w_3),$$

where each $w_j \in \mathbb{B}^{32}$ and $(\cdot \parallel \cdot)$ denotes $(\text{hi} \parallel \text{lo})$.

1. NeoAlzette S-box Layer Π (diagonal pairing). The bijection acts on the four 32-bit words by two parallel NeoAlzette calls with the round constant rc_{index} :

$$(w_0, w_2) \leftarrow \text{NeoAlzette_ForwardLayer}(w_0, w_2, rc_{index}), \quad (w_1, w_3) \leftarrow \text{NeoAlzette_ForwardLayer}(w_1, w_3, rc_{index}).$$

Repacking yields updated `lane0`, `lane1`. For decryption,

$$(w_1, w_3) \leftarrow \text{NeoAlzette_BackwardLayer}(w_1, w_3, rc_{index}), \quad (w_0, w_2) \leftarrow \text{NeoAlzette_BackwardLayer}(w_0, w_2, rc_{index}),$$

followed by repacking. (See Section 5.)

2. Mix Linear Transform Layer Θ . On each 64-bit lane $x \in \mathbb{B}^{64}$ with selector cf_i , rotation amount $\beta_i \in [0, 63]$, and subkey half $sk_i^{(\ell)}$:

$$\Theta_i(x) = \begin{cases} x \oplus sk_i^{(\ell)} & \text{if } cf_i = 0, \\ (\neg x) \oplus sk_i^{(\ell)} & \text{if } cf_i = 1, \\ x \lll_{64} \beta_i & \text{if } cf_i = 2, \\ x \ggg_{64} \beta_i & \text{if } cf_i = 3. \end{cases}$$

Then (on 32-bit halves) apply the binary pseudo-Hadamard transform independently:

$$(w_0, w_1) \mapsto (w_0 \oplus w_1, w_0 \oplus (w_1 \ll_{32} 1)), \quad (w_2, w_3) \mapsto (w_2 \oplus w_3, w_2 \oplus (w_3 \ll_{32} 1)),$$

and repack. Finally apply the nonlinear *Random Bit Tweak* using α_i :

$$\text{lane0} \leftarrow \text{lane0} \oplus 2^{\alpha_i \bmod 64}, \quad \text{lane1} \leftarrow \text{lane1} \oplus 2^{63 - (\alpha_i \bmod 64)}.$$

Its inverse uses the obvious reversals:

$$\Theta_i^{-1}(x) = \begin{cases} x \oplus sk_i^{(\ell)} & cf_i = 0, \\ (\neg x) \oplus sk_i^{(\ell)} & cf_i = 1, \\ x \ggg_{64} \beta_i & cf_i = 2, \\ x \lll_{64} \beta_i & cf_i = 3, \end{cases}$$

bit-tweak reapplied (same mask), and the inverse PHT on (w_2, w_3) and (w_0, w_1) respectively.

3. Key Mixing with ARX Γ . Let lane0 use k_0 and $sk^{(0)}$, and lane1 use k_1 and $sk^{(1)}$. The forward map on a lane x is

$$\Gamma_i(x) = ((x \boxplus_{64} (k \oplus sk)) \oplus k) \ggg_{64} 16 \oplus ((k \boxplus_{64} sk) \lll_{64} 48),$$

and the inverse is

$$\Gamma_i^{-1}(y) = \left(((y \oplus ((k \boxplus_{64} sk) \lll_{64} 48)) \lll_{64} 16) \oplus k \right) \boxplus_{64} (k \oplus sk).$$

Key Schedule (overview). For each round i , the schedule produces KeyState_i with:

- $sk_i^{(0)} = k_0 \oplus \text{key_prng}(\text{number_once})$, $sk_i^{(1)} = k_1 \oplus \text{prng}(\text{number_once} \oplus i)$;
- $cf_i = \text{key_prng}(sk_i^{(0)} \oplus (k_0 \gg_{64} 1)) \wedge 3$;
- β_i, α_i obtained from $\text{prng}(sk_i^{(1)} \oplus cf_i)$ as two 6-bit amounts;
- $index_i \equiv i \bmod 16$ for ROUND_CONSTANTS .

Here `key_prng` is seeded from a key-derived 64-bit `seed` and `prng` from the cipher instance seed; both advance with `number_once`.

6.3 XCR CSPRNG-based Key Schedule

The key schedule is driven by two XCR instances: a *key-seeded* generator and an *instance-seeded* generator. Let the 128-bit key be $Key = (Key_L, Key_R) \in (\{0, 1\}^{64})^2$. The key-seeded XCR is initialized from

$$a \leftarrow Key_L \oplus Key_R, \quad b \leftarrow \neg \text{ghash_multiply}(Key_L, Key_R),$$

$$\text{mix}(a, b) \text{ (binary PHT)}, \quad seed \leftarrow (a \oplus (b \lll_{64} 1)) \oplus (a \lll_{64} 13),$$

and then $\mathcal{G}_{\text{XCR}}^{\text{key}} \leftarrow \text{XCR}.\text{Init}(seed)$. The instance generator $\mathcal{G}_{\text{XCR}}^{\text{sys}}$ is the class member PRNG seeded by the cipher's global 64-bit `seed`. Here `ghash_multiply` denotes carry-less multiplication in $\mathbb{F}_{2^{64}}$ modulo $x^{64} + x^4 + x^3 + x + 1$.

The schedule proceeds strictly sequentially as:

```

Step 0:  $rc\_index^{raw} \leftarrow 0$ 
Step 1:  $\forall i \in [0, \text{rounds}] : \text{KeyState}_i \triangleq \{ sk, cf, \alpha \lll \gg, \beta \lll \gg, index \}$ 
Step 2:  $sk \equiv (sk^{(0)}, sk^{(1)})$ ,  $\begin{cases} sk^{(0)} \leftarrow Key_L \oplus \mathcal{G}_{\text{XCR}}^{\text{key}}(\text{number\_once}), \\ sk^{(1)} \leftarrow Key_R \oplus \mathcal{G}_{\text{XCR}}^{\text{sys}}(\text{number\_once} \oplus i), \end{cases}$  (two-limb ephemeral subkey)
Step 3:  $cf \leftarrow \mathcal{G}_{\text{XCR}}^{\text{key}}(sk^{(0)} \oplus (Key_L \gg_{64} 1))$  (key-feedback selector)
Step 4:  $\alpha_{\lll \gg}^{raw} \leftarrow \mathcal{G}_{\text{XCR}}^{\text{sys}}(sk^{(1)} \oplus (cf \bmod 4))$  (bijective mixing)
Step 5:  $\beta_{\lll \gg} \leftarrow (\alpha_{\lll \gg}^{raw} \gg 6) \bmod 64$  (high 6 bits for  $\beta$ )
Step 6:  $\alpha_{\lll \gg} \leftarrow \alpha_{\lll \gg}^{raw} \bmod 64$  (constrain  $\alpha$  to  $[0, 63]$ )
Step 7:  $cf \leftarrow cf \bmod 4$  (limit selector to 2 bits)
Step 8:  $index \leftarrow (rc\_index^{raw} \gg 1) \bmod 16$  (NeoAlzette constant selector)
Step 9:  $rc\_index^{raw} \leftarrow rc\_index^{raw} + 2$  (advance with stride 2).

```

- **Step 2** yields a pairwise subkey composed of a key-masked limb from $\mathcal{G}_{\text{XCR}}^{\text{key}}$ and a round-salted limb from $\mathcal{G}_{\text{XCR}}^{\text{sys}}$ ($\text{salt} = \text{number_once} \oplus i$).
- **Steps 3–6** implement the confusion–diffusion cascade using XCR outputs, rotations, and modular reduction for bit-width control.
- **Steps 8–9** realize the round-constant progression $rc_index^{raw} : 0 \rightarrow 2 \rightarrow 4 \rightarrow \dots \rightarrow 2(\text{rounds} - 1)$ with $index$ taken after a right shift and masked to 16 entries.

6.4 XCR Round Constant Generation

Let $\{a\} \triangleq a - \lfloor a \rfloor$ denote the fractional-part operator. The array `ROUND_CONSTANTS` consists of $|\text{ROUND_CONSTANTS}| = 300$ words in $\{0, 1\}^{64}$ constructed as:

1. **Seed quartet.** Prepend four public, fixed 64-bit words derived from (i) Fibonacci concatenation, (ii) π , (iii) $\phi = \frac{1+\sqrt{5}}{2}$, and (iv) e (least-significant 64 bits of their standard “nothing-up-my-sleeve” encodings).
2. **Transcendental stream.** For consecutive integers $x = 1, 2, \dots$ evaluate

$$f(x) = \underbrace{(e^x - \cos(\pi x))}_{\text{exp-osc}} \cdot \underbrace{(\phi x^2 - \phi x - 1)}_{\text{quadratic}} \cdot \prod_{k \in \{\sqrt{2}, \sqrt{3}, \delta, \rho\}} \{kx\} \cdot \ln(1+x),$$

where δ is Feigenbaum's constant and ρ the plastic number. Append

$$\text{RC_next} \leftarrow \lfloor \{f(x)\} \cdot 2^{64} \rfloor$$

until the array is filled (together with the seed quartet, totaling 300 words).

Round-time selection for XCR. At call `StateIteration(number_once)`, dynamically select

$$RC0 \leftarrow RC[\text{number_once} \bmod |\text{RC}|], \quad RC1 \leftarrow RC[(\text{counter} + \text{number_once}) \bmod |\text{RC}|], \quad RC2 \leftarrow RC[\text{state} \bmod |\text{RC}|].$$

Algorithm 4 Generator — Computing ROUND_CONSTANTS (64-bit words)

```

1: function GENERATEROUNDCONSTANTS( $N=300$ )
2:    $S \leftarrow [\text{Fibonacci64}, \pi\cdot64, \phi\cdot64, e\cdot64]$                                  $\triangleright$  fixed  $4 \times 64$ -bit seeds
3:    $x \leftarrow 1$ 
4:   while  $|S| < N$  do
5:      $r \leftarrow (e^x - \cos(\pi x)) \cdot (\phi x^2 - \phi x - 1) \cdot \{\sqrt{2}x\} \cdot \{\sqrt{3}x\} \cdot \ln(1+x) \cdot \{x\delta\} \cdot \{x\rho\}$ 
6:      $w \leftarrow \lfloor(r - \lfloor r \rfloor) \cdot 2^{64} \rfloor$ 
7:      $S.\text{APPEND}(w); \quad x \leftarrow x + 1$ 
8:   end while
9:   return  $S$ 
10:  end function

```

Algorithm 5 Cryptographically Secure Pseudo-Random Number Generator — XorConstantRotation (XCR)

```

1: function XCR INITIALIZE( $seed$ )
2:    $x, y, counter \leftarrow 0; \quad state \leftarrow seed;$  if  $state = 0$  then  $state \leftarrow 1$ 
3:    $state\_0 \leftarrow state; \quad random \leftarrow state$ 
4:   for  $round \leftarrow 0$  to  $3$  do                                               $\triangleright$  4-round GGM-style expansion
5:      $next\_random \leftarrow 0$ 
6:     for  $bit\_index \leftarrow 0$  to  $63$  do
7:        $random \leftarrow \text{XCR GENERATION}(random)$ 
8:       if  $random \wedge 1 = 1$  then
9:          $next\_random \leftarrow (next\_random \ll_{64} 1) \vee 0$ 
10:      else
11:         $next\_random \leftarrow (next\_random \ll_{64} 1) \vee 1$ 
12:      end if
13:    end for
14:     $random \leftarrow next\_random$ 
15:  end for
16:   $state \leftarrow state \oplus (state\_0 \boxplus_{64} random)$                                                $\triangleright$  whitening
17: end function

18: function XCR GENERATION( $number\_once$ )
19:    $RC0, RC1, RC2 \leftarrow \text{Dynamic Round Constant Selection}$                                  $\triangleright$  as above
20:   if  $x = 0$  then
21:      $x \leftarrow RC0$ 
22:   else
23:      $y \leftarrow y \oplus (x \ll_{64} 19) \oplus (x \ll_{64} 32)$ 
24:      $state \leftarrow state \oplus (y \ll_{64} 32) \oplus (y \ll_{64} 47) \oplus (y \ll_{64} 63) \oplus counter$ 
25:      $x \leftarrow x \oplus (state \ll_{64} 7) \oplus (state \ll_{64} 19) \oplus RC0 \oplus number\_once$ 
26:   end if
27:    $state \leftarrow state \boxplus_{64} (y \oplus (y \gg_{64} 1) \oplus RC0)$ 
28:    $x \leftarrow x \oplus (state \boxplus_{64} (state \gg_{64} 1) \boxplus_{64} RC1)$ 
29:    $y \leftarrow y \boxplus_{64} (x \oplus (x \gg_{64} 1) \oplus RC2)$ 
30:    $counter \leftarrow counter + 1$ 
31:   return  $y$ 
32: end function

```

Algorithm 6 LittleOaldresPuzzle_Cryptic Class Implementation (Aligned with XCR CSPRNG and LOPC Wrapper)

```

1:  $seed \leftarrow \text{initial 64-bit seed}$ 
2:  $\text{prng} \leftarrow \text{XorConstantRotation}(seed)$ 
3:  $rounds \in \{4, 8, 16, 32, 64, \dots\}$ 
4: function KEYSTATE( $subkey, choice\_function, bit\_rotation\_amount\_a, bit\_rotation\_amount\_b, constant\_index$ )
5:   return  $\{subkey = (subkey^{(0)}, subkey^{(1)}), choice\_function, bit\_rotation\_amount\_a, bit\_rotation\_amount\_b, round\_constant\_index = constant\_index\}$ 
6: end function
7: function GENERATEANDSTOREKEYSTATES( $Key = (k_0, k_1), number\_once$ )
8:    $a \leftarrow k_0 \oplus k_1, \quad b \leftarrow \neg \text{GHASH-MULTIPLY}(k_0, k_1)$                                  $\triangleright$  GF(2) carry-less product, polynomial  $x^{64} + x^4 + x^3 + x + 1$ 
9:    $\text{MIX}(a, b)$                                                                 $\triangleright$  Binary Pseudo-Hadamard Transform:  $(a', b') = (a \oplus b, a \oplus (b \ll 1))$ 
10:   $seed^* \leftarrow (a \oplus (b \ll_{64} 1)) \oplus (a \ll_{64} 13)$ 
11:   $\text{key\_prng} \leftarrow \text{XorConstantRotation}(seed^*)$ 
12:   $rc\_raw \leftarrow 0$ 
13:  for  $round \leftarrow 0$  to  $rounds - 1$  do
14:     $ks \leftarrow \text{KeyState}[round]$ 
15:     $sk0 \leftarrow \text{key\_prng}(number\_once)$ 
16:     $para \leftarrow \text{prng}(number\_once \oplus round)$ 
17:     $ks.\text{subkey}^{(0)} \leftarrow k_0 \oplus sk0, \quad ks.\text{subkey}^{(1)} \leftarrow k_1 \oplus para$ 
18:     $ks.\text{choice\_function} \leftarrow \text{key\_prng}(ks.\text{subkey}^{(0)} \oplus (k_0 \gg 1)) \wedge 3$ 
19:     $br \leftarrow \text{prng}(ks.\text{subkey}^{(1)} \oplus ks.\text{choice\_function})$ 
20:     $ks.\text{bit\_rotation\_amount\_a} \leftarrow br \bmod 64, \quad ks.\text{bit\_rotation\_amount\_b} \leftarrow (br \gg 6) \bmod 64$ 
21:     $ks.\text{round\_constant\_index} \leftarrow (rc\_raw \gg 1) \bmod 16, \quad rc\_raw \leftarrow rc\_raw + 2$ 
22:  end for
23: end function
24: function ADDROUNDKEY( $v, k, sk$ )                                               $\triangleright v, k, sk \in \{0, 1\}^{64}$ 
25:    $v \leftarrow v \boxplus_{64} (k \oplus sk)$ 
26:    $v \leftarrow (v \oplus k) \gg_{64} 16$ 

```

```

27:    $v \leftarrow v \oplus ((k \boxplus_{64} sk) \lll_{64} 48)$ 
28: end function
29: function SUBTRACTROUNDKEY( $v, k, sk$ )
30:    $v \leftarrow v \oplus ((k \boxminus_{64} sk) \lll_{64} 48)$ 
31:    $v \leftarrow (v \lll_{64} 16) \oplus k$ 
32:    $v \leftarrow v \boxminus_{64} (k \oplus sk)$ 
33: end function
34: function BITSPLIT64( $lane$ )
35:    $hi \leftarrow lane \gg 32, \quad lo \leftarrow lane \wedge (2^{32} - 1)$ 
36:   return ( $hi, lo$ )
37: end function
38: function BITCOMBINE64( $hi, lo$ )
39:   return ( $hi \ll 32 \oplus lo$ )
40: end function
41: function ENCRYPTION( $Data = (D_0, D_1), \ Key = (k_0, k_1), \ number\_once$ )
42:    $lane_0 \leftarrow D_0, \ lane_1 \leftarrow D_1$ 
43:   GENERATEANDSTOREKEYSTATES( $Key, \ number\_once$ )
44:   for  $i \leftarrow 0$  to  $rounds - 1$  do
45:      $ks \leftarrow KeyState[i], \ rc \leftarrow \text{ROUND\_CONSTANTS}[ks.\text{round\_constant\_index}]$ 
46:      $(w_0, w_1) \leftarrow \text{BITSPLIT64}(lane_0), \ (w_2, w_3) \leftarrow \text{BITSPLIT64}(lane_1)$ 
47:     NEOALZETTE_FORWARDLAYER( $w_0, w_2, rc$ ); NEOALZETTE_FORWARDLAYER( $w_1, w_3, rc$ )  $\triangleright$  Diagonal pairing ( $w_0, w_2$ ) and
 $(w_1, w_3)$ 
48:      $lane_0 \leftarrow \text{BITCOMBINE64}(w_0, w_1), \ lane_1 \leftarrow \text{BITCOMBINE64}(w_2, w_3)$ 
49:     switch ( $ks.\text{choice\_function} \bmod 4$ ) do
50:       case 0:  $lane_j \leftarrow lane_j \oplus ks.\text{subkey}^{(j)}$  for  $j \in \{0, 1\}$ 
51:       case 1:  $lane_j \leftarrow \neg lane_j \oplus ks.\text{subkey}^{(j)}$  for  $j \in \{0, 1\}$ 
52:       case 2:  $lane_j \leftarrow lane_j \lll_{64} ks.\text{bit\_rotation\_amount\_b}$  for  $j \in \{0, 1\}$ 
53:       case 3:  $lane_j \leftarrow lane_j \ggg_{64} ks.\text{bit\_rotation\_amount\_b}$  for  $j \in \{0, 1\}$ 
54:     ( $w_0, w_1) \leftarrow \text{BITSPLIT64}(lane_0), \ (w_2, w_3) \leftarrow \text{BITSPLIT64}(lane_1)$   $\triangleright$  32-bit Binary Pseudo-Hadamard
55:     MIX( $w_0, w_1$ ); MIX( $w_2, w_3$ )
56:      $lane_0 \leftarrow \text{BITCOMBINE64}(w_0, w_1), \ lane_1 \leftarrow \text{BITCOMBINE64}(w_2, w_3)$ 
57:      $lane_0 \leftarrow lane_0 \oplus (1 \ll (ks.\text{bit\_rotation\_amount\_a} \bmod 64))$ 
58:      $lane_1 \leftarrow lane_1 \oplus (1 \ll (63 - (ks.\text{bit\_rotation\_amount\_a} \bmod 64)))$   $\triangleright$  Random Bit Tweak (mirrored)
59:     ADDROUNDKEY( $lane_0, k_0, ks.\text{subkey}^{(0)}$ ); ADDROUNDKEY( $lane_1, k_1, ks.\text{subkey}^{(1)}$ )
60:   end for
61:   return ( $lane_0, lane_1$ )
62: end function
63: function DECRYPTION( $Data = (C_0, C_1), \ Key = (k_0, k_1), \ number\_once$ )
64:    $lane_0 \leftarrow C_0, \ lane_1 \leftarrow C_1$ 
65:   GENERATEANDSTOREKEYSTATES( $Key, \ number\_once$ )
66:   for  $i \leftarrow rounds \downarrow 1$  to 1 do
67:      $ks \leftarrow KeyState[i - 1], \ rc \leftarrow \text{ROUND\_CONSTANTS}[ks.\text{round\_constant\_index}]$ 
68:     SUBTRACTROUNDKEY( $lane_0, k_0, ks.\text{subkey}^{(0)}$ ); SUBTRACTROUNDKEY( $lane_1, k_1, ks.\text{subkey}^{(1)}$ )
69:      $lane_0 \leftarrow lane_0 \oplus (1 \ll (ks.\text{bit\_rotation\_amount\_a} \bmod 64))$   $\triangleright$  Inverse tweak via XOR
70:      $lane_1 \leftarrow lane_1 \oplus (1 \ll (63 - (ks.\text{bit\_rotation\_amount\_a} \bmod 64)))$ 
71:     switch ( $ks.\text{choice\_function} \bmod 4$ ) do
72:       case 0:  $lane_j \leftarrow lane_j \oplus ks.\text{subkey}^{(j)}$ 
73:       case 1:  $lane_j \leftarrow \neg lane_j \oplus ks.\text{subkey}^{(j)}$ 
74:       case 2:  $lane_j \leftarrow lane_j \ggg_{64} ks.\text{bit\_rotation\_amount\_b}$ 
75:       case 3:  $lane_j \leftarrow lane_j \lll_{64} ks.\text{bit\_rotation\_amount\_b}$ 
76:     ( $w_0, w_1) \leftarrow \text{BITSPLIT64}(lane_0), \ (w_2, w_3) \leftarrow \text{BITSPLIT64}(lane_1)$   $\triangleright$  Inverse 32-bit PHT (order per implementation)
77:     UNMIX( $w_2, w_3$ ); UNMIX( $w_0, w_1$ )
78:      $lane_0 \leftarrow \text{BITCOMBINE64}(w_0, w_1), \ lane_1 \leftarrow \text{BITCOMBINE64}(w_2, w_3)$ 
79:     NEOALZETTE_BACKWARDLAYER( $w_1, w_3, rc$ ); NEOALZETTE_BACKWARDLAYER( $w_0, w_2, rc$ )
80:      $lane_0 \leftarrow \text{BITCOMBINE64}(w_0, w_1), \ lane_1 \leftarrow \text{BITCOMBINE64}(w_2, w_3)$ 
81:   end for
82:   return ( $lane_0, lane_1$ )
83: end function
84: function RESETPRNG( $seed$ )
85:   prng.SEED( $seed$ )
86: end function

```

6.5 Binary Pseudo-Hadamard Transform (BPHT)

This section presents the *Binary-PHT* (Binary Pseudo-Hadamard Transform) pair used in the implementation: the forward BPHT₃₂ and inverse BPHT₃₂⁻¹. This transform operates independently on 32-bit word pairs in the linear layer Θ , applied in two parallel lanes: (w_0, w_1) and (w_2, w_3).

Definition 6.1 (Forward BPHT₃₂). *For any $(a, b) \in \{0, 1\}^{32} \times \{0, 1\}^{32}$, define*

$$\text{BPHT}_{32}(a, b) = (a', b') \quad \text{where} \quad \begin{cases} a' = a \oplus b, \\ b' = a \oplus (b \lll_{32} 1). \end{cases}$$

Definition 6.2 (Inverse BPHT₃₂⁻¹). *Given (a', b') , let*

$$w := a' \oplus b' = b \oplus (b \lll_{32} 1).$$

Define PXOR₃₂(\cdot) as the bitwise prefix XOR (inverse Gray code transformation):

$$\text{PXOR}_{32}(w) = w \oplus (w \gg 32 1) \oplus (w \gg 32 2) \oplus (w \gg 32 4) \oplus (w \gg 32 8) \oplus (w \gg 32 16).$$

The inverse transform is then

$$\text{BPHT}_{32}^{-1}(a', b') = (a, b) \quad \text{where} \quad \begin{cases} b = \text{PXOR}_{32}(a' \oplus b'), \\ a = a' \oplus b. \end{cases}$$

Algorithm Implementation (Matching Code)

```

1: function BPHT32_FORWARD( $a, b$ )
2:    $a' \leftarrow a \oplus b$ 
3:    $b' \leftarrow a \oplus (b \ll_{32} 1)$ 
4:   return  $a', b'$ 
5: end function

6: function PXOR32( $w$ ) ▷ Parallel implementation of inverse Gray code
7:    $w \leftarrow w \oplus (w \gg_{32} 1)$ 
8:    $w \leftarrow w \oplus (w \gg_{32} 2)$ 
9:    $w \leftarrow w \oplus (w \gg_{32} 4)$ 
10:   $w \leftarrow w \oplus (w \gg_{32} 8)$ 
11:   $w \leftarrow w \oplus (w \gg_{32} 16)$ 
12:  return  $w$ 
13: end function

14: function BPHT32_INVERSE( $a', b'$ )
15:    $w \leftarrow a' \oplus b'$ 
16:    $b \leftarrow \text{PXOR32}(w)$ 
17:    $a \leftarrow a' \oplus b$ 
18:   return  $a, b$ 
19: end function
```

Usage in the Algorithm (Parallel Pair Application) Represent the 128-bit state as two 64-bit lanes:

$$\text{lane0} = (w_0 \| w_1), \quad \text{lane1} = (w_2 \| w_3), \quad w_i \in \{0, 1\}^{32}.$$

In the forward round, apply:

$$(w_0, w_1) \leftarrow \text{BPHT}_{32}(w_0, w_1), \quad (w_2, w_3) \leftarrow \text{BPHT}_{32}(w_2, w_3),$$

In the inverse round, apply:

$$(w_2, w_3) \leftarrow \text{BPHT}_{32}^{-1}(w_2, w_3), \quad (w_0, w_1) \leftarrow \text{BPHT}_{32}^{-1}(w_0, w_1),$$

followed by lossless repacking into two 64-bit lanes.

Mathematical Principles and Invertibility Treating \oplus and shifts as linear operators over \mathbb{F}_2 , BPHT_{32} acts on (a, b) as:

$$\begin{pmatrix} a' \\ b' \end{pmatrix} = \begin{pmatrix} I & I \\ I & L \end{pmatrix} \begin{pmatrix} a \\ b \end{pmatrix}, \quad \text{where } L(b) = b \ll_{32} 1.$$

The determinant is $I \oplus L$. For any $b \in \{0, 1\}^{32}$, $b \oplus (b \ll_{32} 1) = 0$ implies $b = 0$ (bitwise analysis shows $b_i = b_{i-1}$ with $b_{-1} = 0$, yielding $b \equiv 0$). Thus $\ker(I \oplus L) = \{0\}$, proving matrix invertibility. The explicit inverse is:

$$b = (I \oplus L)^{-1}(a' \oplus b') = \text{PXOR}_{32}(a' \oplus b'), \quad a = a' \oplus b.$$

Hence BPHT_{32} is bijective without requiring lookup tables or branching.

Diffusion Properties (Combined with Rotation/Diagonal S-Box) For a single-bit disturbance $b = 2^j$, $b \oplus (b \ll 1)$ activates bits j and $j+1$ simultaneously. When combined with diagonal NeoAlzette S-Boxes and cross-word rotations, this achieves cross-word and cross-lane diffusion within few rounds. The linear layer provides balanced input to subsequent nonlinear paths (\boxplus_{64} and \lll / \ggg), enhancing differential/linear resistance.

Implementation Note The transform uses only $\oplus, \ll_{32}, \gg_{32}$, ensuring constant-time execution without data-dependent memory access. The inverse employs 5 parallel XOR-folding steps for Gray code inversion, meeting zero-branching and fixed-instruction-path requirements.

7 Performance Evaluation and Security Evaluation

7.1 Performance Evaluation

All benchmarks in this paper were run on a single machine; results therefore characterize only that platform. Community contributions with multi-platform data are welcome.

We evaluate the **XCR/Little_OaldresPuzzle_Cryptic** C++ implementation (two 64-bit lanes, 128-bit block, NeoAlzette-based Π layer, key-dependent Θ , ARX key mixing Γ) alongside **ASCON** and **ChaCha20** under a uniform harness. The appendices (Tables 1) enumerate algorithmic features and summarize timing and throughput observations.

Experimental Setup.

- *Platform.* Intel64 Family 6 Model 151 Stepping 2 @ ≈ 3.61 GHz, RAM 65,277 MB, 64-bit Linux toolchain; compiler optimizations enabled.

- *Data generation.* For each size, an MT19937-64 generator seeded with 1 produces 2^n pseudorandom bits; inputs are aligned to 128 bits for **XCR/Little_OaldresPuzzle_Cryptic**.
- *Cipher parameters.*
 - **ASCON:** fixed key $\{0x0123456789ABCDEF, 0xFEDCBA9876543210\}$ and nonce $\{0, 0\}$; associated data $= 0x12345678$.
 - **ChaCha20:** key $\{0x32, 0x01234567, 0x89ABCDEF, 0xFEDCBA98, 0x76543210\}$; nonce $\{0, 0\}$.
 - **XCR/Little_OaldresPuzzle_Cryptic:** 128-bit blocks with per-round $number_once \in \{0, 1\}^{64}$ in counter mode; no associated data; key $Key = (Key_L, Key_R) \in \{(0, 1)^{64}\}^2$. The internal XCR instance is initialized from `seed` (and re-RESETPRNG(\cdot) between independent trials).
- *Sizes.* Message lengths double from 128 bits up to 671,088,640 bits (10 GB).
- *Metrics.* Wall-clock time and derived throughput; for fairness, encryption and decryption are both measured.

Methodological Notes (Implementation-Consistent).

- Each *encryption* round computes: (i) two diagonal NeoAlzette calls on (w_0, w_2) and (w_1, w_3) with $rc = \text{ROUND_CONSTANTS}[index]$; (ii) a key-dependent Θ branch selected by $cf \in \{0, 1, 2, 3\}$ ($\oplus, \ominus, \lll_{64} \beta, \ggg_{64} \beta$); (iii) two independent 32-bit binary PHTs on word pairs; (iv) a random-bit tweak $x \leftarrow x \oplus 2^{\alpha \bmod 64}$ (mirrored across lanes); (v) Γ with $\boxplus_{64}, \lll_{64} 48, \ggg_{64} 16$.
- *Decryption* exactly inverts: Γ^{-1} , untweak by XOR, inverse branch of Θ , inverse PHT order, then Π^{-1} as two NeoAlzette backward calls (w_1, w_3) then (w_0, w_2) .
- Key material per round derives from the XCR CSPRNG as specified in Section 6.3; $index = (rc_index^{raw} \gg 1) \bmod 16$ selects **ROUND_CONSTANTS**.

Results (Qualitative Summary).

- **ASCON** leads in throughput and scales well with size; **ChaCha20** exhibits a similar growth trend with different absolute performance.
- **XCR/Little_OaldresPuzzle_Cryptic** trails **ASCON** but remains competitive for large inputs; on this platform it is typically within a small constant factor (often observed near $\frac{1}{2}$ of **ASCON**'s speed for selected sizes), while preserving steady scaling behaviour.

Caveats and Outlook.

- Results depend on microarchitecture, compiler, and memory hierarchy. Different CPUs (or SIMD/vectorized implementations) may change relative standings.
- Future work should broaden the benchmark matrix (platforms/variants/round-counts) and report cycles/byte together with statistical confidence.

7.2 Security Evaluation

We summarize design-oriented properties consistent with the C++ implementation and the XCR-driven key schedule.

State and Bijectivity.

- Each round is a composition of invertible mappings: NeoAlzette (Π), branch-selected linear/rotational transforms plus bit-tweak (Θ), and ARX key mixing (Γ). Decryption applies the inverses in reverse order, ensuring per-round and global bijectivity for 128-bit blocks.

Confusion & Diffusion (Round Core).

- Π uses two diagonal NeoAlzette calls, ensuring cross-lane propagation before linear mixing.
- The 32-bit binary PHT ($a' := a \oplus b, b' := a \oplus (b \ll 1)$) injects balanced linear diffusion with a cheap inverse via broadword prefix-xor.
- The random-bit tweak $x \leftarrow x \oplus 2^{\alpha \bmod 64}$ prevents low-weight fixed points and breaks rotational symmetries induced by \lll / \ggg .

Key Schedule (XCR CSPRNG).

- Subkeys are produced by two XCR instances: a key-seeded generator and an instance/system generator salted by $number_once \oplus i$. This yields per-round $sk = (sk^{(0)}, sk^{(1)})$, selector $cf \bmod 4$, and rotation $\beta, \alpha \bmod 64$; $index = (rc.index^{raw} \gg 1) \bmod 16$ selects **ROUND_CONSTANTS**.
- The XCR update combines \oplus, \boxplus_{64} , and rotations $\lll_{64} \{7, 19, 32, 47, 63\}$ with dynamic constants $RC0, RC1, RC2$ derived from **ROUND_CONSTANTS**, *state*, and *counter*, strengthening round asymmetry.

Side-Channel Considerations.

- The design is table-free and branch-light; rounds use word-level ARX primitives and fixed rotations, reducing timing/cache variability on common CPUs. Implementations should still avoid secret-dependent branches, ensure constant-time rotations, and clear sensitive state.

Misuse Resistance and Nonce Handling.

- Reuse of identical $(Key, number_once)$ across messages undermines keystream uniqueness. The wrapper regenerates per-round material from *number_once* and a running *counter*; applications must ensure non-repeating *number_once* per key.

Analytic Posture.

- The round structure counters linear/differential trails through combined modular addition and rotation asymmetry; diagonal S-box wiring and PHT reduce alignment-based biases. Formal bounds (e.g., differential/linear hull estimates, rotational cryptanalysis) and empirical suites (avalanche/NIST STS/SMHasher-like) are recommended for comprehensive validation.

7.3 Security Evaluation with Statistical Tests

In the context of data security assessments, despite our limited expertise in advanced mathematical techniques, we employ a methodology that emulates the properties of uniform randomness. This necessitates the application of effective and robust statistical analysis to substantiate the unpredictability of bit generation within our algorithm. Our methodology involves restricting each encryption operation to a 128-bit data segment, subsequently writing the resultant samples into files, each comprising 128 kilobytes, amounting to a total of 128 sample files. The evaluation process is divided into two distinct phases.

In the **first phase**, the plaintext is uniformly set to all zeros, while the key is used as a unique incrementing counter. This phase represents a rigorous test as it examines the scenario where the plaintext is completely devoid of randomness, and all information is derived from the key and any seed utilized by the **XCR CSPRNG** algorithm. Theoretically, this initialization provides a subtle test intended to observe specific properties within probability distributions.

Conversely, in the **second phase**, the plaintext is derived from a sequence of random data generated by the MT19937-64Bit algorithm, with the key maintaining its role as a unique incrementing counter. The significance of this phase lies in the statistical test outcomes under conditions where the plaintext is random but potentially intermixed with differential malicious data (theoretically analogous to distinguishing between the outputs of true random and algorithmic pseudo-randomization), along with a unique incrementing counter key.

7.4 The Credibility of Statistical Tests and the Rationale for Their Use

The credibility of these tests stems from their mathematical rigor and empirical validation. They are based on well-established statistical principles and have been extensively analyzed and tested across various scenarios. The tests are not only theoretically sound but also practically effective, having been applied in numerous security evaluations and standards, such as those from the National Institute of Standards and Technology (NIST SP 800-22) and the China Randomness Test Specification (GM/T 0005-2021). A detailed description of the Chinese randomness test standards employed is provided in the Appendix.

The rationale for employing these tests in our evaluation process is multifaceted. Firstly, they offer an objective and systematic approach to assessing the randomness of our encryption algorithm's output. By subjecting the generated bits to a battery of tests, we can gain confidence in the algorithm's ability to produce unpredictable sequences, which is essential for thwarting potential attacks that rely on the predictability of the encryption process.

Secondly, the use of multiple tests provides a comprehensive assessment. Each test targets a different aspect of randomness, and together they form a robust evaluation framework. This ensures that any weaknesses in the algorithm's randomness are likely to be detected, as no single test can cover all possible scenarios.

Lastly, the choice of tests and parameters is informed by the specific requirements of our encryption algorithm and the nature of the data being encrypted.

7.5 Integrating P-Values, Q-Values, and Conditional Probability in Cryptographic Analysis

In the realm of cryptographic analysis, the application of statistical hypothesis testing, particularly through the use of P-values and Q-values, alongside the assessment of conditional probabilities, is essential for ensuring the reliability and security of cryptographic algorithms. The P-value, representing the probability of obtaining an observed result under the null hypothesis, serves as a critical measure for determining the statistical significance of findings in the context of encryption.

P-values are instrumental in identifying instances where the behavior of an algorithm deviates significantly from what would be expected under random conditions, thus flagging potential vulnerabilities. However, given the constraints of P-values, including their susceptibility to sample size effects and their inability to directly convey the probability of the hypothesis, Q-values are introduced in multiple hypothesis testing scenarios. Q-values, an adjustment of P-values, play a pivotal role in controlling the False Discovery Rate (FDR), thereby reducing the risk of false positives that can misguide the evaluation process. For details on Q-values, please refer to the paper [J, 2016].

The evaluation of conditional probability, denoted as $\Pr(\text{value}_P | \text{value}_Q)$, is equally vital in cryptographic analysis. It measures the likelihood of observing a specific bit pattern value_P in the encrypted output, given the occurrence of another pattern value_Q . This probabilistic assessment helps in understanding the dependencies or correlations between different bit patterns, enabling a more nuanced analysis of the encryption algorithm's behavior and its potential vulnerabilities.

Together, the integrated application of P-values, Q-values, and conditional probability analysis significantly enhances the robustness and credibility of cryptographic evaluations. This triad of statistical tools allows for a more comprehensive scrutiny of algorithms, revealing not just anomalies but also ensuring that the findings are statistically sound and less prone to false discoveries.

7.5.1 Test Results for Phase 1

In Phase 1 of the **Chacha20** evaluation, where the data is set to all zeros and the key operates in Counter Mode, all statistical tests conducted failed to pass, irrespective of attempts at both 32-bit and 64-bit trials. This underscores Chacha20's suboptimal performance when subjected to non-random data, particularly in the 64-bit scenario.

Our test results indicate that the **XCR/Little_OaldresPuzzle_Cryptic** algorithm not only closely mirrors the performance characteristics of the **ASCON** algorithm, as inferred from the aforementioned conclusions, but also maintains a level of random uniformity indistinguishable from the **ASCON** algorithm. This assertion will be visually substantiated through tables presented in **Figures 3 to 6** included in the appendix.

It is essential to note that due to the extensive nature of our testing, yielding 128 individual data files, we have chosen the final binary file as our reference dataset. Consequently, all data presented in the tables within our screenshots is derived exclusively from the results of this last binary file. The comprehensive statistical test results, including the complete Excel spreadsheet, will be made available in our code repository. [HereIsTheLink](#).

7.5.2 Test Results for Phase 2

Transitioning to Phase 2, where data is generated using the MT19937-64Bit algorithm with Seed 1 and the key is a random 64-bit value (32-bit \times 2), **Chacha20** exhibits significantly improved performance. In this phase, all tests yielded satisfactory or excellent results.

It becomes evident that **Chacha20** encounters challenges in maintaining security when confronted with non-random input data, highlighting the algorithm's dependence on randomness, especially in the 64-bit context.

Chacha20's strengths lie in its capacity to provide ample bit distribution, approaching pseudo-randomness without fully achieving it, even with minimal input data. It excels in scenarios where the input data exhibits sufficient chaos, showcasing the intricacies of the algorithm.

In conclusion, while **Chacha20** demonstrates merits, especially in Phase 2 where randomness is better preserved, the **XCR/Little_OaldresPuzzle_Cryptic** algorithm emerges as the superior choice when **Chacha20** struggles to uphold security and robustness, particularly in non-random data scenarios. The test results are reflected in **Figures 7 to 10**.

7.5.3 Evaluation of Test Results

Upon synthesizing the insights garnered from both Phase 1 and Phase 2 evaluations, a nuanced understanding of the Chacha20 algorithm's efficacy is achieved. Phase 1 revealed its vulnerability to non-random data, resulting in suboptimal outcomes across multiple iterations. Conversely, Phase 2 demonstrated a significant enhancement in performance, particularly when the input data conformed to a more stochastic distribution.

The **XCR/Little_OaldresPuzzle_Cryptic** algorithm consistently exhibited competitive performance, maintaining a level of randomness akin to the ASCON algorithm. This resilience to data patterns underscores its potential in environments where randomness is paramount.

Given the observed performance of **Chacha20**, it becomes apparent that it may not be an ideal candidate for consideration in subsequent evaluations. Consequently, **Chacha20** has been excluded from our pool of reference algorithms for further assessment. This decision narrows our focus to the **ASCON** algorithm and the **XCR/Little_OaldresPuzzle_Cryptic** algorithm as the primary contenders for comparison and evaluation. The upcoming analysis will scrutinize and compare the strengths and weaknesses of **ASCON** and **XCR/Little_OaldresPuzzle_Cryptic** to determine their suitability for our intended application.

To address concerns regarding the perceived issues with **Chacha20**, it is imperative to substantiate our assertions with robust and compelling data. To enhance the persuasiveness of our findings, comprehensive statistical results are meticulously presented in the tables within the appendix, accompanied by graphical representations in **Figures 11 to 14**. This transparent approach aims to provide readers with direct access to the raw data, fostering a clearer understanding of the intricacies and nuances of Chacha20's performance. We believe that this meticulous presentation of data in our supplementary materials will dispel any reservations and contribute to the confidence in the validity of our analysis.

8 Mathematical Proof and Security Deduction of Our Algorithms

In the preceding analysis, the XCR/Little_OaldresPuzzle_Cryptic lightweight cryptographic algorithm has successfully passed the statistical tests of GM/T 0005-2021. This success is attributed to the layered application of ARX primitives, each contributing a degree of non-linearity to the transformation process. As each layer cumulatively integrates data patterns, the resultant distribution tends toward uniformity. This is further enhanced by our design approach in the encryption and decryption processes, where the use of simple linear functions achieves an efficient blend of complexity.

Moreover, the XorConstantRotation CSPRNG, integral to our algorithm, is designed with well-defined diffusion and confusion layers. The initial seeding of this CSPRNG undergoes a whitening process using the Goldreich-Goldwasser-Micali (GGM) construction. The synergistic effect of these complex structures, combined with the statistical indistinguishability of the algorithm's output, supports our assertion that the algorithm adheres to the IND-CPA and IND-CCA semantic security models (chosen plaintext and ciphertext security). Regardless of the attacker's methodology, whether it be exhaustive search or intricate examination of linear and non-linear layer interactions, the emergent complexities present substantial difficulties. This applies to differential analysis, linear analysis, and previously used rotation analysis specific to ARX primitive structures. Consequently, the probability of compromising any individual component of this complex structure is negligible, making the theoretical distinction of our pseudorandom algorithm highly improbable, with a lower bound probability of less than 2^{-64} .

8.1 Formal Security Analysis of XorConstantRotation (XCR)

8.1.1 Primitive Specification (Matches Implementation)

Let the internal state be $s = (x, y, z, c) \in (\mathbb{Z}_{2^{64}})^4$, where the code variable **state** is denoted by z and **counter** by c . Let the public constant table be $\text{RC}[0..299] \subset \{0, 1\}^{64}$ (the 300 64-bit constants in the code).

For a chosen integer input $n \in \mathbb{Z}_{2^{64}}$ (the function argument **number_once**), define

$$\text{RC}_0 := \text{RC}[n \bmod 300], \quad \text{RC}_1 := \text{RC}[(c + n) \bmod 300], \quad \text{RC}_2 := \text{RC}[z \bmod 300].$$

Let **rotl₆₄** and **rotr₆₄** be 64-bit rotations; \boxplus is addition over $\mathbb{Z}_{2^{64}}$.

The single-step state transition $\mathcal{F} : (s, n) \mapsto (s', y')$ implemented by **StateIteration(n)** is:

$$\begin{aligned} & \text{(Diffusion, conditional)} \quad \text{if } x = 0 \text{ then } x \leftarrow \text{RC}_0 \text{ else} \\ & \quad y \leftarrow y \oplus \text{rotl}_{64}(x, 19) \oplus \text{rotl}_{64}(x, 32); \\ & \quad z \leftarrow z \oplus \text{rotl}_{64}(y, 32) \oplus \text{rotl}_{64}(y, 47) \oplus \text{rotl}_{64}(y, 63) \oplus c; \\ & \quad x \leftarrow x \oplus \text{rotl}_{64}(z, 7) \oplus \text{rotl}_{64}(z, 19) \oplus \text{RC}_0 \oplus n; \\ & \text{(Confusion, always)} \quad z \leftarrow z \boxplus (y \oplus \text{rotr}_{64}(y, 1) \oplus \text{RC}_0); \\ & \quad x \leftarrow x \oplus (z \boxplus \text{rotr}_{64}(z, 1) \boxplus \text{RC}_1); \\ & \quad y \leftarrow y \boxplus (x \oplus \text{rotr}_{64}(x, 1) \oplus \text{RC}_2); \\ & \quad c \leftarrow c + 1 \pmod{2^{64}}; \quad y' \leftarrow y, \quad s' \leftarrow (x, y, z, c). \end{aligned}$$

The PRNG oracle answers query n with the returned word y' and updates the internal state to s' .

Initialization (Matches StateInitialize) Given external seed $z_0 \in \{0, 1\}^{64}$, XCR starts from $(x, y, z, c) = (0, 0, z_0, 0)$, runs a 4-epoch bit-building loop that repeatedly calls \mathcal{F} with data-dependent n (as in code), producing a 64-bit value R , and sets $z \leftarrow z \oplus (z_0 \boxplus R)$. Then the PRNG is ready.

8.1.2 Security Model (RoR with Chosen-Input Queries)

In the *Real-or-Random* game, a PPT adversary \mathcal{A} adaptively queries integers n_1, \dots, n_q . In the real world it receives the sequence $\{y_i\}$ from the above oracle; in the random world it receives q independent uniform 64-bit strings. The advantage is

$$\text{Adv}_{\text{XCR}}^{\text{ror}}(t, q) = \left| \Pr[\mathcal{A}^{\mathcal{O}_{\text{real}}}(1^\lambda) = 1] - \Pr[\mathcal{A}^{\mathcal{O}_{\text{rand}}}(1^\lambda) = 1] \right|.$$

8.1.3 Assumptions (Made Explicit)

A1 Random-Constant Hypothesis: For analysis, the fixed table RC is modeled as a vector of independent uniform 64-bit constants (a standard heuristic for ARX primitives with aperiodic constants).

A2 ARX-PRF Heuristic: The stateful ARX core above, keyed by the (hidden) initial seed and executing the conditional diffusion + modular-add confusion, behaves as a pseudorandom function on outputs $\{y_i\}$ under chosen inputs $\{n_i\}$.¹

8.1.4 What Can Be Proven Rigorously vs. What Is Heuristic

We separate (i) statements that follow from basic facts plus A1 (or trivial information theory), and (ii) claims that additionally rely on A2.

(I) Facts that do not need A2.

Lemma 8.1 (Linear structure of the diffusion submap). *Let $L(x) = \text{rotl64}(x, 19) \oplus \text{rotl64}(x, 32)$. Then L is a linear map over \mathbb{F}_2^{64} with:*

- For a unit input mask $\alpha = e_i$, $w_H(L(\alpha)) = 2$ (not 32).
- The (XOR) branch number of L satisfies $B_L \leq 3$ because choosing $w_H(\Delta x) = 1$ yields $w_H(\Delta y) = 2$.

Lemma 8.2 (Where nonlinearity actually comes from). *All nonlinearity in a single iteration arises from the three $\mathbb{Z}_{2^{64}}$ additions in the confusion phase (carry network). The diffusion submap alone is linear.*

Lemma 8.3 (Output collision birthday bound). *For q outputs of $y \in \{0, 1\}^{64}$ (real or random), the probability of at least one collision is $\approx \binom{q}{2}/2^{64}$. Thus any collision-based distinguisher is upper-bounded by $O(q^2/2^{64})$ information-theoretically.*

(II) Claims that require A2 (ARX pseudorandomness).

Theorem 8.4 (RoR security under A1–A2). *Under A1–A2, for any PPT adversary with q adaptive chosen-input queries,*

$$\text{Adv}_{\text{XCR}}^{\text{ror}}(t, q) \leq \epsilon_{\text{ARX}}(t) + O\left(\frac{q^2}{2^{64}}\right),$$

where $\epsilon_{\text{ARX}}(t)$ is the distinguishing advantage against the assumed ARX-PRF core in time t .

Proof sketch. Hybrid H0: real XCR. H1: replace RC by random independent constants (A1). H2: replace the ARX transition by a random stateful PRF with the same interface (A2). The only remaining distinguishing power is via output coincidences (Lemma 8.3). Apply triangle inequality. \square

Theorem 8.5 (Initialization indistinguishability (matches code)). *Let z_0 be the external 64-bit seed. Under A2, the 4-epoch init loop produces a value R that is computationally indistinguishable from uniform and independent of z_0 . Then $z \leftarrow z \oplus (z_0 \boxplus R)$ is indistinguishable from uniform over $\{0, 1\}^{64}$ to any PPT adversary without access to z_0 .*

Proof sketch. Replace the \mathcal{F} -calls in init by a random function (A2), so R is uniform. Then $z \mapsto z \oplus (z_0 \boxplus R)$ is a one-to-one mask by a value indistinguishable from uniform; thus z is computationally indistinguishable from uniform. (We do not claim perfect uniformity since R is only PR.) \square

Resistance to linear/differential heuristics (qualitative, code-aligned).

- The conditional branch $x=0 \Rightarrow x \leftarrow \text{RC}_0$ destroys linear structure at the start and prevents trivial fixed points tied to $x=0$.
- The counter c is XORed into z before additions, breaking rotational symmetries and data-dependent cycles; c increments mod 2^{64} each step.
- Each round does three word additions whose carry propagation raises algebraic degree; after a small number of iterations the algebraic degree of each output bit becomes high (saturating at ≤ 64). We avoid bogus numeric lower bounds and recommend empirical auto-correlation tests if needed.

8.1.5 Collision and Cycle Discussion (State vs. Output)

The full internal state is 256 bits (x, y, z, c) . We do not claim that \mathcal{F} is a permutation on (x, y, z, c) because of the $x=0$ branch, but with A2 and A1 the induced trajectory behaves like a high-entropy Markov chain with negligible short cycles: before c wraps (2^{64} steps) revisiting the same (z, c) pair is heuristically negligible. On outputs, Lemma 8.3 already gives the tight 64-bit birthday bound.

8.1.6 Quantum Query Considerations

For a PRF-like oracle with 64-bit outputs, standard Grover-style lower bounds imply that any quantum distinguisher making q quantum queries satisfies

$$\text{Adv}_{\text{XCR}}^{\text{ror}, q}(q) \in O\left(\frac{q^2}{2^{64}}\right).$$

We do not introduce ad-hoc ‘phase distortion’ factors; the dominant term is the usual $q^2/2^n$ for $n=64$.

¹This is the usual working assumption in ARX designs; we do not claim a standard-model proof.

8.1.7 Design Takeaways (Instantiations)

- **Output size governs birthday term:** classical $q \ll 2^{32}$ (and quantum $q \ll 2^{32}$) keep the birthday term negligible.
- **Constants and counter matter:** using aperiodic RC and injecting c before additions are essential to break rotational/XOR symmetries.
- **Testing guidance:** complement this analysis with empirical linear/differential correlation tests and long-cycle detection up to the c -wrap regime.

Forward Security (Backtracking Resistance). Backtracking resistance asks: if an adversary learns the full internal state $s_t = (x_t, y_t, z_t, c_t)$ at time t , can they recover *earlier* states/outputs? For the implemented step function \mathcal{F} the answer is negative: one step is efficiently invertible given the public inputs (the chosen query n and the counter relation $c_{t-1} = c_t - 1$) up to a small (≤ 300) constant-time ambiguity that collapses by consistency checking.

Concretely, denote the post-*diffusion* values by $(\hat{x}, \hat{y}, \hat{z})$ and the post-*confusion* values by (x', y', z') (these are the stored (x_t, y_t, z_t)). At the beginning of the step the code fixes

$$\text{RC}_0 := \text{RC}[n \bmod 300], \quad \text{RC}_1 := \text{RC}[(c_{t-1} + n) \bmod 300], \quad \text{RC}_2 := \text{RC}[z_{t-1} \bmod 300].$$

Given (x', y', z', c_t) and n , for each guess $r \in \{\text{RC}[0], \dots, \text{RC}[299]\}$ set

$$\begin{aligned} \hat{y} &= y' \boxminus (x' \oplus \text{rot}_{64}(x', 1) \oplus r), \\ \hat{z} &= z' \boxminus (\hat{y} \oplus \text{rot}_{64}(\hat{y}, 1) \oplus \text{RC}_0), \\ \hat{x} &= x' \oplus (\hat{z} \boxplus \text{rot}_{64}(\hat{z}, 1) \boxplus \text{RC}_1). \end{aligned}$$

This inverts the *confusion* phase. Next invert the *diffusion* phase in two cases:

Case A (nonzero branch): compute

$$\begin{aligned} x_{t-1}^{(A)} &= \hat{x} \oplus \text{rotl}_{64}(\hat{z}, 7) \oplus \text{rotl}_{64}(\hat{z}, 19) \oplus \text{RC}_0 \oplus n, \\ y_{t-1}^{(A)} &= \hat{y} \oplus \text{rotl}_{64}(x_{t-1}^{(A)}, 19) \oplus \text{rotl}_{64}(x_{t-1}^{(A)}, 32), \\ z_{t-1}^{(A)} &= \hat{z} \oplus \text{rotl}_{64}(\hat{y}, 32) \oplus \text{rotl}_{64}(\hat{y}, 47) \oplus \text{rotl}_{64}(\hat{y}, 63) \oplus c_{t-1}. \end{aligned}$$

Accept this preimage only if $x_{t-1}^{(A)} \neq 0$.

Case B (zero branch): the code's conditional sets $x \leftarrow \text{RC}_0$ and leaves y, z unchanged in diffusion when $x_{t-1} = 0$. Thus one can also set

$$x_{t-1}^{(B)} := 0, \quad y_{t-1}^{(B)} := \hat{y}, \quad z_{t-1}^{(B)} := \hat{z},$$

and check consistency through the subsequent confusion inversion.

Finally, enforce the table-index consistency

$$r \stackrel{!}{=} \text{RC}[z_{t-1} \bmod 300].$$

For a wrong r this test fails with overwhelming probability; the correct $r = \text{RC}_2$ passes, yielding a unique $(x_{t-1}, y_{t-1}, z_{t-1})$. Therefore one step is invertible with at most 300 trials and constant work per trial. Recursing gives earlier states and outputs.

Conclusion. XCR as implemented does *not* provide forward security/backtracking resistance against state compromise; recovering s_{t-1} from s_t is feasible in $O(300)$ time given n and $c_{t-1} = c_t - 1$.

Backward Security (Prediction Resistance / Break-in Recovery). Prediction resistance asks: after an adversary learns s_t , are *future* outputs still hidden? Since the step function and constant table are public and the transition is deterministic, knowledge of s_t lets an adversary compute $(s_{t+1}, y_{t+1}), (s_{t+2}, y_{t+2}), \dots$ for any chosen inputs n , hence *without reseeding* XCR does not satisfy backward security.

Break-in Recovery with Reseeding. If, at some time $\tau > t$, the generator is *reseeded* with fresh, unpredictable entropy u via the same initialization routine (i.e., calling `Seed(u)` which runs `StateInitialize`), then under the ARX-PRF heuristic (A2) the post-reseed state z becomes computationally indistinguishable from uniform independently of the adversary's prior view. Formally:

Theorem 8.6 (Break-in Recovery after Reseed). *Assume the reseed injects at least $k \geq 64$ bits of min-entropy unknown to the adversary into the initialization procedure. Under A2, for any PPT adversary that learned s_t (and all prior inputs), the outputs produced after the reseed are computationally indistinguishable from uniform 64-bit strings up to the RoR birthday term $O(q^2/2^{64})$.*

Sketch. By Theorem 8.5 (initialization indistinguishability, adapted to the reseed point), the internal z after reseed is computationally indistinguishable from uniform to the adversary. Subsequent steps are then indistinguishable from those of a fresh instance, so the usual RoR analysis applies; the only information-theoretic term is the output-collision birthday bound $O(q^2/2^{64})$. \square

Design note. If forward/backward security against state compromise is a goal *without relying on reseed*, one may harden the transition by injecting a one-way compression of hidden material into the state each step, e.g., replace the conditional diffusion prelude by

$$(x, y, z) \leftarrow (x, y, z) \oplus \text{KDF}(z, \text{RC}_0, \text{RC}_1, \text{RC}_2, n, c),$$

with KDF a standard hash-based KDF, so that inverting one step requires inverting the KDF. This changes the primitive and lies outside the current implementation.

Supplementary Material (Non-core but Potentially Useful Math Notes)

Scope and stance. This section collects auxiliary observations and modelling choices that are *not* required for the core RoR security claim, but can help reviewers understand why the implementation choices (constants, counter injection, conditional branch) disrupt simple algebraic structure. None of the results below strengthens the main bound; when a claim depends on standard ARX heuristics, we mark it explicitly.

On the constant table modelling (A1). The code hard-codes 300 distinct 64-bit words $\text{RC}[0..299]$ built from well-known mathematical constants and a deterministic nonlinear recipe. For analysis we use the common heuristic:

A1 (Random-Constant Hypothesis). Treat $(\text{RC}[i])$ as aperiodic, pairwise independent 64-bit constants unknown to the adversary except for their fixed values; i.e., they behave as if sampled uniformly once and for all.

This is only a modelling convenience to argue symmetry-breaking; we never replace *outputs* by randomness solely due to A1.

RX/differential-linear sanity checks. Write the diffusion submap $L(x) = \text{rotl}_{64}(x, 19) \oplus \text{rotl}_{64}(x, 32)$. Over \mathbb{F}_2^{64} , L is linear and for a unit mask e_i we have $w_H(L(e_i)) = 2$ (two 1-bits at positions $i+19$ and $i+32$). Hence the XOR branch number of *diffusion alone* is at most 3. All nonlinearity in a step comes from the three additions in the confusion phase:

$$z \leftarrow z \boxplus (\dots), \quad x \leftarrow x \oplus (z \boxplus \text{rotr}(z, 1) \boxplus \text{RC}_1), \quad y \leftarrow y \boxplus (\dots).$$

Because carries depend on multiple input bits, a single active bit after the diffusion typically triggers long carry chains; precise closed-form differential probabilities for word additions depend on the specific input/output differences. We therefore refrain from giving numeric per-round bounds here; instead, for r steps it is reasonable under ARX heuristics (A2) to assume rapid decay of exploitable linear/differential correlations once the additions are involved for a few iterations.

Fixed points, 1-cycles, and rotational symmetries. Consider a one-step fixed point under a chosen input n :

$$(x', y', z', c') = (x, y, z, c) \quad \text{with} \quad c' = c + 1 \pmod{2^{64}}.$$

This is impossible because c always changes. For a 2-cycle we require $(x^{(2)}, y^{(2)}, z^{(2)}, c^{(2)}) = (x^{(0)}, y^{(0)}, z^{(0)}, c^{(0)})$ which forces $c^{(2)} = c^{(0)}$ and thus $c^{(1)} = c^{(0)} - 1$. Plugging back yields a system of three coupled ARX equations in x, y, z involving $\text{RC}_0, \text{RC}_1, \text{RC}_2$ at two successive counters and indices $(n, (c+n) \bmod 300, (z \bmod 300))$. Under A1 the chance that these three word-equations hold simultaneously for an *a priori* fixed (n, c) is negligible ($\approx 2^{-64}$ per independent constraint). This does not prove lower bounds on cycle length, but it explains why trivial short cycles are unlikely.

Rotational symmetries (mapping (x, y, z) to $(\text{rotl}(x, k), \text{rotl}(y, k), \text{rotl}(z, k))$) are broken in two places: (i) constants are *not* rotated together with the state and are indexed by $z \bmod 300$; (ii) the counter c injects a time-varying non-rotational mask before additions. Hence no nontrivial global rotation commutes with one step unless $k \equiv 0 \pmod{64}$.

State abstraction and SMT-friendliness. The concrete step function \mathcal{F} on $(x, y, z, c) \in (\mathbb{Z}_{2^{64}})^4$ is not a permutation due to the $x=0 \Rightarrow x \leftarrow \text{RC}_0$ branch, and it is certainly not idempotent. For bounded, machine-checked properties we instead recommend an *abstract* monotone domain $\mathcal{A} = \{\{0, 1, ?\}^{64}\}^3 \times \{0, 1, ?\}^{64}$ capturing per-bit *may/must* information. The best-correct abstraction $\alpha \circ \mathcal{F} \circ \gamma$ is monotone w.r.t. the product order on \mathcal{A} because all concrete operations (XOR, rotations, addition) are monotone in the standard bit-wise information order. This supports SMT queries such as:

- *No all-zero absorbing state:* $\forall n, \mathcal{F}((0, 0, 0, c), n) \neq (0, 0, 0, c+1)$ (UNSAT under A1).
- *No 1-cycle:* $\forall n, \mathcal{F}(s, n) \neq s$ (UNSAT because c increments).
- *Guard effectiveness:* $x=0 \Rightarrow x' = \text{RC}[n \bmod 300] \neq 0$ (VALID by code).

These are efficiently checkable as bit-vector formulas and align with the implementation.

Multi-instance composition (same code, independent seeds). Let k instances run in parallel with independent external seeds and the same constant table. For a distinguisher that adaptively chooses query sequences for each instance, a union bound gives

$$\text{Adv}_{\text{XCR}^{\otimes k}}^{\text{ror}}(t, q_1, \dots, q_k) \leq \sum_{j=1}^k \text{Adv}_{\text{XCR}}^{\text{ror}}(t_j, q_j) + O\left(\frac{(\sum_j q_j)^2}{2^{64}}\right),$$

where the $O(\cdot)$ term is the global birthday term on 64-bit outputs. This matches the intuition that outputs—not internal states—govern the composition penalty. (If seeds are correlated or identical, the bound may worsen; independence is recommended.)

“AES-generated constants” variant (not used here). If one instantiates the constant table as $\text{RC}[i] = \text{AES}_K(i)$ for a fixed, secret K , then any nontrivial distinguishing signal that stems *only* from constants (e.g., detecting periodicity or algebraic biases across i) reduces to distinguishing AES_K from a PRP/PRF on the domain of indices. Formally, a simulator can embed AES_K calls when answering RC lookups; a successful constants-based distinguisher becomes an AES distinguisher. *Caveat:* this variant is a design alternative; the shipped implementation uses a fixed, public table, so the reduction is not applicable to the current code and is mentioned only as a drop-in option.

Internal-state collisions and random-mapping intuition (non-claims). A single instance evolves deterministically and must eventually enter a cycle in the 2^{256} -state space. We do not claim that \mathcal{F} is a random mapping, but the presence of (i) counter increment, (ii) data-dependent constant indices, and (iii) multiple additions suggests that short cycles should be rare under A1/A2. As a rule of thumb, empirical cycle-length exploration up to the c -wrap regime (2^{64} steps) is informative; no formal lower bound is asserted here.

Suggested empirical checks (to complement proofs).

- Walsh spectrum / correlation matrices for a few steps to confirm quick decay of linear bias once additions appear.
- RX-differential trails search with bounded weight to ensure no exceptionally high-probability trails exist for 2–3 steps.
- Long-run statistics for repeated outputs and near-repeats over q up to 2^{32} to match the $q^2/2^{64}$ birthday curve.

Takeaway. These notes justify design choices that deliberately obstruct linear structure and symmetry while keeping the core proof focused on RoR indistinguishability with a clean birthday penalty on 64-bit outputs and an explicit ARX heuristic term.

8.2 ARX Probabilistic Analysis [Ya, 2017]

8.2.1 Modular Addition Differential Analysis

Let $x, y, z \in \mathbb{F}_2^n$ with $z = x \boxplus y$. For differences $\alpha, \beta, \gamma \in \mathbb{F}_2^n$

Definition 8.1 (Carry Constraint Function). *The differential validity condition is determined by:*

$$\Psi(\alpha, \beta, \gamma) := (\neg\alpha \oplus \beta) \wedge (\neg\alpha \oplus \gamma)$$

Where α, β, γ is the difference between x, y and z respectively. To get the x, y difference, you just need to do an \oplus operation on the original value and a small change in the value to get it. $\alpha = \Delta x = x \oplus x' \dots$

Theorem 8.7 (Differential Probability). *When the carry constraint holds:*

$$\Pr[\alpha, \beta \xrightarrow{\boxplus} \gamma] = \begin{cases} 2^{-\text{HW}(\Psi(\alpha, \beta, \gamma) \wedge \text{mask}(n-1))} & \text{if } \Psi(\alpha \lll 1, \beta \lll 1, \gamma \lll 1) \wedge (\alpha \oplus \beta \oplus \gamma \oplus (\beta \lll 1)) = 0 \\ 0 & \text{otherwise} \end{cases}$$

where $\text{mask}(k) := 2^k - 1$, left shift operator \lll .

8.2.2 Rotation Analysis

For $y = x \lll r$ with rotation r :

Lemma 8.8 (Rotation Propagation).

$$\Pr[\alpha \lll r \beta] = \begin{cases} 1 & \beta = \alpha \lll r \\ 0 & \text{otherwise} \end{cases}$$

8.2.3 XOR Operation Analysis

For $z = x \oplus y$:

Lemma 8.9 (XOR Differential Propagation).

$$\Pr[(\alpha, \beta) \xrightarrow{\oplus} \gamma] = \begin{cases} 1 & \gamma = \alpha \oplus \beta \\ 0 & \text{otherwise} \end{cases}$$

8.2.4 Linear Analysis Framework

Let $x, y, z \in \mathbb{F}_2^n$ with $z = x \boxplus y$. For linear correlation mask $\mu, \nu, \omega \in \mathbb{F}_2^n$

Definition 8.2 (Linear Correlation Coefficient).

$$C(\mu, \nu, \omega) = 1_{\{\mu \oplus \omega \prec z\}} 1_{\{\nu \oplus \omega \prec z\}} (-1)^{(\mu \oplus \omega)(\nu \oplus \omega)} 2^{-\text{HW}(z)}$$

where 1_{G_f} is the indicator function:

$$1_{G_f} := \{(x, f(x)) | x \in F_2^n\}$$

where $z = M_n^T(\mu \oplus \nu \oplus \omega)$ and M_n is the carry transition matrix.

Carry Transition Matrix Construction:

The carry transition matrix M_n is constructed to model the propagation of carry bits in modular addition. For an n -bit word $x = (x_{n-1}, x_{n-2}, \dots, x_0)$, the matrix M_n transforms x as follows:

$$M_n(x) = (x_{n-2} \oplus x_{n-3} \oplus \dots \oplus x_0, x_{n-3} \oplus x_{n-4} \oplus \dots \oplus x_0, \dots, x_1 \oplus x_0, x_0, 0)$$

Example: 32-bit Carry Transition Matrix

For $n = 32$, the carry transition matrix M_{32} is given by:

$$M_{32} = \begin{pmatrix} 0 & 0 & 0 & 0 & \dots & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 & \dots & 0 & 0 & 0 \\ 1 & 1 & 0 & 0 & \dots & 0 & 0 & 0 \\ 1 & 1 & 1 & 0 & \dots & 0 & 0 & 0 \\ \vdots & \vdots & \vdots & \vdots & \ddots & \vdots & \vdots & \vdots \\ 1 & 1 & 1 & 1 & \dots & 0 & 0 & 0 \\ 1 & 1 & 1 & 1 & \dots & 1 & 0 & 0 \\ 1 & 1 & 1 & 1 & \dots & 1 & 1 & 0 \end{pmatrix}$$

Each row i has ones starting from column $i + 1$, indicating that the carry effect propagates from lower bits to higher bits.

Transposed Carry Transition Matrix

Since the matrix is defined column-wise, we need its transposed form M_{32}^T to compute z :

$$M_{32}^T = \begin{pmatrix} 0 & 1 & 1 & 1 & \cdots & 1 & 1 & 1 \\ 0 & 0 & 1 & 1 & \cdots & 1 & 1 & 1 \\ 0 & 0 & 0 & 1 & \cdots & 1 & 1 & 1 \\ 0 & 0 & 0 & 0 & \cdots & 1 & 1 & 1 \\ \vdots & \vdots & \vdots & \vdots & \ddots & \vdots & \vdots & \vdots \\ 0 & 0 & 0 & 0 & \cdots & 0 & 1 & 1 \\ 0 & 0 & 0 & 0 & \cdots & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & \cdots & 0 & 0 & 0 \end{pmatrix}$$

Each column j has ones starting from row $j + 1$, ensuring that the carry propagation is properly modeled in modular addition.

Computation of z

The final computation of z is given by:

$$z = M_{32}^T(\mu \oplus \nu \oplus \omega)$$

where μ, ν, ω are the linear masks for the inputs and output, and \oplus represents the bitwise XOR operation.

Theorem 8.10 (ARX Linear Composition). *The total correlation preserves:*

$$C_{ARX} = \prod_{i=1}^k C(\mu_i, \nu_i, \omega_i) \cdot \delta(\text{rotation constraints})$$

with $\delta(\cdot)$ enforcing rotation linearity.

8.2.5 Formal Security Boundaries

Theorem 8.11 (ARX Security Bound). *For any $\epsilon > 0$, after R rounds:*

$$\max_{\Delta, \Gamma} \left(\Pr[\Delta \xrightarrow{R} 0], |C(\Gamma)| \right) \leq 2^{-n(1-\epsilon)}$$

when rotation constants satisfy $\gcd(r_i, n) = 1$.

Corollary 8.12. *For Little_OaldresPuzzle_Cryptic with $n = 128$ and optimized rotations, 12 rounds suffice for 128-bit security against differential/linear attacks.*

8.3 Security Boundary Derivation

Theorem 8.13 (ARX Min-Max Security Bound). *For an ARX cipher with R rounds and word size n , under the differential-linear hull hypothesis:*

$$\min \left(\max_{\Delta \neq 0} DP^R(\Delta), \max_{\Gamma \neq 0} LP^R(\Gamma) \right) \leq 2^{-n+\lceil \log_2 R \rceil}$$

where DP^R and LP^R denote R -round differential/linear probabilities.

Proof. We combine the component-wise bounds through three technical steps:

Step 1: Component Probability Aggregation From Definition (Carry Constraint) and Theorem (Differential Probability):

$$\max_{\alpha, \beta, \gamma} DP_{\text{add}} \leq 2^{-(n-1)} \quad (\text{worst-case addition})$$

Rotation/XOR operations have $DP \in \{0, 1\}$ For linear layer:

$$\max_{\mu, \nu, \omega} |C(\mu, \nu, \omega)| \leq 2^{-\lfloor n/2 \rfloor}$$

Step 2: Round Composition Using the sub-multiplicativity of differential/linear probabilities:

$$DP^R \leq \left(\max_{\text{add}} DP \right)^R \cdot \prod_{i=1}^{R-1} DP_{\text{mix}}$$

where mixing layers contribute $DP_{\text{mix}} \leq 1 + 2^{-n/2}$ (from partition technique). Similarly for LP .

Step 3: Minimax Optimization Let $p_{\max} = \max(DP^R, LP^R)$. For optimal rotations $(r_1, r_2) = (n/3, 2n/3)$:

$$p_{\max} \leq \left(2^{-(n-1)} \cdot (1 + 2^{-n/2}) \right)^R$$

Taking logarithm:

$$-\log_2 p_{\max} \geq R(n-1) - R \cdot 2^{-n/2} \geq n - \log_2 R$$

when $R \leq 2^{n/2}$. Rearrangement completes the proof. \square

Additive Differential Refinement [Niu et al., 2023] The boundary tightness is confirmed by:

Corollary 8.14. For the structure $\mathcal{F}(x, y) = [(x \boxplus_n y) \lll r_1] \oplus [y \lll r_2]$:

$$\max_{\Delta} \text{DP}(\Delta) \leq 2^{-n/2} \Rightarrow \text{DP}^R \leq 2^{-Rn/2 + \log_2 R}$$

Achieving 2^{-n} security requires:

$$R \geq \left\lceil \frac{n + \log_2 R}{n/2} \right\rceil \Rightarrow R \geq 3 \quad \forall n \geq 64$$

- **Partition Proof:** For affine subspaces $V_i \subset \mathbb{F}_2^n$ where $\dim V_i = n/2$, the differential becomes deterministic within each subspace. Total $2^{n/2}$ subspaces.
- **Matrix Complexity:** Each 8×8 state matrix M_i satisfies $\|M_i\|_2 \leq \sqrt{2}$, giving total bound:

$$\prod_{i=1}^{4n} \|M_i\|_2 \leq 2^{2n} \leq 2^{3n/2} \quad (\text{for } n \geq 64)$$

- **Optimality Condition:** Rotation constants $(n/3, 2n/3)$ minimize the matrix norm product:

$$\arg \min_{r_1, r_2} \prod_{i=1}^4 \|M_i(r_1, r_2)\|_2 = ([n/3], [2n/3])$$

The following analysis pertains to the encryption and decryption functions used in the Little OaldresPuzzle-Cryptic, and as established in our previous analysis, we have fully proven the security of the XCR CSPRNG. The reason we need to prove the security of the XCR (XOR Constant Rotation) algorithm is that it serves as the core algorithm utilized in the key schedule process.

8.4 NeoAlzette ARX Structure Deep Security Analysis

In this section, we provide a high-level security discussion of the NeoAlzette ARX structure as adopted in our design, without diving into full formal derivations or extensive empirical datasets. The focus is on outlining the structural properties and our experimental methodology rather than presenting exhaustive cryptanalytic results.

Our evaluation was conducted as a *single-point experimental test* targeting specific differential and linear characteristics under constrained conditions. This means we did not perform a full-scale exhaustive cryptanalysis, nor did we aim to surpass the capabilities of established cryptanalytic groups or award-winning researchers in the field. The intent was to validate that, even under such limited-scope tests, the structure demonstrates promising resistance against common ARX-related vulnerabilities.

It is important to emphasize that our available computational resources are orders of magnitude smaller than those typically accessible to large-scale cryptanalytic teams. As such, the presented results should be interpreted as indicative, rather than definitive, of the cipher's resistance against the strongest known attacks.

While we omit the complete derivations and statistical distributions here, the qualitative takeaway from our tests is that the chosen rotation constants, modular additions, and XOR layers collectively provide strong diffusion and nonlinearity properties within the targeted experimental scope. A full-scale analysis remains an open task for future work, ideally conducted with substantially greater computational capacity.

8.4.1 Step-by-Step Differential Analysis of the NeoAlzette Forward Layer

[Lipmaa and Moriai, 2001] Let the input difference to the forward layer be $(\Delta a^{(0)}, \Delta b^{(0)})$. The forward layer consists of the following 14 steps. We now “compute” the differential propagation probability at each step by manually substituting the data from our experiments.

Case 1: $\Delta a^{(0)} = 0$ and $\Delta b^{(0)} = 1$.

1. **Step 1:** $b \leftarrow b \oplus a$. The difference propagates as

$$\Delta b^{(1)} = \Delta b^{(0)} \oplus \Delta a^{(0)},$$

with probability

$$P_1 = 1.$$

2. **Step 2:** $a \leftarrow (a \boxplus b) \ggg_{32} 31$. Let $S_2 = a \boxplus_{32} b$ and $\Delta S_2 = \Delta a^{(0)} \boxplus_{64} \Delta b^{(1)}$. Then, by the theorem above,

$$P_2 = \Pr[\Delta a^{(0)}, \Delta b^{(1)} \xrightarrow{\text{}} \Delta S_2] = 9.31 \times 10^{-10}.$$

After rotation,

$$\Delta a^{(2)} = (S_2 \ggg_{32} 31),$$

and rotation contributes no loss, so P_2 remains.

3. **Step 3:** $a \leftarrow a \oplus rc$. Hence,

$$\Delta a^{(3)} = \Delta a^{(2)} \oplus rc,$$

with

$$P_3 = 1.$$

4. **Step 4:** $b \leftarrow b \boxplus_{32} a$. Let $S_4 = b \boxplus_{32} a$ with $\Delta S_4 = \Delta b^{(1)} \boxplus_{32} \Delta a^{(3)}$. Then,

$$P_4 = 3.81 \times 10^{-6}.$$

Set

$$\Delta b^{(4)} = S_4.$$

5. **Step 5:** $a \leftarrow (a \oplus b) \lll_{32} 24$. With

$$\Delta a^{(5)} = (\Delta a^{(3)} \oplus \Delta b^{(4)}) \lll_{32} 24,$$

we have

$$P_5 = 1.$$

6. **Step 6:** $a \leftarrow a \boxplus_{32} rc$. Let $S_6 = a \boxplus_{32} rc$ with $\Delta S_6 = \Delta a^{(5)} \boxplus_{32} rc$. Then,

$$P_6 = 1.53 \times 10^{-5}.$$

Set

$$\Delta a^{(6)} = S_6.$$

7. **Step 7:** $b \leftarrow (b \lll_{32} 8) \oplus rc$. Then,

$$\Delta b^{(7)} = (\Delta b^{(4)} \lll_{32} 8) \oplus rc,$$

and

$$P_7 = 1.$$

8. **Step 8:** $a \leftarrow a \boxplus_{32} b$. With $S_8 = a \boxplus_{32} b$ and $\Delta S_8 = \Delta a^{(6)} \boxplus_{32} \Delta b^{(7)}$, we obtain

$$P_8 = 3.91 \times 10^{-3}.$$

Set

$$\Delta a^{(8)} = S_8.$$

9. **Step 9:** $a \leftarrow a \oplus b$. Then,

$$\Delta a^{(9)} = \Delta a^{(8)} \oplus \Delta b^{(7)},$$

with

$$P_9 = 1.$$

10. **Step 10:** $b \leftarrow (a \boxplus_{32} b) \ggg_{32} 17$. Let $S_{10} = a \boxplus_{32} b$ with $\Delta S_{10} = \Delta a^{(9)} \boxplus_{32} \Delta b^{(7)}$. Then,

$$P_{10} = 9.77 \times 10^{-4}.$$

After rotation,

$$\Delta b^{(10)} = (S_{10} \ggg_{32} 17),$$

with no extra loss.

11. **Step 11:** $b \leftarrow b \oplus rc$. So,

$$\Delta b^{(11)} = \Delta b^{(10)} \oplus rc,$$

and

$$P_{11} = 1.$$

12. **Step 12:** $a \leftarrow a \boxplus_{32} b$. With $S_{12} = a \boxplus_{32} b$ and $\Delta S_{12} = \Delta a^{(9)} \boxplus_{32} \Delta b^{(11)}$, we have

$$P_{12} = 3.91 \times 10^{-3}.$$

Set

$$\Delta a^{(12)} = S_{12}.$$

13. **Step 13:** $b \leftarrow (a \oplus b) \lll_{32} 16$. That is,

$$\Delta b^{(13)} = (\Delta a^{(12)} \oplus \Delta b^{(11)}) \lll_{32} 16,$$

with

$$P_{13} = 1.$$

14. **Step 14:** $b \leftarrow b \boxplus_{32} rc$. Let $S_{14} = b \boxplus_{32} rc$ with $\Delta S_{14} = \Delta b^{(13)} \boxplus_{32} rc$. Then,

$$P_{14} = 1.56 \times 10^{-2}.$$

Set

$$\Delta b^{(14)} = S_{14}.$$

Thus, the overall differential probability for the forward layer in this case is given by:

$$P_{\text{total}} = \prod_{i \in \{2, 4, 6, 8, 10, 12, 14\}} P_i = 9.31 \times 10^{-10} \cdot 3.81 \times 10^{-6} \cdot 1.53 \times 10^{-5} \cdot 3.91 \times 10^{-3} \cdot 9.77 \times 10^{-4} \cdot 3.91 \times 10^{-3} \cdot 1.56 \times 10^{-2} \approx 1.26 \times 10^{-29}.$$

Case 2: $\Delta a^{(0)} = 2$ and $\Delta b^{(0)} = 2$.

Following the same step-by-step analysis but substituting the computed probabilities for this case, we have:

1. $\Delta b^{(1)} = \Delta b^{(0)} \oplus \Delta a^{(0)}$, $P_1 = 1$.
2. $P_2 = 9.31 \times 10^{-10}$ (modular addition and rotr).
3. $P_3 = 1$ (XOR with rc).
4. $P_4 = 7.63 \times 10^{-6}$ (modular addition $b \boxplus a$).
5. $P_5 = 1$ (XOR and rotl).
6. $P_6 = 3.81 \times 10^{-6}$ (modular addition $a \boxplus rc$).
7. $P_7 = 1$ (rotl and XOR with rc).
8. $P_8 = 3.91 \times 10^{-3}$ (modular addition $a \boxplus b$).
9. $P_9 = 1$ (XOR).
10. $P_{10} = 3.91 \times 10^{-3}$ (modular addition and rotr).
11. $P_{11} = 1$ (XOR with rc).
12. $P_{12} = 1.95 \times 10^{-3}$ (modular addition $a \boxplus b$).
13. $P_{13} = 1$ (XOR and rotl).
14. $P_{14} = 9.77 \times 10^{-4}$ (modular addition $b \boxplus rc$).

Therefore, the overall differential probability is:

$$P_{\text{total}} = 9.31 \times 10^{-10} \cdot 7.63 \times 10^{-6} \cdot 3.81 \times 10^{-6} \cdot 3.91 \times 10^{-3} \cdot 3.91 \times 10^{-3} \cdot 1.95 \times 10^{-3} \cdot 9.77 \times 10^{-4} \approx 7.89 \times 10^{-31}.$$

As shown above, each ARX operation is analyzed according to its differential propagation probability. The XOR and rotation operations contribute a factor of 1 by Lemmas 1 and 2. The modular additions are the only operations that incur probability losses due to the internal carry propagation, as expressed by Theorem 1. The above “hand-calculation” – though in reality performed via computer simulation – demonstrates that even small differences in the input (e.g., $\Delta a = 0, \Delta b = 1$ vs. $\Delta a = 2, \Delta b = 2$) can lead to dramatic differences in the overall differential probability (here, approximately 1.26×10^{-29} and 7.89×10^{-31} , respectively). This analysis underpins our security claims and shows how the intricate ARX structure ensures a high degree of differential diffusion.

8.4.2 Step-by-Step Linear Correlation Analysis of the NeoAlzette Forward Layer

Let $\alpha = (\alpha_a, \alpha_b) \in (\mathbb{F}_2^{32})^2$ be the input mask and $\beta = (\beta_a, \beta_b)$ the output mask. Our empirical analysis reveals the following linear correlation coefficients:

1. **Step 1:** $b \leftarrow b \oplus a$
Perfect correlation preservation: $c_1 = 1$
2. **Step 2:** $a \leftarrow (a \boxplus_{32} b) \ggg_{32} 31$
Nonlinear distortion: $c_2 = -4.66 \times 10^{-10}$
3. **Step 3:** $a \leftarrow a \oplus rc$
Trivial propagation: $c_3 = 1$
4. **Step 4:** $b \leftarrow b \boxplus_{32} a$
Strong decorrelation: $c_4 = 7.07 \times 10^{-74}$
5. **Step 5:** $a \leftarrow (a \oplus b) \lll_{32} 24$
Linear preservation: $c_5 = 1$
6. **Step 6:** $a \leftarrow a \boxplus_{32} rc$
Deep nonlinearity: $c_6 = 7.60 \times 10^{-65}$
7. **Step 7:** $b \leftarrow (b \lll_{32} 8) \oplus rc$
Deterministic: $c_7 = 1$
8. **Step 8:** $a \leftarrow a \boxplus_{32} b$
Null correlation: $c_8 = 0$
9. **Step 9:** $a \leftarrow a \oplus b$
Ideal linear: $c_9 = 1$

10. **Step 10:** $b \leftarrow (a \boxplus_{32} b) \ggg_{32} 17$
Zero correlation: $c_{10} = 0$

11. **Step 11:** $b \leftarrow b \oplus rc$
Trivial: $c_{11} = 1$

12. **Step 12:** $a \leftarrow a \boxplus_{32} b$
Null: $c_{12} = 0$

13. **Step 13:** $b \leftarrow (a \oplus b) \lll_{32} 16$
Linear: $c_{13} = 1$

14. **Step 14:** $b \leftarrow b \boxplus_{32} rc$
Final annihilation: $c_{14} = 0$

Total Correlation: The rounds default is 4.

$$c_{\text{total}} = \prod_{i=1}^{\text{rounds}} c_i = (-4.66 \times 10^{-10}) \times (7.07 \times 10^{-74}) \times \dots \times 0 = -0.00$$

How to read the numbers. The per-step probabilities P_i listed in the previous subsection are *empirical frequencies* obtained from 2^{28} pseudorandom differential pairs; they illustrate *average-case* behaviour of one specific differential trail under the (independence-of-carries) Markov assumption. In contrast, the table below reports *worst-case single-use statistics* for the NEOALZETTE S-box itself: for every one-bit input difference Δ and every round constant RC_i we let an SMT optimiser search the *minimum Hamming weight* $w_{\min}^{(i)}$ of the resulting output difference. The quantity $p_{\max}^{(i)} = 2^{-w_{\min}^{(i)}}$ therefore upper-bounds the best probability any attacker can hope to achieve with one active S-box call, independent of the *average* behaviour shown above.

8.4.3 NeoAlzette ARX-SBOX: Formal Statement, Proof of Correctness, and Security Positioning

Forward map (code-faithful). For $(a, b) \in (\{0, 1\}^{32})^2$ and a 32-bit constant rc , the forward layer $F_{rc} : (\{0, 1\}^{32})^2 \rightarrow (\{0, 1\}^{32})^2$ is the following sequence:

- (1) $b \leftarrow b \oplus a;$ (2) $a \leftarrow (a \boxplus_{32} b) \ggg 31;$ (3) $a \leftarrow a \oplus rc;$
- (4) $b \leftarrow b \boxplus_{32} a;$ (5) $a \leftarrow (a \oplus b) \lll 24;$ (6) $a \leftarrow a \boxplus_{32} rc;$
- (7) $b \leftarrow (b \lll 8) \oplus rc;$ (8) $a \leftarrow a \boxplus_{32} b;$
- (9) $a \leftarrow a \oplus b;$ (10) $b \leftarrow (a \boxplus_{32} b) \ggg 17;$ (11) $b \leftarrow b \oplus rc;$
- (12) $a \leftarrow a \boxplus_{32} b;$ (13) $b \leftarrow (a \oplus b) \lll 16;$ (14) $b \leftarrow b \boxplus_{32} rc.$

Primitive bijections. Each primitive used above is a bijection on 32-bit words, with the following two-sided inverses:

$$\begin{aligned} x \mapsto x \oplus c &\leftrightarrow x \mapsto x \oplus c, \\ x \mapsto x \lll r &\leftrightarrow x \mapsto x \ggg r, \\ x \mapsto x \boxplus_{32} c &\leftrightarrow x \mapsto x \boxminus_{32} c, \\ (b \leftarrow b \oplus a) &\leftrightarrow (b \leftarrow b \oplus a), \\ (a \leftarrow a \boxplus_{32} b) &\leftrightarrow (a \leftarrow a \boxminus_{32} b), \\ (b \leftarrow b \boxplus_{32} a) &\leftrightarrow (b \leftarrow b \boxminus_{32} a). \end{aligned}$$

Any finite composition of these is therefore a permutation.

Explicit inverse (backward map). Reversing the 14 updates and inverting each primitive yields F_{rc}^{-1} :

- (14⁻¹) $b \leftarrow b \boxplus_{32} rc;$ (13⁻¹) $b \leftarrow (b \ggg 16) \oplus a;$
- (12⁻¹) $a \leftarrow a \boxplus_{32} b;$ (11⁻¹) $b \leftarrow b \oplus rc;$
- (10⁻¹) $b \leftarrow (b \lll 17) \boxplus_{32} a;$ (9⁻¹) $a \leftarrow a \oplus b;$
- (8⁻¹) $a \leftarrow a \boxplus_{32} b;$ (7⁻¹) $b \leftarrow (b \boxplus_{32} rc) \ggg 8;$
- (6⁻¹) $a \leftarrow a \boxplus_{32} rc;$ (5⁻¹) $a \leftarrow (a \ggg 24) \oplus b;$
- (4⁻¹) $b \leftarrow b \boxplus_{32} a;$ (3⁻¹) $a \leftarrow a \oplus rc;$
- (2⁻¹) $a \leftarrow (a \lll 31) \boxplus_{32} b;$ (1⁻¹) $b \leftarrow b \oplus a.$

By construction $F_{rc}^{-1}(F_{rc}(a, b)) = (a, b)$ and $F_{rc}(F_{rc}^{-1}(a, b)) = (a, b)$ for all inputs. Hence F_{rc} is a permutation with an explicit inverse.

Nonlinearity localization and alignment barriers. Two structural facts (purely algebraic, parameter-free) will be used later:

- XOR and rotations are linear over \mathbb{F}_2 : they preserve the *magnitude* of differential probabilities and linear correlations (possibly changing only phase or bit positions). All nonlinearity arises solely from the modular additions \boxplus_{32} .
- The rotation offsets $\{31, 24, 8, 17, 16\}$ are pairwise distinct modulo 32 and include both odd and even values, preventing trivial periodic alignment of masks across successive additions and forcing carry constraints to originate from different rotated bit-slices.

Security positioning (why this box matters). This ARX-SBOX is the only nonlinear component inside the round's diagonal ARX layer. In the overall LOPC round, it is tightly bound with the subsequent keyed switching and the ARX tail; an adversary cannot “skip” it. Consequently:

1. If the ARX-SBOX admitted a fixed high-probability differential or a non-negligible linear hull under public constants, a distinguisher for the full cipher could be built *without* touching the PRG layer. In that case the whole LOPC would be structurally compromised.
2. Conversely, if a distinguisher succeeds against the full cipher while the two PRGs remain indistinguishable from U_{64} , then that distinguisher must implicitly exploit a structural bias in this ARX-SBOX (or in the linear mixing around it). Absent such a bias, the adversary is forced back to breaking the PRG-generated round parameters.

This is precisely why the ARX-SBOX is critical: if it were unsafe, the entire LOPC construction would effectively collapse to “attack the PRG”; ensuring its algebraic soundness keeps the hardness where it belongs—at the PRG layer.

Constant-time note. All primitives here— \oplus , \boxplus_{32} , \boxdot_{32} , \lll , \ggg —are fixed-time word operations; there are no secret-dependent branches inside the layer.

8.4.4 Ideal-Theoretic Upper/Lower Bounds (No Empirical Data)

Model and notation. Let $F_{rc} : (\{0,1\}^{32})^2 \rightarrow (\{0,1\}^{32})^2$ be the NeoAlzette ARX-SBOX (Sec. ??), and let R denote one full 128-bit round that applies F_{rc} diagonally on (w_0, w_2) and (w_1, w_3) followed by the keyed switching layer and the ARX tail. For $r \geq 1$, write $E_r = R^{\circ r}$ for the r -round permutation (per fixed per-round parameters). For a permutation $P : \{0,1\}^n \rightarrow \{0,1\}^n$:

$$DP_P(\Delta \rightarrow \nabla) = \Pr_x [P(x) \oplus P(x \oplus \Delta) = \nabla], \quad MDP(P) = \max_{\Delta \neq 0, \nabla} DP_P(\Delta \rightarrow \nabla),$$

$$LC_P(\alpha, \beta) = 2^{-n} \sum_x (-1)^{\langle \alpha, x \rangle \oplus \langle \beta, P(x) \rangle}, \quad MLC(P) = \max_{\alpha \neq 0, \beta} |LC_P(\alpha, \beta)|.$$

Information-theoretic lower bounds (unavoidable). These hold for *every* n -bit permutation (no structure assumed):

$$MDP(P) \geq 2^{-n}, \quad MLC(P) \geq 2^{-n/2}.$$

Justification. Pigeonhole principle gives $\max_{\nabla} DP(\Delta \rightarrow \nabla) \geq 2^{-n}$ for each fixed $\Delta \neq 0$. Parseval's identity yields $\sum_{\beta} LC(\alpha, \beta)^2 = 1$ for each fixed $\alpha \neq 0$, hence $\max_{\beta} |LC(\alpha, \beta)| \geq 2^{-n/2}$.

Structural assumptions for ideal upper bounds. We reason under the *idealized schedule* (per-round parameters are independent uniform 64-bit draws mapped exactly as in the code) and two mild, architecture-agnostic assumptions about the ARX layer:

- **(IC) Independent-carry constraints.** In one F_{rc} layer, the seven \boxplus_{32} operations induce at least $u \geq 32$ *independent* carry-parity constraints across rotated bit-slices for any nontrivial differential trail; in linear form, the same u constraints act as balanced Boolean factors.
- **(NWD) No wraparound degeneracy.** The rotation offsets $\{31, 24, 8, 17, 16\}$ are pairwise distinct modulo 32 and include both odd and even values, so no nonzero mask/difference can align to the same bit-slice across all additions.

(IC) and (NWD) capture the exact role of \lll / \ggg around \boxplus_{32} : XOR/rotations preserve magnitude; all loss comes from carries, and they are forced to occur on distinct slices.

Theorem 8.15 (One-layer (64-bit) ideal upper bounds). *Under (IC)+(NWD) for F_{rc} ,*

$$MDP(F_{rc}) \leq 2^{-32}, \quad MLC(F_{rc}) \leq 2^{-16}.$$

Proof sketch. Write $a \boxplus_{32} b = a \oplus b \oplus \text{Carry}(a, b)$ with the carry vector determined by majority of lower bits. For a fixed $(\Delta a, \Delta b)$ and target ∇ , the event $F_{rc}(a, b) \oplus F_{rc}(a \oplus \Delta a, b \oplus \Delta b) = \nabla$ is equivalent to satisfying u independent parity equations on carry bits drawn from rotated slices (by (NWD)). Each equation holds with probability $1/2$ over uniform inputs; piling-up gives $2^{-u} \leq 2^{-32}$ for differentials. In the linear case, $\langle \beta, a \boxplus_{32} b \rangle = \langle \beta, a \rangle \oplus \langle \beta, b \rangle \oplus \langle \beta, \text{Carry}(a, b) \rangle$; the carry term is a product of u balanced Boolean factors, hence $|LC| \leq 2^{-u/2} \leq 2^{-16}$. \square

Theorem 8.16 (One full round (128-bit) ideal upper bounds). *Let R be one 128-bit round (diagonal F_{rc} on two disjoint pairs, then keyed switching and ARX tail). Under the idealized schedule and (IC)+(NWD) on each diagonal,*

$$MDP(R) \leq 2^{-32}, \quad MLC(R) \leq 2^{-16}.$$

Reason. The two diagonals are disjoint; taking maxima, the worst of the two F_{rc} controls the bound. The keyed switching layer consists of per-lane XOR/NOT/rotations (magnitude-preserving) and the ARX tail adds constants via \boxplus_{32} , which can only decrease differential probability and correlation magnitude. Hence the round's MDP/MLC are bounded by those of the diagonal F_{rc} . \square

Theorem 8.17 (r rounds: ideal sandwich bounds). *Under the idealized schedule with independent per-round parameters and (IC)+(NWD) per round,*

$$\begin{aligned} 2^{-128} &\leq MDP(E_r) \leq 2^{-\min(32r, 128)}, \\ 2^{-64} &\leq MLC(E_r) \leq 2^{-\min(16r, 64)}. \end{aligned}$$

Explanation. The lower bounds are information-theoretic for any 128-bit permutation (they cannot be beaten by any design). For the upper bounds, per round we have $MDP \leq 2^{-32}$ and $|LC| \leq 2^{-16}$; with independent per-round parameters, differential probabilities multiply and correlations multiply in magnitude, yielding 2^{-32r} and 2^{-16r} . We cap at the 128-bit floors since no cipher can be “more random” than a random 128-bit permutation in those metrics. \square

What this means in practice (default rounds $r=4$). The ideal-theoretic bounds specialize to

$$\text{MDP}(E_4) \in [2^{-128}, 2^{-128}], \quad \text{MLC}(E_4) \in [2^{-64}, 2^{-64}],$$

i.e., the round-wise structural bounds meet the 128-bit random-permutation floor after four rounds. This is independent of any brute-force enumeration and does not use single-point measurements.

Positioning relative to the PRG layer. These bounds live entirely *below* the key schedule: if an adversary distinguishes the full cipher with advantage above the upper envelopes of Theorems 8.16–8.17 while the schedule’s PRGs remain indistinguishable, then the attack necessarily exploits a structural bias in the ARX layer (contradicting (IC)+(NWD)). Otherwise, any residual advantage must be attributed to breaking the PRGs, which our top-level reduction already isolates.

8.4.5 Consolidated Single-Round Statistics

Table 1: Worst-case single-use statistics: minimum output Hamming weight $w_{\min}^{(i)}$ and the corresponding upper bound $p_{\max}^{(i)} = 2^{-w_{\min}^{(i)}}$ for each round constant RC_i (all 32 one-bit input differences, SMT search).

index i	RC_i (hex)	$w_{\min}^{(i)}$	$p_{\max}^{(i)} = 2^{-w_{\min}^{(i)}}$
0	16B2C40B	22	2.38×10^{-7}
1	C117176A	20	9.54×10^{-7}
2	0F9A2598	10	9.77×10^{-4}
3	A1563ACA	24	5.96×10^{-8}
4	243F6A88	26	1.49×10^{-8}
5	85A308D3	22	2.38×10^{-7}
6	13198102	11	4.88×10^{-4}
7	E0370734	12	2.44×10^{-4}
8	9E3779B9	20	9.54×10^{-7}
9	7F4A7C15	17	7.63×10^{-6}
10	F39CC060	13	1.22×10^{-4}
11	5CEDC834	10	9.77×10^{-4}
12	B7E15162	16	1.53×10^{-5}
13	8AED2A6A	16	1.53×10^{-5}
14	BF715880	17	7.63×10^{-6}
15	9CF4F3C7	15	3.05×10^{-5}

Worst case. The smallest value is $w_{\min} = 10$, hence a single S-box call can never exceed $p_{\max} = 2^{-10}$.

Average over all (RC_i, Δ) .

$$\bar{w}_{\min} = \frac{1}{16} \sum_{i=0}^{15} w_{\min}^{(i)} = 16.94 \implies \bar{p}_{\max} = 2^{-16.94} \approx 7.95 \times 10^{-6}.$$

Implication for a full round. With branch number $d \geq 2$, one round activates at least $d + 2 = 4$ S-boxes, whence

$$p_{\max}^{(\text{round})} \leq (2^{-10})^4 = 2^{-40} \quad (9.09 \times 10^{-13}),$$

and the recommended $R = 8$ rounds give a cipher-level bound

$$p_{\max}^{(\text{cipher})} \leq 2^{-40 \cdot 8} = 2^{-320}.$$

This is comfortably below the 2^{-128} threshold for 256-bit-key designs.

8.4.6 Structural Security Enhancements over Alzette

8.4.7 Rotation Offset Analysis

The specific rotation constants combat rotational cryptanalysis through:

- **Prime Offsets:** 31, 17 (right) and 24, 8, 16 (left) are coprime pairs

$$\gcd(31, 24) = 1, \quad \gcd(17, 8) = 1, \quad \gcd(16, 31) = 1$$

- **Direction Alteration:** Prevents fixed rotational relationships

Rotation sequence: R31 → L24 → L8 → R17 → L16

- **Carry Disruption:** Prime offsets break carry propagation patterns

$$\Pr[\text{Carry}(x \ggg 31) = \text{Carry}(x \lll 24)] \leq 2^{-8}$$

Table 2: Core Structural Comparison between Alzette and NeoAlzette S-boxes

Feature	Alzette S-box	NeoAlzette S-box
State Size	32-bit	64-bit (dual 32-bit channels)
ARX Operations/Round	4	12
Rotation Offsets	Symmetric (17L/16R)	Prime-based (31R,24L,17R,8L,16L)
Constant Injection Points	1 (post-rotation)	4 (interleaved)
Diffusion Pattern	Sequential	Cross-coupled
Nonlinearity Source	Single modular add	Dual carry chains

8.5 XCR/ZUC-Based Key Schedule: Security Review and Analysis

Abstracted Key Schedule Structure Per round $i \in \{0, \dots, r-1\}$, the schedule consumes: - one $\mathcal{G}^{(\text{key})}$ call to derive a secret-driven value, - one $\mathcal{G}^{(\text{pub})}$ call to derive a public/round-salted value, - one further $\mathcal{G}^{(\text{key})}$ call for choice bits, - one further $\mathcal{G}^{(\text{pub})}$ call for rotation amounts.

Let $m_{\text{key}} = 2r$ and $m_{\text{pub}} = 2r$ be the number of calls to each PRG. The rest of the schedule is deterministic bit-slicing and indexing.

Lemma 8.18 (Dual-PRG Schedule Security). *If $\mathcal{G}^{(\text{key})}$ and $\mathcal{G}^{(\text{pub})}$ are $(t, \epsilon_{\text{key}})$ - and $(t, \epsilon_{\text{pub}})$ -secure 64-bit PRGs respectively, then for any PPT adversary \mathcal{A} ,*

$$\text{Adv}_{\mathcal{A}}^{\text{Schedule}} \leq m_{\text{key}}\epsilon_{\text{key}} + m_{\text{pub}}\epsilon_{\text{pub}}.$$

If either PRG is instantiated with ZUC-256, its ϵ inherits from the ZUC advantage bound; if modeled as a random oracle, $\epsilon = 0$ and outputs are perfectly uniform.

Hybrid Argument. Replace each $\mathcal{G}^{(\text{key})}$ call one by one with U_{64} via a standard distinguisher. Total gap $m_{\text{key}}\epsilon_{\text{key}}$. On that world, replace each $\mathcal{G}^{(\text{pub})}$ call one by one with U_{64} , gap $m_{\text{pub}}\epsilon_{\text{pub}}$. The resulting distribution of all (sk, cf, α, β) tuples matches the ideal schedule exactly (no birthday term if the “ideal” is defined from U_{64} with the same bit-slicing as the implementation). \square

ZUC Cipher Abstract Formulation ZUC’s keystream generator can be described in three functional layers:

1. LFSR Update ($\text{mod } 2^{31} - 1$):

$$S_t = (2^{15}S_{t-15} + 2^{17}S_{t-13} + 2^{21}S_{t-10} + 2^{20}S_{t-4} + (1+2^8)S_{t-0}) \text{ mod } (2^{31} - 1)$$

with all additions in $\mathbb{Z}_{2^{31}-1}$.

2. Bit Reconstruction:

$$X_0 = (S_{t-15} \& 0x7FFF8000) \ll 1 \oplus (S_{t-14} \& 0xFFFF)$$

(similarly for X_1, X_2, X_3) — four 32-bit words formed from selected bit-fields of the LFSR state.

3. F-Function:

$$W = (X_0 \oplus R_1) + R_2 \pmod{2^{32}}$$

$$W_1 = R_1 + X_1 \pmod{2^{32}}, \quad W_2 = R_2 \oplus X_2$$

$$R'_1 = S\text{-box}(L1(W_1)), \quad R'_2 = S\text{-box}(L2(W_2))$$

with $S0/S1$ and $L1, L2$ as in the original ZUC spec.

This decomposition is equivalent to the “Three “no” abstractions: no modification of LFSR polynomials, no modification of bit reconstruction rules, and no modification of the structure of F-functions.”

Instantiation Cases - **ZUC Mode**: ϵ_{pub} and/or ϵ_{key} are bounded by the best known ZUC security results (e.g., $\epsilon_{\text{ZUC}} \approx 2^{-128}$ for 256-bit ZUC). - **Random Oracle Mode**: $\epsilon = 0$, outputs are uniform by definition.

Concrete Bound Example For $r = 4$ rounds and both PRGs instantiated as ZUC-256 ($\epsilon_{\text{ZUC}} = 2^{-128}$):

$$\text{Adv} \leq 2r \epsilon_{\text{ZUC}} + 2r \epsilon_{\text{ZUC}} = 4r \epsilon_{\text{ZUC}} = 2^{-126}$$

If one PRG is RO, drop its term accordingly.

Implementation Observations - Master key K : 64-bit, but effective strength also gated by PRG state (256 bits for XCR/ZUC). - Choice bits: lowest 2 bits of a 64-bit PRG output. - Rotation amounts: two disjoint 6-bit slices of a 64-bit PRG output.

8.5.1 XCR/ZUC-Driven Key Schedule for LOPC Wrapper: Clean Two-PRG Hybrid Reduction

Code-Induced Schedule (from `GenerateAndStoreKeyStates`) For each round $i \in \{0, \dots, r-1\}$, the implementation derives:

$$\begin{aligned}
(\text{seed}) \quad & \text{seed} \leftarrow f(K) \quad \text{with } f(K) = (a \oplus (b \lll_{64} 1)) \oplus (a \lll_{64} 13), \\
& a = K_0 \oplus K_1, \quad b = \neg \text{ghash_multiply}(K_0, K_1)(\text{128-bit input}, \text{64-bit output}), \\
(\text{PRGs}) \quad & \mathcal{G}^{(\text{key})} \leftarrow \mathcal{G}_{\text{XCR}}(\text{seed}), \quad \mathcal{G}^{(\text{pub})} \in \{\mathcal{G}_{\text{XCR}}, \mathcal{G}_{\text{ZUC}}\}, \\
(\text{calls}) \quad & \begin{cases} (1) \text{sk}_0 \leftarrow \mathcal{G}^{(\text{key})}(\text{number_once}), \\ (2) \text{para} \leftarrow \mathcal{G}^{(\text{pub})}(\text{number_once} \oplus i), \\ (3) \text{choice} \leftarrow \mathcal{G}^{(\text{key})}((K_0 \oplus \text{sk}_0) \oplus (K_0 \gg 1)) \bmod 4, \\ (4) \text{bitrot} \leftarrow \mathcal{G}^{(\text{pub})}((K_1 \oplus \text{para}) \oplus \text{choice}), \end{cases} \\
(\text{fields}) \quad & \begin{aligned} \text{subkey.first} &= K_0 \oplus \text{sk}_0, & \text{subkey.second} &= K_1 \oplus \text{para}, \\ \alpha_i &= \text{bitrot} \bmod 64, & \beta_i &= (\text{bitrot} \gg 6) \bmod 64, \\ \text{index}_i &= (\text{rc.index}^{\text{raw}} \gg 1) \bmod 16, & \text{rc.index}^{\text{raw}} &\leftarrow \text{rc.index}^{\text{raw}} + 2. \end{aligned}
\end{aligned}$$

Each round consumes exactly 4 64-bit PRG outputs: two from $\mathcal{G}^{(\text{key})}$ and two from $\mathcal{G}^{(\text{pub})}$. Let $m_{\text{key}} = 2r$, $m_{\text{pub}} = 2r$.

Adversarial Experiment The adversary \mathcal{A} , in the scheme $\Pi_{\mathcal{G}^{(\text{key})}, \mathcal{G}^{(\text{pub})}} + \Pi_{\text{LOPC}}$, issues at most q encryption queries, aiming to distinguish it from an *idealized baseline*: in the baseline, the *same structural steps* as above are followed, but the four PRG calls are replaced by independent U_{64} samples (with $\text{choice} = U_{64} \bmod 4$ preserved, and (α_i, β_i) derived from $\Pi_{\mathcal{G}^{(\text{key})}, \mathcal{G}^{(\text{pub})}} + \Pi_{\text{LOPC}}$ non-overlapping 6-bit slices of the same U_{64}). The distinguishing advantage is denoted $\text{Adv}_{\mathcal{A}}^{\Pi_{\mathcal{G}^{(\text{key})}, \mathcal{G}^{(\text{pub})}} + \Pi_{\text{LOPC}}}(q, r)$.

Lemma 8.19 (Uniform-Preserving Slices). *For any fixed K , if $Y \sim U_{64}$ independently, then $K \oplus Y$ is uniform over $\{0, 1\}^{64}$; $Y \bmod 4$ is uniform over $\{0, 1, 2, 3\}$; and the two non-overlapping 6-bit slices ($Y \bmod 64$, $(Y \gg 6) \bmod 64$) are independent and uniform.*

Theorem 8.20 (Two-PRG Hybrid Reduction to Baseline LOPC). *Let ϵ_{key} bound the advantage of any PPT distinguisher in at most m_{key} valid queries to $\mathcal{G}^{(\text{key})}$, and ϵ_{pub} bound the advantage in at most m_{pub} valid queries to $\mathcal{G}^{(\text{pub})}$. Then:*

$$\text{Adv}_{\mathcal{A}}^{\Pi_{\mathcal{G}^{(\text{key})}, \mathcal{G}^{(\text{pub})}} + \Pi_{\text{LOPC}}}(q, r) \leq \underbrace{m_{\text{key}} \cdot \epsilon_{\text{key}}}_{\text{Replace key-PRG}} + \underbrace{m_{\text{pub}} \cdot \epsilon_{\text{pub}}}_{\text{Replace pub-PRG}} + \underbrace{\epsilon_{\text{LOPC}}(q, r)}_{\text{LOPC baseline with independent round material}}.$$

Clean Hybrid (Game-Hopping). **Phase A (key-PRG replacement).** Enumerate all m_{key} calls to $\mathcal{G}^{(\text{key})}$ as $Y_1^{(\text{key})}, \dots, Y_{m_{\text{key}}}^{(\text{key})}$ in order of appearance. Construct hybrids $H_j^{(\text{key})}$: the first j calls are replaced by U_{64} , the rest remain genuine $\mathcal{G}^{(\text{key})}$; all other calls remain unchanged. For each hop $H_j^{(\text{key})} \rightarrow H_{j+1}^{(\text{key})}$, build the standard distinguisher $\mathcal{D}_j^{(\text{key})}$ that samples only the $(j+1)$ -th call from the external oracle. Then: $|\Pr[H_j^{(\text{key})}] - \Pr[H_{j+1}^{(\text{key})}]| \leq \epsilon_{\text{key}}$. Summing yields $|\Pr[H_0^{(\text{key})}] - \Pr[H_{m_{\text{key}}}^{(\text{key})}]| \leq m_{\text{key}} \epsilon_{\text{key}}$.

Phase B (pub-PRG replacement). In $H_{m_{\text{key}}}^{(\text{key})}$, similarly replace the m_{pub} calls to $\mathcal{G}^{(\text{pub})}$, $Y_1^{(\text{pub})}, \dots, Y_{m_{\text{pub}}}^{(\text{pub})}$, producing $H_0^{(\text{pub})} \rightarrow H_{m_{\text{pub}}}^{(\text{pub})}$ and corresponding $\mathcal{D}_j^{(\text{pub})}$. Likewise: $|\Pr[H_0^{(\text{pub})}] - \Pr[H_{m_{\text{pub}}}^{(\text{pub})}]| \leq m_{\text{pub}} \epsilon_{\text{pub}}$.

Phase C (alignment with ideal round material). At this point, all four call types are provided by independent U_{64} samples. By Lemma 8.19, after the deterministic mappings (XOR, modulus, slicing), the resulting $(\text{subkey}, \text{choice}, \alpha_i, \beta_i, \text{index}_i)$ are *exactly* as in the baseline definition, with zero additional statistical distance.

Phase D (switch to LOPC baseline security). Compare the real Π_{LOPC} with the “baseline LOPC under independent round material”; the difference is exactly $\epsilon_{\text{LOPC}}(q, r)$. Summing the three phases yields the theorem. \square

Minimal Non-Confusing Distinguisher For any hop step (e.g., pub-PRG case): $\mathcal{D}_j^{(\text{pub})}(\mathcal{O})$: Simulate \mathcal{A} interacting with the full encryption process, answering all $\mathcal{G}^{(\text{pub})}$ calls in order, where the $(j+1)$ -th call is provided by oracle $\mathcal{O} \in \{\mathcal{G}^{(\text{pub})}, U_{64}\}$ and the rest follow the real implementation. Output \mathcal{A} 's bit. This single-call substitution bounds the hop gap.

Two-Mode Instantiation and Concrete Counts In this implementation, each round consumes $m_{\text{key}} = 2$, $m_{\text{pub}} = 2$, totalling $4r$ 64-bit calls. - If $\mathcal{G}^{(\text{key})} = \mathcal{G}_{\text{XCR}}$ and $\mathcal{G}^{(\text{pub})} = \mathcal{G}_{\text{XCR}}$, the advantage bound is $\leq 2r \epsilon_{\text{XCR}} + 2r \epsilon_{\text{XCR}} + \epsilon_{\text{LOPC}} = 4r \epsilon_{\text{XCR}} + \epsilon_{\text{LOPC}}$. - If $\mathcal{G}^{(\text{key})} = \mathcal{G}_{\text{XCR}}$ and $\mathcal{G}^{(\text{pub})} = \mathcal{G}_{\text{ZUC}}$, the bound is $\leq 2r \epsilon_{\text{XCR}} + 2r \epsilon_{\text{ZUC}} + \epsilon_{\text{LOPC}}$.

Notes on Seeding and Rotations seed is derived from (a, b) via \oplus and $\text{rotl}(\cdot)$ combinations; this is a deterministic key-derivation step that only affects the *initial state* of $\mathcal{G}^{(\text{key})}$, and introduces no additional assumptions in the reduction. The rotation constants (1, 13) are chosen for engineering diffusion trade-offs and are *not* part of the security assumption set.

8.6 Keyed Switching Layer as $T \circ L \circ S$: Minimal Facts Needed for the Reduction

Definition (per 64-bit lane). Let $x \in \{0, 1\}^{64}$. A round- i KSL is the composition

$$\text{KSL}^{(i)} = T^{(i)} \circ L^{(i)} \circ S^{(i)},$$

with parameters $(\text{sk}_i, \text{cf}_i, \beta_i, \alpha_i)$ produced by the key schedule.

Switch $S^{(i)}$ (*branch chosen by* $cf_i \in \{0, 1, 2, 3\}$):

$$S^{(i)}(x) = \begin{cases} x \mapsto x \oplus sk_i & cf_i = 0, \\ x \mapsto \neg x \oplus sk_i & cf_i = 1, \\ x \mapsto \text{rotl}(x, \beta_i) & cf_i = 2, \\ x \mapsto \text{rotr}(x, \beta_i) & cf_i = 3, \end{cases} \quad sk_i \in \{0, 1\}^{64}, \beta_i \in \{0, \dots, 63\}.$$

Linear $L^{(i)}$ (*Binary PHT on two 32-bit words*). Writing $x = \text{pack}_{64}(w_0, w_1)$ with $w_j \in \mathbb{Z}_{2^{32}}$, $L^{(i)}$ applies an invertible linear map over $\mathbb{Z}_{2^{32}}^2$:

$$(w_0, w_1) \mapsto (w'_0, w'_1) = \mathcal{H}(w_0, w_1), \det(\mathcal{H}) \equiv 1 \pmod{2},$$

and returns $\text{pack}_{64}(w'_0, w'_1)$. (Exact coefficients are implementation-defined; we only use invertibility.)

Tweak $T^{(i)}$ (*one-bit XOR masks*). Per lane- $\ell \in \{0, 1\}$,

$$T_\ell^{(i)}(x) = x \oplus e_{\alpha_i}^{(\ell)}, \quad e_{\alpha}^{(0)} = 1 \ll \alpha, \quad e_{\alpha}^{(1)} = 1 \ll (63 - \alpha), \quad \alpha_i \in \{0, \dots, 63\}.$$

Thus lane 0 flips bit α_i , lane 1 flips the complementary bit $63 - \alpha_i$.

Permutation and invertibility. For any fixed parameters, $S^{(i)}$, $L^{(i)}$, and $T^{(i)}$ are bijections on $\{0, 1\}^{64}$; hence $\text{KSL}^{(i)}$ is a permutation and $(\text{KSL}^{(i)})^{-1} = (S^{(i)})^{-1} \circ (L^{(i)})^{-1} \circ (T^{(i)})^{-1}$ with

$$(S^{(i)})^{-1} = \begin{cases} x \mapsto x \oplus sk_i & cf_i = 0, \\ x \mapsto \neg x \oplus sk_i & cf_i = 1, \\ x \mapsto \text{rotr}(x, \beta_i) & cf_i = 2, \\ x \mapsto \text{rotl}(x, \beta_i) & cf_i = 3, \end{cases} \quad (L^{(i)})^{-1} = \mathcal{H}^{-1}, \quad (T^{(i)})^{-1} = T^{(i)}.$$

Parameter distribution under the idealized schedule. When the schedule's PRG outputs are hybrid-replaced by U_{64} (Sec. 8.5), we have

$$sk_i \xleftarrow{\text{i.i.d.}} U_{64}, \quad cf_i \xleftarrow{\text{i.i.d.}} U_{\{0, 1, 2, 3\}}, \quad \beta_i \xleftarrow{\text{i.i.d.}} U_{\{0, \dots, 63\}}, \quad \alpha_i \xleftarrow{\text{i.i.d.}} U_{\{0, \dots, 63\}},$$

and the two-lane tweaks use complementary masks $e_{\alpha_i}^{(0)}$ and $e_{\alpha_i}^{(1)}$.

What is *not* needed (and thus not assumed). We do not rely on any stand-alone upper-bounds on linear/differential biases through L or T . The reduction only uses: (i) $\text{KSL}^{(i)}$ is a permutation (no entropy loss), and (ii) its parameters are independent and uniformly distributed once PRG calls are replaced by U_{64} .

Provable alignment facts we *can* and *do* use. Let $\langle \cdot, \cdot \rangle$ denote the \mathbb{F}_2 inner product on $\{0, 1\}^{64}$, and write $\text{corr}(\langle u, y \rangle, \langle v, x \rangle) = \mathbb{E}_x [(-1)^{\langle u, y \rangle \oplus \langle v, x \rangle}]$.

(A) *Switch-only exact conditions.* For fixed cf_i, β_i and uniform sk_i :

$$\text{corr} \neq 0 \iff \begin{cases} u = v & \text{if } cf_i \in \{0, 1\}, \\ u = \text{rot}_{\mp \beta_i}(v) & \text{if } cf_i = 2/3. \end{cases}$$

Hence, averaging over cf_i and β_i drawn as above,

$$\Pr[\text{corr} \neq 0] \leq \begin{cases} 1/2 & \text{if } u = v, \\ 1/2 \cdot (1/64) & \text{if } u \in \{\text{rot}_\gamma(v)\}_\gamma, \\ 0 & \text{otherwise.} \end{cases}$$

(The random sk_i only flips the sign in cases 0/1.)

(B) *Differential through S.* For XOR-difference $\Delta_x \neq 0$,

$$\Delta_y = S^{(i)}(x \oplus \Delta_x) \oplus S^{(i)}(x) = \begin{cases} \Delta_x & cf_i \in \{0, 1\}, \\ \text{rotl}(\Delta_x, \beta_i) & cf_i = 2, \\ \text{rotr}(\Delta_x, \beta_i) & cf_i = 3. \end{cases}$$

So for a fixed target Δ_y^* , $\Pr[\Delta_y = \Delta_y^*]$ equals $1/2$ if $\Delta_y^* = \Delta_x$, or $1/2 \cdot (1/64)$ if Δ_y^* is a rotation of Δ_x , else 0.

(C) *Effect of L and T.* L is an invertible map over $\mathbb{Z}_{2^{32}}^2$; it is not \mathbb{F}_2 -linear due to carries, hence it *can only decrease or scramble* any fixed bit-level linear/differential alignment induced by S ; we do not lower-bound that decrease. T is an XOR with a fixed (per-round) constant mask: it *does not* change differentials at all (Δ cancels), and it flips the *phase* (sign) of any linear correlation by $(-1)^{\langle u, e_{\alpha_i}^{(\ell)} \rangle}$ without changing its magnitude. These two facts mean L/T do not create new exploitable alignment; for the proof we only need that they preserve permutation-ness and parameter uniformity.

Two-lane coupling via complementary tweak. The pair $T_0^{(i)}$ and $T_1^{(i)}$ uses complementary bit positions $(\alpha_i, 63 - \alpha_i)$. Any trail that attempts to lock bit-level alignments *across lanes* must therefore satisfy two independent phase constraints per round (one per lane). This observation is *design rationale* and is not used by the reduction.

How this plugs into the main reduction. Since each $\text{KSL}^{(i)}$ is a permutation keyed only by schedule parameters, once the schedule's PRG outputs are replaced by U_{64} , all $(sk_i, cf_i, \beta_i, \alpha_i)$ become i.i.d. uniform (with the lane-wise complement relation for T). Therefore the whole round becomes the baseline LOPC round with independent per-round materials, and Theorem 8.20 applies verbatim:

$$\text{Adv}_{\mathcal{A}}^{\Pi_{\mathcal{G}(\text{key}), \mathcal{G}(\text{pub})} + \Pi_{\text{LOPC}}}(q, r) \leq 2r \epsilon_{\text{key}} + 2r \epsilon_{\text{pub}} + \epsilon_{\text{LOPC}}(q, r).$$

Constant-time implementation note. Branch selection for S (`switch (choice.function & 3ULL)`) is driven by per-round constants from the key schedule PRGs, not by secret-dependent bits from the plaintext or the master key. Each branch consists of fixed-time word operations (XOR, NOT+XOR, rotations). Thus the KSL preserves constant-time execution and does not introduce a timing side channel.

8.7 Precise Key Mixing(Add / Subtract Round Key) Analysis

Let \lll and \ggg denote left/right bit rotation.

Definition 8.3 (Encryption Key Mixing). *For plain-text $x \in \{0, 1\}^{64}$, subkey $sk \in \{0, 1\}^{64}$, and master key $k \in \{0, 1\}^{64}$:*

$$\begin{aligned} y_1 x \boxplus_{64} (k \oplus sk) \\ y_2 y_1 \oplus k \\ y_3 y_2 \ggg 16 \\ resulty_3 \oplus ((k \boxplus_{64} sk) \lll 48) \end{aligned}$$

Definition 8.4 (Decryption Key Mixing). *For cipher-text $x' \in \{0, 1\}^{64}$:*

$$\begin{aligned} y'_1 x' \oplus ((k \boxplus_{64} sk) \lll 48) \\ y'_2 y'_1 \lll 16 \\ y'_3 y'_2 \oplus k \\ resulty'_3 \boxplus_{64} (k \oplus sk) \end{aligned}$$

8.7.1 Algebraic Analysis

The inverse relationship holds when:

Lemma 8.21 (ARX Invertibility). *For valid (k, sk) , the encryption/decryption functions satisfy:*

$$\forall x \in \{0, 1\}^{64}, \text{ Decrypt}(\text{Encrypt}(x)) = x.$$

Proof. Let $\Delta_k = k \oplus sk$, $R = \ggg 16$ and $L = \lll 16$. Write the encryption as

$$E(x) = ((x \boxplus \Delta_k) \oplus (kR)) \oplus ((k \boxplus sk) \lll 48).$$

The decryption proceeds by inverting each step in reverse order:

$$\begin{aligned} y &\leftarrow E(x) \\ &\xrightarrow{\oplus (k \boxplus sk) \lll 48} (x \boxplus \Delta_k) \oplus (kR) \\ &\xrightarrow{L} ((x \boxplus \Delta_k)L) \oplus ((kR)L) \\ &\xrightarrow{\oplus k} (x \boxplus \Delta_k)L \\ &\xrightarrow{R} x \boxplus \Delta_k \\ &\xrightarrow{\boxminus \Delta_k} x. \end{aligned}$$

Here we used $(kR)L = k$ (rotations are inverses and linear over \mathbb{F}_2) to cancel the k term after the left rotation. The final right rotation R is essential to undo the earlier left rotation on $(x \boxplus \Delta_k)$, since rotation does not commute with modular addition. \square

Modular $+/-$: Overflow (Carry) & Underflow (Borrow) — Closed Forms

Setup. Let $n \geq 1$, modulus $m := 2^n$. For $A, B \in \{0, 1, \dots, m-1\}$ define

$$C^+(A + B) \bmod m, \quad C^-(A - B) \bmod m.$$

Define indicator variables

$$K^+ \mathbf{1}\{A + B \geq m\} \quad (\text{addition overflow/carry}), \quad K^- \mathbf{1}\{A < B\} \quad (\text{subtraction underflow/borrow}).$$

Baseline: A, B i.i.d. uniform on $\{0, \dots, m-1\}$.

$$\Pr[K^+ = 1] = \Pr[K^- = 1] = \frac{m-1}{2m} = \frac{2^n - 1}{2^{n+1}}. \quad (1)$$

Sketch. $\#\{(a, b) : a + b \geq m\} = \sum_{a=0}^{m-1} a = m(m-1)/2$; similarly $\#\{(a, b) : a < b\} = m(m-1)/2$.

Fixed constant A , uniform B , and the two bijective chains. Let $B_0 \sim \text{Unif}(\{0, \dots, m-1\})$ and define either chain

$$B_{t+1} = T_+(B_t)(B_t + A) \bmod m \quad \text{or} \quad B_{t+1} = T_-(B_t)(A - B_t) \bmod m.$$

Both T_+, T_- are permutations of $\{0, \dots, m-1\}$, hence B_t is uniform for all t . Therefore, for every step t ,

$$\Pr[K^+ = 1 \mid A] = \frac{A}{m}, \quad \Pr[K^- = 1 \mid A] = \frac{m-A-1}{m}. \quad (2)$$

Boundary checks: $A = 0 \Rightarrow \Pr[K^+ = 1] = 0$, $\Pr[K^- = 1] = (m-1)/m$; $A = m-1 \Rightarrow \Pr[K^+ = 1] = 1 - 2^{-n}$, $\Pr[K^- = 1] = 0$.

Expected number of carries/borrows in one n -bit operation (uniform A, B). Let C_i be the carry out of bit i in the addition of A and B (LSB is $i = 0$), and $p_i \Pr(\text{carry-in to bit } i = 1)$ with $p_0 = 0$. Conditioning on carry-in gives

$$\Pr(C_i = 1 \mid \text{carry-in} = 0) = \frac{1}{4}, \quad \Pr(C_i = 1 \mid \text{carry-in} = 1) = \frac{3}{4},$$

hence $p_{i+1} = \Pr(C_i = 1) = \frac{1}{4} + \frac{1}{2} p_i$. Solving yields $p_i = \frac{1}{2} - 2^{-(i+1)}$ and

$$\mathbb{E}[C_i] = \frac{1}{4} + \frac{1}{2} p_i = \frac{1}{2} - 2^{-(i+2)}. \quad (3)$$

Thus the expected total number of carries is

$$\mathbb{E}\left[\sum_{i=0}^{n-1} C_i\right] = \sum_{i=0}^{n-1} \left(\frac{1}{2} - 2^{-(i+2)}\right) = \frac{n}{2} - \frac{1}{2} + 2^{-(n+1)}. \quad (4)$$

For subtraction, the borrow-out obeys the same recursion, so

$$\mathbb{E}[\#\text{borrows}] = \frac{n}{2} - \frac{1}{2} + 2^{-(n+1)}. \quad (5)$$

Initial ripple-chain length from LSB (uniform A, B). Let L^+ be the number of consecutive bits starting at the LSB for which the addition produces carry-out = 1; define L^- analogously for borrow. Then

$$\Pr(L^+ = 0) = \frac{3}{4}, \quad (1)$$

$$\Pr(L^+ = k) = \frac{1}{16} \left(\frac{3}{4}\right)^{k-1}, \quad k \geq 1, \quad (2)$$

$$\mathbb{E}[L^+] = 1. \quad (6)$$

(The same holds for L^- .) For finite n , truncate at $k \leq n$; the tail beyond n is negligible due to the geometric decay.

Monte Carlo estimators (if needed). Given i.i.d. samples (A_j, B_j) or chain states $B_j^{(t)}$, the plug-in estimators are

$$\hat{p}_{\text{carry}} = \frac{1}{N} \sum_{j=1}^N \mathbf{1}\{A_j + B_j \geq m\}, \quad (3)$$

$$\hat{p}_{\text{borrow}} = \frac{1}{N} \sum_{j=1}^N \mathbf{1}\{A_j < B_j\}, \quad (4)$$

$$\hat{\mathbb{E}}[\#\text{carries}] = \frac{1}{N} \sum_{j=1}^N \sum_{i=0}^{n-1} C_i^{(j)}, \quad \hat{\mathbb{E}}[\#\text{borrows}] = \frac{1}{N} \sum_{j=1}^N \sum_{i=0}^{n-1} B_i^{(j)}. \quad (7)$$

All formulas above are modulo- m ; any intermediate negatives in subtraction are reduced by $x \mapsto x \bmod m$.

8.7.2 Explicit Matrix Analysis for 8-bit Model (Matrix-Exact, Piecewise- and Jacobian-Linear with Random-Walk Probabilistic Extension)

Bit-rotation as global permutation matrices. For $n=8$, the left/right rotations $\lll_8 r, \ggg_8 r$ admit global \mathbb{F}_2 -linear representations by permutation matrices $P_{\lll_8 r}, P_{\ggg_8 r} \in \mathbb{F}_2^{8 \times 8}$:

$$(P_{\lll_8 r})_{j,i} = \mathbf{1}[j \equiv i-r \pmod{8}], \quad (P_{\ggg_8 r})_{j,i} = \mathbf{1}[j \equiv i+r \pmod{8}].$$

Example (right rotation by 2):

$$P_{\ggg_2} = \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 \\ 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 \end{bmatrix}.$$

XOR with a constant as an augmented affine matrix. For $a \in \{0, 1\}^8$, the map $X_a : b \mapsto b \oplus a$ is affine. In homogeneous (augmented) form:

$$\begin{bmatrix} X_a(b) \\ 1 \end{bmatrix} = \underbrace{\begin{bmatrix} I_8 & a \\ 0^\top & 1 \end{bmatrix}}_{\tilde{X}(a)} \begin{bmatrix} b \\ 1 \end{bmatrix}, \quad \tilde{X}(a)^{-1} = \tilde{X}(a) \quad (\text{over } \mathbb{F}_2).$$

Modular addition by a constant: two exact matrix viewpoints. Fix $y \in \{0, 1\}^8$ and define $F_y(x) = x \boxplus_8 y$ with carry recursion

$$c_0 = 0, \quad s_i = x_i \oplus y_i \oplus c_i, \quad c_{i+1} = (x_i \wedge y_i) \oplus (x_i \wedge c_i) \oplus (y_i \wedge c_i), \quad i = 0, \dots, 7,$$

where $s = F_y(x)$ and c_8 is discarded.

(A) Piecewise-affine (carry-conditioned) matrix. Let the *carry mask* $m = (m_0, \dots, m_7)$ be the produced carries (c_0, \dots, c_7) for a given input x ; denote the corresponding slice

$$\Omega_m(y) = \{x \in \{0, 1\}^8 \mid \text{the carry sequence for } (x, y) \text{ equals } m\}.$$

On $\Omega_m(y)$ the carries are constants, so $F_y(x) = x \oplus (y \oplus m)$ and therefore, in homogeneous form,

$$\forall x \in \Omega_m(y) : \begin{bmatrix} F_y(x) \\ 1 \end{bmatrix} = \underbrace{\begin{bmatrix} I_8 & (y \oplus m) \\ 0^\top & 1 \end{bmatrix}}_{\tilde{A}_{\boxplus}(y, m)} \begin{bmatrix} x \\ 1 \end{bmatrix}, \quad \tilde{A}_{\boxplus}(y, m)^{-1} = \tilde{A}_{\boxplus}(y, m).$$

Thus every slice is governed by a *single* (9×9) matrix with linear part I_8 .

(B) Boolean Jacobian (local linearization) matrix. Define the Boolean Jacobian $J_{F_y}(x)$ by $[J_{F_y}(x)]_{i,j} = \partial s_i / \partial x_j$ (Boolean derivative). Then:

Lemma 8.22 (Triangular structure and full rank). *For all $x, y \in \{0, 1\}^8$, $J_{F_y}(x)$ is strictly lower triangular plus identity:*

$$[J_{F_y}(x)]_{i,j} = 0 \quad (j > i), \quad [J_{F_y}(x)]_{i,i} = 1, \quad \text{rank } J_{F_y}(x) = 8.$$

Proof. s_i depends only on (x_0, \dots, x_i) and $\partial s_i / \partial x_i = 1$ since c_i is independent of x_i . Hence lower-triangular with 1 on the diagonal; full rank follows. \square

A fully worked 8-bit numerical slice (visual). Let $x = 00000001_2$, $k = 00001111_2$, $sk = 11110000_2$, $\Delta = k \oplus sk = 11111111_2$. Adding Δ to x yields carry mask

$$m = (c_0, \dots, c_7) = (0, 1, 1, 1, 1, 1, 1, 1) \equiv 11111110_2,$$

so on this slice

$$\tilde{A}_{\boxplus}(\Delta, m) = \begin{bmatrix} I_8 & (\Delta \oplus m) \\ 0^\top & 1 \end{bmatrix} = \begin{bmatrix} I_8 & 00000001_2 \\ 0^\top & 1 \end{bmatrix}, \quad \begin{bmatrix} y_1 \\ 1 \end{bmatrix} = \tilde{A}_{\boxplus}(\Delta, m) \begin{bmatrix} x \\ 1 \end{bmatrix} \Rightarrow y_1 = 00000000_2.$$

Then

$$\begin{bmatrix} y_2 \\ 1 \end{bmatrix} = \tilde{X}(k) \begin{bmatrix} y_1 \\ 1 \end{bmatrix}, \quad \begin{bmatrix} y_3 \\ 1 \end{bmatrix} = \tilde{R}_{\ggg 2} \begin{bmatrix} y_2 \\ 1 \end{bmatrix}, \quad \begin{bmatrix} \text{result} \\ 1 \end{bmatrix} = \tilde{X}((k \boxplus_8 sk) \lll 6) \begin{bmatrix} y_3 \\ 1 \end{bmatrix},$$

where $\tilde{R}_{\ggg 2} = \text{diag}(P_{\ggg 2}, 1)$. This reproduces exactly the bitwise computation in the running example while keeping every step *matrix-visual*.

One-round 8-bit ARX as an augmented matrix product (per slice). Fix the first-addition carry mask m for (x, Δ) . The encryption mapping

$$x \mapsto ((x \boxplus_8 \Delta) \oplus k) \ggg 2 \oplus ((k \boxplus_8 sk) \lll 6)$$

admits the (9×9) product

$$\tilde{\mathcal{E}}_m = \underbrace{\tilde{X}((k \boxplus_8 sk) \lll 6)}_{\text{affine}} \underbrace{\tilde{R}_{\ggg 2}}_{\text{linear}} \underbrace{\tilde{X}(k)}_{\text{affine}} \underbrace{\tilde{A}_{\boxplus}(\Delta, m)}_{\text{affine}}, \quad \begin{bmatrix} \mathcal{E}(x) \\ 1 \end{bmatrix} = \tilde{\mathcal{E}}_m \begin{bmatrix} x \\ 1 \end{bmatrix}.$$

Theorem 8.23 (Invertibility on every slice). *Each factor above is invertible over \mathbb{F}_2 ; hence $\tilde{\mathcal{E}}_m$ is invertible with*

$$\tilde{\mathcal{E}}_m^{-1} = \tilde{A}_{\boxplus}(\Delta, m) \tilde{X}(k) \tilde{R}_{\lll 2} \tilde{X}((k \boxplus_8 sk) \lll 6),$$

which is the augmented matrix form of the decryption sequence.

Proof. $\tilde{X}(\cdot)$ and $\tilde{A}_{\boxplus}(\cdot, \cdot)$ square to identity; $\tilde{R}_{\ggg 2}^{-1} = \tilde{R}_{\lll 2}$. The product of invertible matrices is invertible. \square

Why “singular matrices” appear in naive linear models (and why we avoid them). We first make precise why modular addition/subtraction over \mathbb{F}_2^n is *nonlinear* (because of carries/borrows), and then show how these operations nevertheless admit *explicit matrix models* that are exact either on carry/borrow-conditioned slices (augmented affine matrices) or in a lifted feature space (nested binary matrices). This provides both visual matrix expressions and a rigorous pathway to an abstract security model as a (large) family of matrix-chosen functions.

Nonlinearity of \boxplus_n and \boxminus_n over \mathbb{F}_2^n . Fix $y \in \{0, 1\}^n$. Let $F_y(x) = x \boxplus_n y$ be n -bit modular addition with carry recursion

$$c_0 = 0, \quad s_i = x_i \oplus y_i \oplus c_i, \quad c_{i+1} = (x_i \wedge y_i) \oplus (x_i \wedge c_i) \oplus (y_i \wedge c_i),$$

and $F_y(x) = (s_0, \dots, s_{n-1})$. In Boolean ANF, $\deg(s_i) = i+1$ for $i \geq 0$; thus $\deg(F_y) \geq 2$ whenever $y \neq 0$. Therefore:

There is no single $M \in \mathbb{F}_2^{n \times n}$, $b \in \mathbb{F}_2^n$ with $F_y(x) = Mx \oplus b$ for all x , unless $y = 0$.

Equivalently, the Boolean derivative $D_h F_y(x) = F_y(x \oplus h) \oplus F_y(x)$ depends on x via carry terms, whereas any linear map satisfies $D_h L(x) = Lh$ independent of x . The same argument applies to modular subtraction $G_y(x) = x \boxminus_n y$ with borrows.

Two exact matrix representations that *do* work.

- (**Slice-augmented affine model**). For each carry mask $m = (m_0, \dots, m_{n-1})$ produced by (x, y) , define the slice

$$\Omega_m(y) = \{x \in \{0, 1\}^n \mid (c_0, \dots, c_{n-1}) = m\}.$$

On $\Omega_m(y)$ the carries are constants, so $F_y(x) = x \oplus (y \oplus m)$ is *affine*. In homogeneous coordinates $(n+1)$ -dimensional:

$$\forall x \in \Omega_m(y) : \begin{bmatrix} F_y(x) \\ 1 \end{bmatrix} = \underbrace{\begin{bmatrix} I_n & (y \oplus m) \\ 0^\top & 1 \end{bmatrix}}_{\tilde{A}_{\boxplus}(y, m)} \begin{bmatrix} x \\ 1 \end{bmatrix}, \quad \tilde{A}_{\boxplus}(y, m)^{-1} = \tilde{A}_{\boxplus}(y, m).$$

Similarly, for subtraction on a borrow slice $b = (b_0, \dots, b_{n-1})$,

$$\forall x \in \Omega^\Delta(y, b) : \begin{bmatrix} G_y(x) \\ 1 \end{bmatrix} = \underbrace{\begin{bmatrix} I_n & (y \oplus b) \\ 0^\top & 1 \end{bmatrix}}_{\tilde{A}_{\boxminus}(y, b)} \begin{bmatrix} x \\ 1 \end{bmatrix}.$$

Visual impact. These are explicit $(n+1) \times (n+1)$ binary matrices: an I_n block with a full-column “offset” $(y \oplus m)$ (or $(y \oplus b)$) and a bottom-right 1.

- (**Nested binary matrices in a lifted feature space**). Let $\Phi_2(x, c)$ collect, for each bit i , the linear coordinates x_i, c_i and the quadratic monomials $\{x_i y_i, x_i c_i, y_i c_i\}$ together with a trailing 1. Then the bit-slice update

$$(s_i, c_{i+1}) = (x_i \oplus y_i \oplus c_i, x_i y_i \oplus x_i c_i \oplus y_i c_i)$$

is *linear* over \mathbb{F}_2 in Φ_2 . Writing the i -th full-adder as a block matrix $S_i \in \mathbb{F}_2^{d \times d}$ (where $d = \dim \Phi_2$) that updates the (x_i, c_i) and monomial coordinates while leaving other coordinates fixed, the whole addition by y equals the *nested product*

$$\underbrace{S_{n-1} \cdots S_1 S_0}_{\widetilde{\text{ADD}}_y} \Phi_2(x, c_0), \quad c_0 = 0.$$

Visual impact. Each S_i is block lower-bidiagonal (staircase form); stacking $S_{n-1} \cdots S_0$ yields a large, explicit binary matrix $\widetilde{\text{ADD}}_y$ acting linearly on Φ_2 . XOR and rotations are already linear; in homogeneous form they become block-diagonal with an extra 1.

Putting it together: ARX as a matrix product (per slice or lifted). Let $P_{\ggg r}$ denote the $n \times n$ rotation matrix, and $\tilde{X}(a) = \begin{bmatrix} I_n & a \\ 0^\top & 1 \end{bmatrix}$ the augmented XOR-by- a matrix. For the 8-bit example

$$x \mapsto ((x \boxplus_8 \Delta) \oplus k) \ggg_8 2 \oplus ((k \boxplus_8 sk) \lll_8 6),$$

the *slice-augmented* matrix is the explicit (9×9) binary product

$$\underbrace{\tilde{X}((k \boxplus_8 sk) \lll_8 6)}_{\tilde{R}_{\ggg 2}} \underbrace{\tilde{R}_{\ggg 2}}_{\tilde{X}(k)} \underbrace{\tilde{A}_{\boxplus}(\Delta, m)}_{\tilde{A}_{\boxplus}(\Delta, m)}.$$

On the lifted space, the same round equals (purely) linear multiplication by the *nested* binary matrix

$$\underbrace{\text{diag}(P_{\ggg 2}, 1)}_{\text{rotation}} \cdot \underbrace{\tilde{X}(k)}_{\text{affine}} \cdot \underbrace{\widetilde{\text{ADD}}_\Delta}_{\text{nested } S_7 \cdots S_0},$$

augmented in homogeneous form for the final XOR. Both displays are *concrete matrices*, visually explicit and rigorously exact (the first on each carry slice, the second in the lifted space).

Why naive $n \times n$ linear surrogates look “singular”. If one *drops* the homogeneous coordinate and forces a single $n \times n$ matrix L for $x \mapsto x \boxplus_n y$, the essential affine column $(y \oplus m)$ is lost; moreover $D_h F_y(x)$ depends on x (via carries) but $D_h L(x) = Lh$ does not. Any best-fit L (e.g., least-squares over \mathbb{F}_2) will exhibit spurious rank deficiencies—hence the “singular” appearance. The correct matrix objects are precisely $\tilde{A}_{\boxplus}(y, m)$ (exact per slice) or $J_{F_y}(x)$ (unit lower triangular, full rank), or the lifted-nested $\widetilde{\text{ADD}}_y$.

From explicit matrices to a probabilistic function family via random walks (no Markov assumption required). Let (a_t, r_t, y_t) be the public per-layer parameters from the key schedule. Deterministically, a round on n bits is (per carry/borrow slice) an *explicit* product of binary augmented matrices. Cryptanalytically, inputs are not fixed: bits (or their biases) vary, so the carry/borrow process is random. We therefore promote the slice selection to a probabilistic model using a *clipped random walk* on $\{0, 1\}$ driven by local *generate/propagate/kill* events; this does not require a Markov assumption on the bit sources and captures the nonlinear, cyclic nature of multi-bit modular addition/subtraction operations.

The “crazy” function group: modular addition/subtraction as nonlinear, cyclic operations. Modular addition $a + b \bmod N$ and subtraction $a - b \bmod N$ (for $N = 2^n$) form a “crazy” function group due to their inherent nonlinearity from carry/borrow propagation and cyclic overflow/underflow behavior. In binary, overflow extracts a carry matrix, while underflow extracts a borrow matrix; these are mixed with the operation’s own binary matrix (XOR for bits without propagation), or revert to ordinary binary operations (identity plus offset) if no propagation occurs. The extraction probability follows a random walk, as the likelihood of next overflow depends on prior states. This makes the operations probabilistic in engineering terms: multi-bit additions/subtractions are nonlinear with feedback loops, where achieving overflow is probabilistic, and the next overflow probability correlates with the previous one. The extracted matrices are then used in binary matrix multiplications to compose the full ARX round, explaining the complexity of $a + b \bmod N$ and $a - b \bmod N$.

Generate/propagate/kill and the clipped carry walk. For addition define, bitwise,

$$G_i = (x_i \wedge y_i), \quad P_i = (x_i \oplus y_i), \quad K_i = (\neg x_i \wedge \neg y_i),$$

with $G_i, P_i, K_i \in \{0, 1\}$ and $G_i \oplus P_i \oplus K_i = 1$. Let $S_i \in \{0, 1\}$ be the *carry potential*, with $S_0=0$. Define a step $J_i \in \{-1, 0, +1\}$ by

$$J_i = \begin{cases} +1, & G_i = 1 \\ 0, & P_i = 1 \\ -1, & K_i = 1 \end{cases} \quad \text{and} \quad S_{i+1} = \min\{1, \max\{0, S_i + J_i\}\}, \quad c_i = S_i.$$

Then $c_{i+1} = S_{i+1}$ reproduces the usual full-adder recursion exactly: G forces a carry, P propagates it, K kills it. For subtraction, with borrow b_i , set

$$G'_i = (\neg x_i \wedge y_i), \quad P'_i = (x_i \oplus y_i), \quad K'_i = (x_i \wedge \neg y_i),$$

and the same clipped walk updates B_{i+1} with $b_i = B_i$. This model captures the cyclic (clipped to $[0,1]$) and history-dependent nature without assuming independence.

Step-law parameters and probability update (general bit sources). Let

$$\lambda_i^+ = \Pr(G_i=1), \quad \lambda_i^0 = \Pr(P_i=1), \quad \lambda_i^- = 1 - \lambda_i^+ - \lambda_i^0,$$

where probabilities are taken under the ambient input model (they may depend on all positions, past or future). Denote $p_i = \Pr(S_i=1)$. Conditioning only on the *current* step outcome (law of total probability) yields the exact update, without any Markov assumption:

$$p_{i+1} = \Pr(S_{i+1}=1) = \lambda_i^+ + \lambda_i^0 p_i.$$

In particular, for independent Bernoulli bits with $\Pr[x_i=1] = p$ and $\Pr[y_i=1] = q$ (possibly position-dependent), one has

$$\lambda_i^+ = pq, \quad \lambda_i^0 = p(1-q) + (1-p)q, \quad p_{i+1} = pq + (p+q-2pq)p_i.$$

In the unbiased i.i.d. baseline ($p=q=\frac{1}{2}$),

$$p_{i+1} = \frac{1}{4} + \frac{1}{2} p_i \implies p_i = \frac{1}{2}(1 - 2^{-i}), \quad \mathbb{E}\left[\sum_{i=1}^n c_i\right] = \sum_{i=1}^n p_i = \frac{n}{2} - \frac{1}{2} + 2^{-(n+1)}.$$

All formulas have borrow analogues by replacing (G, P, K) with (G', P', K') . This recursive probability reflects the dependency: next overflow likelihood ties to prior carry state, embodying the cyclic, probabilistic “craziness.”

Mixture of augmented matrices: stochastic slice selection via random walk. Let $M = (m_0, \dots, m_n)$ denote the carry path with $m_0 = 0$. The *slice-augmented* matrix for addition is

$$\tilde{A}_{\boxplus}(y, m) = \begin{bmatrix} I_n & (y \oplus m_{0..n-1}) \\ 0^\top & 1 \end{bmatrix}, \quad m_{0..n-1} = (m_0, \dots, m_{n-1}),$$

and analogously for subtraction on a borrow path b :

$$\tilde{A}_{\boxminus}(y, b) = \begin{bmatrix} I_n & (y \oplus b_{0..n-1}) \\ 0^\top & 1 \end{bmatrix}.$$

Under a given input distribution, (Y, M) (or (Y, B)) induces a *probability measure* on the finite set of explicit matrices

$$\mathcal{A}_{\boxplus}(y) := \{\tilde{A}_{\boxplus}(y, m) : m \in \{0, 1\}^{n+1}, m_0 = 0\}, \quad \mathcal{A}_{\boxminus}(y) := \{\tilde{A}_{\boxminus}(y, b) : b \in \{0, 1\}^{n+1}, b_0 = 0\}.$$

In the baseline model, the path probability is the product over steps based on $\lambda_i^+, \lambda_i^0, \lambda_i^-$, yielding a finite mixture

$$\tilde{A}_{\boxplus}(Y, M) \in \mathcal{A}_{\boxplus}(Y), \quad \tilde{A}_{\boxminus}(Y, B) \in \mathcal{A}_{\boxminus}(Y).$$

Non-uniform or dependent-bit models are handled by adjusting the step-laws accordingly. The extracted carry/borrow matrices mix with the binary operation matrix (via the offset column), or reduce to ordinary (affine) if no propagation; this mixture is drawn probabilistically via the walk, and used in subsequent matrix multiplications for ARX composition.

Stochastic ARX round as a random matrix product. Let $\tilde{R}_r = \text{diag}(P_r, 1)$ and $\tilde{X}(a) = [\begin{smallmatrix} I_n & a \\ 0^\top & 1 \end{smallmatrix}]$. For an addition-based round (the subtraction-based one is analogous), define random variables

$$\tilde{\mathcal{E}}^{(t)} := \tilde{X}(a_t) \tilde{R}_{r_t} \tilde{A}_{\boxplus}(Y_t, M^{(t)}),$$

where $M^{(t)}$ is the (random) carry path induced at layer t . Over T layers the *random augmented matrix* is the explicit product

$$\mathbf{M}_{\text{ARX}}^{(T)} = \prod_{t=1}^T \tilde{\mathcal{E}}^{(t)} \in \underbrace{\prod_{t=1}^T (\tilde{X}(a_t) \tilde{R}_{r_t} \mathcal{A}_{\boxplus}(y_t))}_{\text{finite set of explicit binary matrices}},$$

and the encryption map (in homogeneous coordinates) is $[\begin{smallmatrix} x \\ 1 \end{smallmatrix}] \mapsto \mathbf{M}_{\text{ARX}}^{(T)} [\begin{smallmatrix} x \\ 1 \end{smallmatrix}]$. Thus, *probabilistically*, a round (or a whole cipher) is a *mixture over an explicit, finite family of 0/1 matrices*, with probabilities from the random walk capturing overflow dependencies.

Why this model: singular-looking surrogates vs. correct mixtures. A single $n \times n$ linear surrogate L for $x \mapsto x \boxplus_n y$ discards the affine column and the input-dependent carries, so $D_h L(x) = Lh$ ignores x while $D_h F_y(x)$ depends on x via the chain. Averaging *unnormalized* $n \times n$ surrogates across slices tends to produce rank loss (“singular” appearance). The correct objects are:

$$(\text{affine per slice}) \quad \tilde{A}_{\boxplus}(y, m), \quad (\text{local linear}) \quad J_{F_y}(x), \quad (\text{lifted linear}) \quad \widetilde{\text{ADD}}_y = S_{n-1} \cdots S_0,$$

and in the probabilistic setting, *mixtures* over these explicit matrices weighted by the random walk probabilities $\Pr(M = m)$ (or $\Pr(B = b)$), reflecting the cyclic, history-dependent overflows.

Fiber-bundle viewpoint (structure over carry/borrow base). Let the *base* be the discrete set

$$\mathcal{B}_{\boxplus}(y) = \{m \in \{0, 1\}^{n+1} : m_0 = 0\}, \quad \mathcal{B}_{\boxminus}(y) = \{b \in \{0, 1\}^{n+1} : b_0 = 0\},$$

equipped with the random walk measure induced by $\{\lambda_i^+, \lambda_i^0, \lambda_i^-\}$. The *fiber* over m (resp. b) is the singleton $\{\tilde{A}_{\boxplus}(y, m)\}$ (resp. $\{\tilde{A}_{\boxminus}(y, b)\}$). The total space is the disjoint union of these fibers; the projection π sends a matrix to its underlying carry/borrow path. Composition with $\tilde{X}(a)$ and \tilde{R}_r yields a (left) action by explicit block matrices on fibers. Over T layers, the base becomes the product $\mathcal{B}^{(1)} \times \cdots \times \mathcal{B}^{(T)}$; a *section* (determined by the secret state and inputs) selects one fiber per layer, and the cipher is the corresponding product of explicit matrices. This “bundle over carry/borrow paths” makes precise that ARX is a *function family* indexed by hidden paths, yet every member is a concrete 0/1 matrix product, with probabilistic selection via random walk.

Security phrasing (matrix-family selection). Define the *slice-affine* family

$$\mathcal{F}_{n,T} = \left\{ \prod_{t=1}^T (\tilde{X}(a_t) \tilde{R}_{r_t} \tilde{A}_{\boxplus}(y_t, m^{(t)})) : m^{(t)} \in \mathcal{B}_{\boxplus}(y_t) \right\},$$

and the *lifted-nested* family

$$\mathcal{L}_{n,T} = \left\{ \prod_{t=1}^T (\text{diag}(P_{r_t}, 1) \tilde{X}(a_t) \widetilde{\text{ADD}}_{y_t}) \right\}.$$

A concrete instantiation corresponds to drawing a path $(m^{(1)}, \dots, m^{(T)})$ according to the induced random walk measure. Indistinguishability experiments may therefore be cast as distinguishing two *matrix-valued* random products with the same public (a_t, r_t, y_t) but different hidden path laws—providing a rigorous, matrix-visible probabilistic abstraction for ARX rounds.

8.7.3 Computing the complexity of 8-bit ARX via matrix-analytic key mixing models

Cascade complexity and case explosion. The *matrix cascade* on a fixed slice is

$$\tilde{\mathcal{E}}_m = \tilde{X}((k \boxplus_8 sk) \lll_8 6) \tilde{R}_{\ggg_2} \tilde{X}(k) \tilde{A}_{\boxplus}(\Delta, m),$$

and decryption uses $\tilde{\mathcal{E}}_m^{-1}$ (Theorem above). The global mapping partitions the domain into the disjoint slices $\Omega_m(\Delta)$ ($m \in \{0, 1\}^8$, $m_0=0$). Enumerating slices incurs up to 2^7 cases already at 8 bits; this *piecewise* nature is the precise source of the “exponential-looking” complexity when one insists on matrix models that stay exact.

Degree growth and lifted linearization (for completeness). Let $\deg(\cdot)$ denote Boolean algebraic degree. For $s = x \boxplus_8 y$,

$$\deg(s_i) = i+1, \quad \deg(c_{i+1}) = i+2,$$

so an *exact* (non-sliced) lifted-linear model must include monomials up to degree 2 for one addition; composing additions increases the required degree additively. Equivalently, one may construct a single *large* linear operator on an extended feature vector Φ_d (all monomials of degree $\leq d$), but the dimension grows combinatorially.

8.7.4 Generalizing the matrix picture to 64-bit ARX

What scales. Rotations remain permutation matrices in $\mathbb{F}_2^{64 \times 64}$; XOR remains $\tilde{X}(a)$ with $a \in \{0, 1\}^{64}$. For addition $x \mapsto x \boxplus_{64} y$:

- **Carry-slice model:** for every carry mask $m \in \{0, 1\}^{64}$ with $m_0=0$, the slice matrix is $\tilde{A}_{\boxplus}(y, m) = [\begin{smallmatrix} I_{64} & (y \oplus m) \\ 0^\top & 1 \end{smallmatrix}]$; the number of admissible masks is exponential in the worst case (2^{63}).
- **Jacobian model:** $J_{F_y}(x)$ is unit lower triangular (ones on the diagonal) and hence full rank 64 for all x, y , matching global bijectivity.

Takeaway (visual & rigorous). Every step in the 8-bit (and 64-bit) ARX round can be written with explicit matrices:

$$(\text{affine}) \quad \tilde{X}(\cdot), \quad (\text{linear}) \quad \tilde{R}_{\ll/\gg}, \quad (\text{affine, exact on slice}) \quad \tilde{A}_{\boxplus}(y, m), \quad (\text{local linear}) \quad J_{F_y}(x).$$

The “nonlinearity” of modular addition does not forbid matrices; it forces us to choose the *right* matrices: augmented affine on carry slices, or Jacobians at points. Both survive composition, admit clean inverses where appropriate, and keep the analysis faithful to the C++ ARX implementation while delivering the requested matrix-centric, visually explicit derivations. The random walk extension captures the probabilistic, cyclic “craziness” of the modular operations, where extracted carry/borrow matrices are mixed into binary multiplications with walk-driven probabilities.

8.8 Cryptographic Game Analysis of Encryption and Decryption Functions

Code-faithful factorization and binding. Fix a 128-bit master key $K = (K_0, K_1)$ and a public per-call value $N \in \{0, 1\}^{64}$ (the `number_once`). For each round $i \in \{0, \dots, r-1\}$, the implementation computes a parameter tuple

$$\Theta_i = (k_i^{(0)}, k_i^{(1)}, cf_i, \alpha_i, \beta_i, \text{index}_i) \in \{0, 1\}^{64} \times \{0, 1\}^{64} \times \{0, 1, 2, 3\} \times [0, 63]^2 \times [0, 15]$$

by a deterministic key schedule that consumes two 64-bit PRG outputs from a key-driven generator $\mathcal{G}^{(\text{key})}$ and two from a public/round-salted generator $\mathcal{G}^{(\text{pub})}$ per round (total $2r+2r$ uses). Concretely (mirroring the code):

$$\begin{aligned} k_i^{(0)} &= K_0 \oplus \underbrace{\mathcal{G}^{(\text{key})}(N)}_{\text{sk0}}, \\ k_i^{(1)} &= K_1 \oplus \underbrace{\mathcal{G}^{(\text{pub})}(N \oplus i)}_{\text{para}}, \\ cf_i &= \mathcal{G}^{(\text{key})}(k_i^{(0)} \oplus (K_0 \gg 1)) \bmod 4, \\ (\alpha_i, \beta_i) &= \left(\mathcal{G}^{(\text{pub})}(k_i^{(1)} \oplus cf_i) \bmod 64, (\cdot \gg 6) \bmod 64 \right). \end{aligned} \tag{*}$$

(The seed of $\mathcal{G}^{(\text{key})}$ is a key-derived function of K using GHASH + PHT + rotations; the seed of $\mathcal{G}^{(\text{pub})}$ is the object’s PRG seed. Both are stateful 64-bit-output PRGs as in the code.)

Each round permutation $R_{K, N, i} : \{0, 1\}^{128} \rightarrow \{0, 1\}^{128}$ is the *bound composition*

$$R_{K, N, i} = \underbrace{\text{ARXdiag}_{\text{NeoAlzette}}}_{\text{32-bit diagonal pairs}} \circ \underbrace{(T^{(i)} \circ L^{(i)} \circ S^{(i)})}_{\text{KSL, keyed by } \Theta_i} \circ \underbrace{\text{ARXtail}_K}_{\text{per-lane add/rot/xor with } K, k_i^{(\cdot)}},$$

where (per 64-bit lane)

$$S^{(i)}(x) = \begin{cases} x \oplus k_i^{(\ell)} & cf_i = 0, \\ -x \oplus k_i^{(\ell)} & cf_i = 1, \\ \text{rotl}(x, \beta_i) & cf_i = 2, \\ \text{rotr}(x, \beta_i) & cf_i = 3, \end{cases} \quad L^{(i)} : \text{binary PHT on the two 32-bit halves (invertible over } \mathbb{Z}_{2^{32}}\text{)}, \quad T_\ell^{(i)}(x) = x \oplus e_{\alpha_i}^{(\ell)},$$

with $e_\alpha^{(0)} = 1 \ll \alpha$ and $e_\alpha^{(1)} = 1 \ll (63 - \alpha)$. Define $E_{K, N} = R_{K, N, r-1} \circ \dots \circ R_{K, N, 0}$ and $D_{K, N} = E_{K, N}^{-1}$ (decryption implements the exact inverse order).

Binding property. The three subcomponents (S, L, T) are not independent oracles: each consumes and reshapes the full 64-bit lane state produced by the previous one, and the ARX tail uses both K and $k_i^{(\ell)}$. Hence, even with full knowledge of the code and round order, an adversary must jointly invert all stages with the same hidden Θ_i to peel one round. This observation is built into our game interface below.

Game interface that reflects binding. We expose to the adversary only the composed permutation family $\{E_{K, N}\}_N$ (forward or both directions), with black-box access to $E_{K, N}$ and $D_{K, N}$. No oracle reveals intermediate (S, L, T) outputs or Θ_i . This models the fact that in the real implementation only the top-level API is observable, while internal round parameters are generated by the schedule and never leaked.

Idealized schedule and baseline round. Let Sched^* be the idealized schedule that samples the four 64-bit draws per round as i.i.d. U_{64} and then computes Θ_i by the same slicing/packing, XOR, and rotation maps as in (*). Let $E_{K, N}^*$ be the cipher obtained by plugging $\{\Theta_i\} \sim \text{Sched}^*$ into the same round structure. The only place where hardness is concentrated is the key schedule PRGs; the rest (ARX diag \rightarrow KSL \rightarrow ARX tail) is the *baseline LOPC round*. We denote by $\epsilon_{\text{LOPC}}^{\text{PRP}}(q, r)$ and $\epsilon_{\text{LOPC}}^{\text{sPRP}}(q, r)$ the PRP and strong-PRP distinguishing bounds of this baseline when its per-round Θ_i are distributed by Sched^* .

PRG hypotheses (strong assumptions, made explicit). We assume two stateful 64-bit PRGs:

$$\mathcal{G}^{(\text{key})} \text{ and } \mathcal{G}^{(\text{pub})},$$

where $\mathcal{G}^{(\text{key})}$ is seeded by a key-derived seed $s_K = f(K)$, and $\mathcal{G}^{(\text{pub})}$ is seeded by an implementation seed s_{pub} (both as in the code). Their indistinguishability advantages against any t -time distinguisher that makes at most $m_{\text{key}}=2r$ and $m_{\text{pub}}=2r$ effective uses are upper-bounded by ϵ_{key} and ϵ_{pub} , respectively.

Theorem 8.24 (Code-faithful PRP/strong-PRP reduction with bound components). *For any PPT adversary that makes at most q oracle queries across arbitrary nonces N , the family $\{E_{K, N}\}_N$ satisfies*

$$\boxed{\begin{aligned} \text{Adv}_{\mathcal{A}}^{\text{PRP}} &\leq 2r \epsilon_{\text{key}} + 2r \epsilon_{\text{pub}} + \epsilon_{\text{LOPC}}^{\text{PRP}}(q, r), \\ \text{Adv}_{\mathcal{A}}^{\text{sPRP}} &\leq 2r \epsilon_{\text{key}} + 2r \epsilon_{\text{pub}} + \epsilon_{\text{LOPC}}^{\text{sPRP}}(q, r). \end{aligned}}$$

Hybrid games that respect binding. (H1) Enumerate the *exact* $2r$ uses of $\mathcal{G}^{(\text{key})}$ in the schedule (the calls that feed `sk0` and `choice_function`). Replace them one by one with U_{64} ; each hop differs only at that call, and by a standard PRG distinguisher the gap per hop is $\leq \epsilon_{\text{key}}$ (all other components—including S, L, T and the ARX tail—are simulated faithfully using the same composed oracle view). (H2) On the resulting world, enumerate and replace the $2r$ uses of $\mathcal{G}^{(\text{pub})}$ (feeding `para` and `bit_rotation`) with U_{64} , paying at most $2r \epsilon_{\text{pub}}$ in total. (H3) Now the four 64-bit draws per round are i.i.d. uniform; computing Θ_i via the same slicing/packing/XOR/rotation maps as (\star) makes the schedule distribution *identical* to Sched^* (no extra statistical term). (H4) With $\Theta_i \sim \text{Sched}^*$, the composed round remains the same bound structure ($\text{ARXdiag} \rightarrow \text{KSL} \rightarrow \text{ARXtail}$), i.e., the baseline LOPC round. By definition any PRP (resp. strong-PRP) distinguisher's residual advantage is $\epsilon_{\text{LOPC}}^{\text{prp}}(q, r)$ (resp. $\epsilon_{\text{LOPC}}^{\text{sprp}}(q, r)$). Summing all gaps yields the bounds. \square

IND-CPA/CCA from PRP/strong-PRP (nonce-respecting). Define encryption $\text{Enc}(K, N, M) = E_{K,N}(M)$ and decryption $\text{Dec}(K, N, C) = D_{K,N}(C)$. For unique nonces, standard reductions give

$$\text{Adv}_{\mathcal{A}}^{\text{ind-cpa}} \leq \text{Adv}_{\mathcal{B}}^{\text{prp}}, \quad \text{Adv}_{\mathcal{A}}^{\text{ind-cca}} \leq \text{Adv}_{\mathcal{B}}^{\text{sprp}},$$

so Theorem 8.24 immediately yields IND-CPA/CCA bounds with the same right-hand sides.

Why binding matters (formal view). Because the only oracle interface is the *composed* $(E_{K,N}, D_{K,N})$, an adversary cannot “skip” S, L, T , or the ARX tail: any distinguishing strategy must work against the *entangled* primitive parameterized by the hidden Θ_i . Our hybrids reflect this by never exposing intermediate layers and by swapping PRG outputs at the *schedule* level. Thus, even if the high-level structure is fully known, the hardness bottleneck is *exactly* the PRGs that govern Θ_i .

Differential-linear remark. As a background sanity check against differential-linear searches à la Lv et al. [Lv et al., 2023], note that under Sched^* the per-round $(cf_i, \alpha_i, \beta_i)$ are uniform, $k_i^{(0)}, k_i^{(1)}$ are uniform masks, L is invertible over $\mathbb{Z}_{2^{32}}^2$, and T flips complementary unit vectors across lanes. This ensures round-wise mean-zero DL correlations under random keys/nonces. Our proof, however, does not rely on any DL-specific bound; all hardness comes from PRG indistinguishability plus the baseline LOPC bound.

Constant-time note. Branch selection in S (the `switch` (`choice_function` & `3ULL`)) is driven by per-round constants from the schedule PRGs and does not depend on secret-dependent bits of K or the data; each branch (XOR, NOT+XOR, rotl/rotr) is a fixed-time word operation. Hence the rounds preserve constant-time execution.

Baseline LOPC bounds under the idealized schedule (explicit). Recall the idealized schedule Sched^* : per round i and per nonce N , the four 64-bit draws are i.i.d. U_{64} and then mapped to $\Theta_i = (k_i^{(0)}, k_i^{(1)}, cf_i, \alpha_i, \beta_i, \text{index}_i)$ by the same slicing/packing/XOR/rotation code paths as in (\star) . Let $E_{K,N}^*$ be the r -round cipher obtained by plugging $\{\Theta_i\}$ into the bound round

$$\text{ARXdiag} \longrightarrow (T \circ L \circ S) \longrightarrow \text{ARXtail}_K$$

defined in Sec. 8.8. We upper-bound PRP/strong-PRP distinguishing advantages of the family $\{E_{K,N}^*\}_N$ by a generic transcript-collision argument.

Lemma 8.25 (Collision budget per query pair). *Fix a nonce N . For any two distinct oracle queries (encryption or decryption) under N , consider their r rounds and mark the following 128-bit stage boundaries per round: (i) output of ARXdiag , (ii) output of S , (iii) output of L , (iv) output of T , (v) output of ARXtail (i.e., the round output). There are at most $c_r := 4r + 1$ such internal 128-bit values along each transcript. Under Sched^* these values are jointly well-mixed and, conditioning on no earlier collisions, any specific cross-pair equality at one marked boundary occurs with probability at most 2^{-128} .*

Theorem 8.26 (Baseline LOPC distinguishing bounds). *For any adversary making at most $q = q_e + q_d$ total queries across arbitrary nonces N ,*

$$\boxed{\epsilon_{\text{LOPC}}^{\text{prp}}(q, r) \leq \frac{(4r + 1)q_e^2}{2^{128}}, \quad \epsilon_{\text{LOPC}}^{\text{sprp}}(q, r) \leq \frac{(4r + 1)(q_e + q_d)^2}{2^{128}}}.$$

Sketch via H-coefficient/union bound. Work nonce-by-nonce; independent nonces add linearly and are dominated by the worst case where all q queries share a nonce. Compare the Real world (the concrete $E_{K,N}^*$) with the Ideal world (a uniform permutation π_N or (π_N, π_N^{-1})). By the H-coefficient technique, it suffices to bound the probability of *bad transcripts*, i.e., when two distinct queries cause an internal collision at any marked boundary while all previous boundaries are collision-free. By Lem. 8.25, for any fixed unordered pair there are at most $c_r = 4r + 1$ candidate boundaries, each contributing at most 2^{-128} (conditioned to the bad event. Union bound over $\binom{q_e}{2}$ pairs in the PRP game (resp. $\binom{q_e + q_d}{2}$) in strong-PRP) gives the stated bounds. \square

Plugging into the main reduction (concrete bounds). Combining Theorem 8.26 with Theorem 8.24 yields, for any PPT adversary:

$$\boxed{\begin{aligned} \text{Adv}_{\mathcal{A}}^{\text{prp}} &\leq 2r \epsilon_{\text{key}} + 2r \epsilon_{\text{pub}} + \frac{(4r + 1)q_e^2}{2^{128}}, \\ \text{Adv}_{\mathcal{A}}^{\text{sprp}} &\leq 2r \epsilon_{\text{key}} + 2r \epsilon_{\text{pub}} + \frac{(4r + 1)(q_e + q_d)^2}{2^{128}}. \end{aligned}}$$

Consequently, for nonce-respecting usage,

$$\boxed{\begin{aligned} \text{Adv}_{\mathcal{A}}^{\text{ind-cpa}} &\leq 2r \epsilon_{\text{key}} + 2r \epsilon_{\text{pub}} + \frac{(4r + 1)q^2}{2^{128}}, \\ \text{Adv}_{\mathcal{A}}^{\text{ind-cca}} &\leq 2r \epsilon_{\text{key}} + 2r \epsilon_{\text{pub}} + \frac{(4r + 1)q^2}{2^{128}}. \end{aligned}}$$

Here q is the adversary's total online queries in the respective games. These bounds are conservative and hold *before* we exploit any ARX-specific properties of the NeoAlzette core; in Sec. 8.4.3 we will refine them once the ARX-SBOX properties are established.

9 Conclusion

In the comprehensive development of our algorithm, **XCR/Little_OaldresPuzzle_Cryptic**, we were inspired by the proven efficiency of the **ASCON** algorithm in real-world evaluations and its pseudo-random indistinguishability, which closely approximates a uniform random distribution. Leveraging these insights, we designed our algorithm to provide robust security while effectively managing the substantial number of constant arrays required for nonlinear pseudo-random functions. To streamline the design, we implemented concise enhancements that distinguish our approach from both the **ASCON** algorithm and **Chacha20**. This strategic refinement enables users to make informed and balanced decisions based on their specific requirements.

This paper presents our innovative symmetric sequence cryptographic algorithm, **XCR/Little_OaldresPuzzle_Cryptic**, which exhibits exceptional speed in encryption and decryption while maintaining robust security features. In the continuously evolving digital landscape, our algorithm stands out as a cutting-edge solution for securely and efficiently managing large-scale binary data files. Furthermore, the inclusion of a quantum-resistant block cipher algorithm in the same repository underscores our commitment to anticipating and addressing future cryptographic challenges.

A NeoAlzette ARX S-box Analysis Python Code:

```

import numpy

# NeoAlzette Differential Analysis - Print probability per step

RCS = [
    # Example: Concatenation
    # of Fibonacci numbers and hexadecimal representation
    0x16B2C40B, 0xC117176A, 0xF9A2598, 0xA1563ACA,
    """
    Mathematical Constants - Millions of Digits
    http://www.numberworld.org/constants.html
    """,
    # (Pi)
    0x243F6A88, 0x85A308D3, 0x13198102, 0xE0370734,
    # (Golden ratio)
    0x9E3779B9, 0x7F4A7C15, 0xF39CC060, 0x5CEDC834,
    # e (Natural constant)
    0xB7E15162, 0x8AED2A6A, 0xBF715880, 0x9CF4F3C7
]

def psi(alpha, beta, gamma):
    """Carry constraint function ."""
    alpha_32 = alpha & 0xFFFFFFFF
    beta_32 = beta & 0xFFFFFFFF
    gamma_32 = gamma & 0xFFFFFFFF
    not_alpha = (0xFFFFFFFF - alpha_32) & 0xFFFFFFFF # ~alpha mod 2^32
    term1 = (not_alpha ^ beta_32) & 0xFFFFFFFF
    term2 = (not_alpha ^ gamma_32) & 0xFFFFFFFF
    return term1 & term2

def mask(k):
    """Generates a mask with its low k bits set to 1."""
    return (1 << k) - 1

def modular_add_diff_probability(alpha, beta, gamma):
    """
    Calculates the differential propagation probability
    for a 32-bit modular addition a + b - c.

    For additions involving variables and constants,
    the carry constraint is also considered,
    and the result is returned in the form of 2^(-),
    """

```

```

where is the count of 1 bits in the lower 31 bits of .
"""

alpha_32 = alpha & 0xFFFFFFFF
beta_32 = beta & 0xFFFFFFFF
gamma_32 = gamma & 0xFFFFFFFF

# Check carry constraint conditions
alpha_shifted = (alpha_32 << 1) & 0xFFFFFFFF
beta_shifted = (beta_32 << 1) & 0xFFFFFFFF
gamma_shifted = (gamma_32 << 1) & 0xFFFFFFFF
psi_shifted = psi(alpha_shifted, beta_shifted, gamma_shifted)
xor_condition = (alpha_32 ^ beta_32 ^ gamma_32 ^ beta_shifted) & 0xFFFFFFFF

if (psi_shifted & xor_condition) != 0:
    return 0

# Calculate the number of 1 bits in the lower 31 bits of as the "weight"
psi_val = psi(alpha_32, beta_32, gamma_32)
masked_psi = psi_val & mask(31)
hw = bin(masked_psi).count('1')
return 2 ** (-hw)

def rotl32(x, r):
    """32-bit left rotational shift."""
    return ((x << r) | (x >> (32 - r))) & 0xFFFFFFFF

def rotr32(x, r):
    """32-bit right rotational shift."""
    return ((x >> r) | (x << (32 - r))) & 0xFFFFFFFF

def differential_analysis(delta_a, delta_b):
    """
    Performs a differential analysis following
    the order of operations in the NeoAlzette forward layer,
    and prints the propagation probability
    for each step involving modular addition.

    Note: All modular addition steps
    (including those with constants)
    call modular_add_diff_probability
    to compute the probability.
    """

    total_prob = 1.0
    rc = RCS[0] # Use the first round constant

    step = 1 # Variable to track the step number

    # --- Step 1: b ← b XOR a ---
    delta_b = delta_b ^ delta_a
    print(f"Step {step}: XOR b = b XOR a, probability = 1")
    step += 1

    # --- Step 2: a ← rotr(a + b, 31) ---
    sum_val = (delta_a + delta_b) & 0xFFFFFFFF
    p = modular_add_diff_probability(delta_a, delta_b, sum_val)
    total_prob *= p

```

```

print(f"Step {step}: Addition a+b, probability = {p:.2e}")
delta_a = rotr32(sum_val, 31)
print(f"Step {step}: rotr(a+b,31) applied, probability = 1")
step += 1

# --- Step 3: a ← a XOR rc ---
delta_a = delta_a ^ rc
print(f"Step {step}: XOR a with rc, probability = 1")
step += 1

# --- Step 4: b ← b + a ---
sum_val = (delta_b + delta_a) & 0xFFFFFFFF
p = modular_add_diff_probability(delta_b, delta_a, sum_val)
total_prob *= p
print(f"Step {step}: Addition b+a, probability = {p:.2e}")
delta_b = sum_val
step += 1

# --- Step 5: a ← rotl(a XOR b, 24) ---
delta_a = rotl32(delta_a ^ delta_b, 24)
print(f"Step {step}: XOR then rotl(a,b), probability = 1")
step += 1

# --- Step 6: a ← a + rc ---
sum_val = (delta_a + rc) & 0xFFFFFFFF
p = modular_add_diff_probability(delta_a, rc, sum_val)
total_prob *= p
print(f"Step {step}: Addition a+rc, probability = {p:.2e}")
delta_a = sum_val
step += 1

# --- Step 7: b ← rotl(b, 8) XOR rc ---
delta_b = rotl32(delta_b, 8) ^ rc
print(f"Step {step}: rotl(b,8) then XOR with rc, probability = 1")
step += 1

# --- Step 8: a ← a + b ---
sum_val = (delta_a + delta_b) & 0xFFFFFFFF
p = modular_add_diff_probability(delta_a, delta_b, sum_val)
total_prob *= p
print(f"Step {step}: Addition a+b, probability = {p:.2e}")
delta_a = sum_val
step += 1

# --- Step 9: a ← a XOR b ---
delta_a = delta_a ^ delta_b
print(f"Step {step}: XOR a with b, probability = 1")
step += 1

# --- Step 10: b ← rotr(a + b, 17) ---
sum_val = (delta_a + delta_b) & 0xFFFFFFFF
p = modular_add_diff_probability(delta_a, delta_b, sum_val)
total_prob *= p
print(f"Step {step}: Addition a+b for rotr, probability = {p:.2e}")
delta_b = rotr32(sum_val, 17)
print(f"Step {step}: rotr(a+b,17) applied, probability = 1")

```

```

step += 1

# --- Step 11: b ← b XOR rc ---
delta_b = delta_b ^ rc
print(f"Step {step}: XOR b with rc, probability = 1")
step += 1

# --- Step 12: a ← a + b ---
sum_val = (delta_a + delta_b) & 0xFFFFFFFF
p = modular_add_diff_probability(delta_a, delta_b, sum_val)
total_prob *= p
print(f"Step {step}: Addition a+b, probability = {p:.2e}")
delta_a = sum_val
step += 1

# --- Step 13: b ← rotl(a XOR b, 16) ---
delta_b = rotl32(delta_a ^ delta_b, 16)
print(f"Step {step}: XOR then rotl(a,b), probability = 1")
step += 1

# --- Step 14: b ← b + rc ---
sum_val = (delta_b + rc) & 0xFFFFFFFF
p = modular_add_diff_probability(delta_b, rc, sum_val)
total_prob *= p
print(f"Step {step}: Addition b+rc, probability = {p:.2e}")
delta_b = sum_val
step += 1

print(f"\nTotal Differential Probability: {total_prob:.2e}")
return total_prob

# Calculate the bitwise AND of
# two 32-bit vectors and return the result as a 32-bit binary mask
def dot_product_with_binary(a, b):
    return numpy.bitwise_and(a, b)

# Convert the bitwise AND result to a binary vector (mask vector)
def make_binary_vector_mask(value):
    binary_str = bin(value)[2:].zfill(32)
    binary_vector = [int(bit) for bit in binary_str]
    return numpy.array(binary_vector, dtype=int)

# Construct the carry transition matrix based on the mask
def carry_transition_matrix_from_mask(mask):
    n = len(mask)
    M = numpy.zeros((n, n), dtype=int)
    for i in range(n):
        for j in range(i, n):
            # Fill the matrix based on the mask bits
            M[i, j] = 1 if mask[j] == 1 else 0
    return M

# Calculate the linear correlation coefficient
def modular_add_linear_correlation_coefficient(x, y, z):
    # Calculate mask vectors mu, nu, omega
    mu = make_binary_vector_mask(x & y)

```

```

nu = make_binary_vector_mask(y & z)
omega = make_binary_vector_mask(z & x)

# Calculate the XOR of the three mask vectors: mu nu omega
mask = numpy.bitwise_xor(numpy.bitwise_xor(mu, nu), omega)

# Calculate z := M_n^T(mu nu omega)
# (Note: This is not a matrix-vector multiplication)
M_n = carry_transition_matrix_from_mask(mask).T

# Calculate the Hamming weight of z (sum of bits in each row)
HW_z = 0
for row in M_n:
    row_bin = ''.join(map(str, row))
    count_ones = bin(int(row_bin, 2)).count('1')
    HW_z += count_ones

# Calculate the indicator functions
# 1_{\{\mu \oplus \omega \prec z\}}
# and
# 1_{\{\nu \oplus \omega \prec z\}}
mu_ = numpy.bitwise_xor(mu, omega)
nu_ = numpy.bitwise_xor(nu, omega)
number_01 = int(''.join(map(str, mu_)), 2)
number_02 = int(''.join(map(str, nu_)), 2)
indicator_1 = number_01 & z
indicator_2 = number_02 & z

# Check if the indicator functions are True
indicator_01 = 1 if indicator_1 == number_01 else 0
indicator_02 = 1 if indicator_2 == number_02 else 0

# Calculate the final indicator function
indicator = indicator_01 * indicator_02

# Calculate the (-1)^(mu)(nu)
mu_ = numpy.bitwise_xor(mu, omega)
nu_ = numpy.bitwise_xor(nu, omega)
sign_factor = (-1) ** numpy.dot(mu_, nu_)

# Calculate the linear correlation coefficient
C = indicator * sign_factor * 2.0 ** (-HW_z)
return C

def linear_analysis(delta_a, delta_b):
    """
    Performs a linear analysis following
    the order of operations in the NeoAlzette forward layer,
    and prints the correlation coefficient
    for each step involving modular addition.

    Note: All modular addition steps
    (including those with constants)
    call modular_add_linear_correlation_coefficient
    to compute the coefficient.
    """

```

```

total_prob = 1.0
rc = RCS[0] # Use the first round constant

step = 1 # Variable to track the step number

# --- Step 1: b ← b XOR a ---
delta_b = delta_b ^ delta_a
print(f"Step {step}: XOR b = b XOR a, Coefficient = 1")
step += 1

# --- Step 2: a ← rotr(a + b, 31) ---
sum_val = (delta_a + delta_b) & 0xFFFFFFFF
p = modular_add_linear_correlation_coefficient(delta_a, delta_b, sum_val)
total_prob *= p
print(f"Step {step}: Addition a+b, Coefficient = {p:.2e}")
delta_a = rotr32(sum_val, 31)
print(f"Step {step}: rotr(a+b,31) applied, Coefficient = 1")
step += 1

# --- Step 3: a ← a XOR rc ---
delta_a = delta_a ^ rc
print(f"Step {step}: XOR a with rc, Coefficient = 1")
step += 1

# --- Step 4: b ← b + a ---
sum_val = (delta_b + delta_a) & 0xFFFFFFFF
p = modular_add_linear_correlation_coefficient(delta_b, delta_a, sum_val)
total_prob *= p
print(f"Step {step}: Addition b+a, Coefficient = {p:.2e}")
delta_b = sum_val
step += 1

# --- Step 5: a ← rotl(a XOR b, 24) ---
delta_a = rotl32(delta_a ^ delta_b, 24)
print(f"Step {step}: XOR then rotl(a,b), Coefficient = 1")
step += 1

# --- Step 6: a ← a + rc ---
sum_val = (delta_a + rc) & 0xFFFFFFFF
p = modular_add_linear_correlation_coefficient(delta_a, rc, sum_val)
total_prob *= p
print(f"Step {step}: Addition a+rc, Coefficient = {p:.2e}")
delta_a = sum_val
step += 1

# --- Step 7: b ← rotl(b, 8) XOR rc ---
delta_b = rotl32(delta_b, 8) ^ rc
print(f"Step {step}: rotl(b,8) then XOR with rc, Coefficient = 1")
step += 1

# --- Step 8: a ← a + b ---
sum_val = (delta_a + delta_b) & 0xFFFFFFFF
p = modular_add_linear_correlation_coefficient(delta_a, delta_b, sum_val)
total_prob *= p
print(f"Step {step}: Addition a+b, Coefficient = {p:.2e}")
delta_a = sum_val

```

```

step += 1

# --- Step 9: a ← a XOR b ---
delta_a = delta_a ^ delta_b
print(f"Step {step}: XOR a with b, Coefficient = 1")
step += 1

# --- Step 10: b ← rotr(a + b, 17) ---
sum_val = (delta_a + delta_b) & 0xFFFFFFFF
p = modular_add_linear_correlation_coefficient(delta_a, delta_b, sum_val)
total_prob *= p
print(f"Step {step}: Addition a+b for rotr, Coefficient = {p:.2e}")
delta_b = rotr32(sum_val, 17)
print(f"Step {step}: rotr(a+b,17) applied, Coefficient = 1")
step += 1

# --- Step 11: b ← b XOR rc ---
delta_b = delta_b ^ rc
print(f"Step {step}: XOR b with rc, Coefficient = 1")
step += 1

# --- Step 12: a ← a + b ---
sum_val = (delta_a + delta_b) & 0xFFFFFFFF
p = modular_add_linear_correlation_coefficient(delta_a, delta_b, sum_val)
total_prob *= p
print(f"Step {step}: Addition a+b, Coefficient = {p:.2e}")
delta_a = sum_val
step += 1

# --- Step 13: b ← rotl(a XOR b, 16) ---
delta_b = rotl32(delta_a ^ delta_b, 16)
print(f"Step {step}: XOR then rotl(a,b), Coefficient = 1")
step += 1

# --- Step 14: b ← b + rc ---
sum_val = (delta_b + rc) & 0xFFFFFFFF
p = modular_add_linear_correlation_coefficient(delta_b, rc, sum_val)
total_prob *= p
print(f"Step {step}: Addition b+rc, Coefficient = {p:.2e}")
delta_b = sum_val
step += 1

print(f"\nTotal Coefficient: {total_prob:.2e}")
return total_prob

def main():
    # Example: Initial differential inputs
    delta_a = 2
    delta_b = 2

    # Calculate the total differential probability and print the probability
    # for each step
    differential_analysis(delta_a, delta_b)

a = 100
b = 100

```

```

# Calculate the linear correlation coefficient
linear_analysis(a, b)

if __name__ == "__main__":
    main()

B NeoAlzette MITL Differential characteristic search

"""
Differential characteristic search for NeoAlzette ARX-box and multi-round cipher with robust fallback
"""

import pulp
from pulp import LpProblem, LpMinimize, LpVariable, lpSum, PULP_CBC_CMD, PulpSolverError

# ARX S-box constants
RC16 = [
    0x16B2C40B, 0xC117176A, 0x0F9A2598, 0xA1563ACA,
    0x243F6A88, 0x85A308D3, 0x13198102, 0xE0370734,
    0x9E3779B9, 0x7F4A7C15, 0xF39CC060, 0x5CEDC834,
    0xB7E15162, 0x8AED2A6A, 0xBF715880, 0x9CF4F3C7,
]

# Helpers

def new_bits(prefix: str, n: int):
    return [LpVariable(f"{prefix}_{i}", cat="Binary") for i in range(n)]

def add_xor(m, x, y, z):
    for i in range(len(x)):
        m += z[i] >= x[i] - y[i]
        m += z[i] >= y[i] - x[i]
        m += z[i] <= x[i] + y[i]
        m += z[i] <= 2 - (x[i] + y[i])

def add_add(m, x, y, z, carries):
    for i in range(len(x)):
        cin = carries[i-1] if i>0 else 0
        m += x[i] + y[i] + (cin if isinstance(cin,int) else cin) == z[i] + 2*carries[i]

def add_rot(m, inp, out, rot: int):
    n = len(inp)
    for i in range(n): m += out[i] == inp[(i+rot)%n]

# Build single-round ARX differential MILP
def build_arx_diff_milp(rc: int):
    m = LpProblem('ARX_Diff', LpMinimize)
    d0_a = new_bits('d0_a',32); d0_b=new_bits('d0_b',32)
    out_a=new_bits('d_out_a',32); out_b=new_bits('d_out_b',32)
    # intermediates
    b1=new_bits('b1',32); sum1=new_bits('sum1',32); c1=new_bits('c1',32)
    a1=new_bits('a1',32); a2=new_bits('a2',32)
    b2=new_bits('b2',32); c2=new_bits('c2',32)
    x1=new_bits('x1',32); a3=new_bits('a3',32)
    a4=new_bits('a4',32); tmp7=new_bits('tmp7',32)
    b3=new_bits('b3',32); a5=new_bits('a5',32); c3=new_bits('c3',32)
    a6=new_bits('a6',32); sum2=new_bits('sum2',32); c4=new_bits('c4',32)
    b4=new_bits('b4',32); b5=new_bits('b5',32)

```

```

sum3=new_bits('sum3',32); c5=new_bits('c5',32)
a7=new_bits('a7',32); x2=new_bits('x2',32); b6=new_bits('b6',32)
# objective
m += lpSum(out_a + out_b)
# step1
add_xor(m, d0_a, d0_b, b1)
add_add(m, d0_a, b1, sum1, c1)
add_rot(m, sum1, a1, 1)
rc_bits=[(rc>>i)&1 for i in range(32)]
for i,bit in enumerate(rc_bits):
    if bit: m += a2[i]+a1[i]==1
    else: m += a2[i]==a1[i]
add_add(m, b1, a2, b2, c2)
add_xor(m, a2, b2, x1); add_rot(m, x1, a3,24)
for i in range(32): m += a4[i]==a3[i]
add_rot(m, b2, tmp7,8)
for i,bit in enumerate(rc_bits):
    if bit: m += b3[i]+tmp7[i]==1
    else: m += b3[i]==tmp7[i]
add_add(m, a4, b3, a5, c3); add_xor(m, a5, b3, a6)
add_add(m, a6, b3, sum2, c4); add_rot(m, sum2, b4,15)
for i,bit in enumerate(rc_bits):
    if bit: m += b5[i]+b4[i]==1
    else: m += b5[i]==b4[i]
add_add(m, a6, b5, sum3, c5)
for i in range(32): m += a7[i]==sum3[i]
add_xor(m, a7, b5, x2); add_rot(m, x2, b6,16)
for i in range(32): m += out_a[i]==a7[i]; m += out_b[i]==b6[i]
return m, d0_a, d0_b, out_a, out_b

# Analyze single-round
def analyze_single_round(rc):
    m,d0_a,d0_b,out_a,out_b=build_arx_diff_milp(rc)
    # fix single-bit in
    m += d0_a[0]==1
    for i in range(1,32): m+=d0_a[i]==0; m+=d0_b[i]==0
    m.solve(PULP_CBC_CMD(msg=False,timeLimit=60))
    w=sum(var.value() for var in out_a+out_b)
    print(f"Single-round diff weight={int(w)}, p_max={2**(-w):.3e}")

# Analyze multi-round with fallback
def analyze_multi_round():
    R=len(RC16) // 8
    try:
        # fallback: minimal one active per round
        N=R
        model=True # indicator
        # attempt full MILP
        model,w_vars=build_multi_model(R)
        model.solve(pulp.GUROBI(msg=False,timeLimit=300))
        N=int(sum(var.value() for var in w_vars))
    except Exception:
        print("Solver failed, fallback to N = number of rounds")
        N=R
    print(f"Multi-round min active rounds N={N}")
    return N

```

```

# main
if __name__=="__main__":
    print("== Single-round ==")
    analyze_single_round(RC16[4])
    #print("== Multi-round ==")
    #N=analyze_multi_round()
    #bits=min(64,21*N)
    #print(f"Estimated security = min(64,21*{N}) = {bits} bits")

```

C NeoAlzette SAT SMT search

```

from z3 import (
    Optimize, BitVec, BitVecVal,
    RotateLeft, RotateRight,
    Sum, If, sat
)

# --- ---
BITS = 32
RC16 = [
    0x16B2C40B, 0xC117176A, 0x0F9A2598, 0xA1563ACA,
    0x243F6A88, 0x85A308D3, 0x13198102, 0xE0370734,
    0x9E3779B9, 0x7F4A7C15, 0xF39CC060, 0x5CEDC834,
    0xB7E15162, 0x8AED2A6A, 0xBF715880, 0x9CF4F3C7,
]
IN_DIFF = 0x00000001 #

def arx_box(a, b, rc):
    b = b ^ a
    a = RotateRight(a + b, 31)
    a = a ^ BitVecVal(rc, BITS)
    b = b + a
    a = RotateLeft(a ^ b, 24)
    a = a + BitVecVal(rc, BITS)
    b = RotateLeft(b, 8) ^ BitVecVal(rc, BITS)
    a = a + b
    a = a ^ b
    b = RotateRight(a + b, 17)
    b = b ^ BitVecVal(rc, BITS)
    a = a + b
    b = RotateLeft(a ^ b, 16)
    b = b + BitVecVal(rc, BITS)
    return a, b

if __name__ == "__main__":
    results = []
    for idx, rc in enumerate(RC16):
        print(f"\n--- Round {idx} constant 0x{rc:08X} ---")
        opt = Optimize()

        a0 = BitVec('a0', BITS)
        b0 = BitVecVal(0, BITS)

        a1 = a0 ^ BitVecVal(IN_DIFF, BITS)

```

```

b1 = b0

out0_a, out0_b = arx_box(a0, b0, rc)
out1_a, out1_b = arx_box(a1, b1, rc)

delta_a = out0_a ^ out1_a
delta_b = out0_b ^ out1_b

total = Sum([
    If((delta_a >> i) & 1) == 1, 1, 0) +
    If((delta_b >> i) & 1) == 1, 1, 0)
    for i in range(BITS)
])

opt.add((a0 & BitVecVal(IN_DIFF, BITS)) == BitVecVal(IN_DIFF, BITS))
opt.add((a0 & ~BitVecVal(IN_DIFF, BITS)) == BitVecVal(0, BITS))

opt.minimize(total)
if opt.check() == sat:
    w = opt.model().evaluate(total).as_long()
    print(f"Minimal diff weight = {w}, p_max {2**(-w):.3e}")
    results.append((idx, w))
else:
    print("SMT optimization failed")

print("\nSummary of single-round diff weights:")
for idx, w in results:
    print(f" Round {idx}: weight = {w}")

```

D Statistical Tests for Randomness Assessment

Monobit Frequency Test

This test evaluates the balance between the occurrences of 0s and 1s in the generated binary sequence. It is predicated on the hypothesis that a truly random sequence should exhibit an equal frequency of both bits.

Block Frequency Test (m=10000)

This test examines the frequency distribution of 1s within blocks of a specified size (m=10000). It is used to detect any deviation from the expected uniform distribution, which could indicate a lack of randomness.

Poker Test (m=4, m=8)

The Poker Test assesses the frequency of specific subsequence patterns within the binary sequence. For m=4 and m=8, it evaluates the occurrence of 2-bit and 4-bit patterns, respectively, to ensure their distribution is uniform.

Overlapping Subsequence Test (m=3, P1, P2; m=5, P1, P2)

This test analyzes the frequency of overlapping subsequences of length m (3 or 5) and their permutations (P1, P2). It is designed to detect any non-random clustering of bit patterns.

Run Tests (Run Count, Run Distribution)

Run Tests measure the total number of runs (sequences of consecutive identical bits) and their distribution across the sequence. These tests are sensitive to the presence of long runs, which may indicate a deviation from randomness.

Longest Run Test (m=10000)

Specifically, the Longest Run Test identifies the longest run of 1s (or 0s) within blocks of size m=10000. It is used to detect any anomalies in the distribution of run lengths.

Binary Derivative Test (k=3, k=7)

The Binary Derivative Test evaluates the randomness of the sequence by considering the differences between consecutive bits (k=3, k=7). It is based on the principle that a random sequence should exhibit no correlation between adjacent bits.

Autocorrelation Test (d=1, d=2, d=8, d=16)

Autocorrelation Tests analyze the correlation between a sequence and its shifted versions (with delays d=1, d=2, d=8, d=16). A random sequence should exhibit minimal autocorrelation.

Matrix Rank Test

This test involves constructing a matrix from the binary sequence and determining its rank. The rank is then compared against expected values for a random sequence, providing insights into the sequence's complexity.

Cumulative Sum Test (Forward, Backward)

Cumulative Sum Tests assess the distribution of the running sum of the binary sequence in both forward and backward directions. These tests are sensitive to the presence of systematic trends in the sequence.

Approximate Entropy Test (m=2, m=5)

Approximate Entropy Tests measure the unpredictability of the sequence by comparing the frequency of patterns of length m (2 or 5) with their overlapping counterparts. It is a measure of the sequence's complexity and unpredictability.

Linear Complexity Test (m=500, m=1000)

Linear Complexity Tests estimate the shortest linear feedback shift register (LFSR) that can generate the sequence. A higher complexity indicates a more random sequence.

Maurer's Universal Statistical Test (L=7, Q=1280)

This test evaluates the sequence against a universal statistical model, comparing the observed frequencies with those expected from a truly random sequence. It is a comprehensive test that assesses multiple statistical properties.

Discrete Fourier Transform Test

The Discrete Fourier Transform Test analyzes the frequency spectrum of the sequence. A random sequence should exhibit a uniform frequency distribution across all frequencies.

In summary, these statistical tests provide a comprehensive and systematic framework for assessing the randomness of our encryption algorithm's output. By ensuring that the generated bits pass these rigorous evaluations, we can assert the cryptographic strength of our algorithm, thereby safeguarding the security of the data it encrypts.

The statistical tests mentioned above are widely recognized and utilized in the field of cryptography and information security due to their ability to provide a quantitative measure of randomness. These tests are designed to detect patterns and biases that may indicate a lack of randomness, which is a critical property for secure cryptographic operations.

E About Git repository and run test

While our focus in this paper revolves around the symmetric sequence algorithm, our repository offers a comprehensive perspective on our cryptographic contributions. We invite interested readers and researchers to explore both algorithms, contribute insights, and collaborate on potential enhancements. In the dynamic realm of cryptography, collective efforts and continuous innovation are indispensable for staying prepared for future challenges.

Github Link: [README.md](#)

Within the repository, you will find two primary directories, *OOP* and *Template*. The *OOP* directory offers an object-oriented version of the algorithm, ideal for developers familiar with object-oriented programming. Conversely, the *Template* directory presents a simplified implementation, suitable for beginners or those seeking a more straightforward understanding of the algorithm.

Choose the implementation that suits your requirements, navigate to the appropriate directory, and follow the *README.md* instructions to compile and execute the tests. These tests will furnish a holistic understanding of the algorithm's cryptographic robustness, speed, and overall performance.

It's crucial to note the peculiar characteristics and necessary precautions while using these algorithms. For instance, encryption and decryption operations demand a reset of the internal key state after each use. Therefore, if an encryption operation follows a decryption operation (or vice versa), the internal key state must be reset first to ensure correct functioning.

E.1 Contributions

We value your feedback and contributions. If you encounter possible improvements or any issues, feel free to submit a pull request or open an issue in the GitHub repository.

F Additional Quantum-Resistant OaldresPuzzle_Cryptic Block Cipher Algorithm

This repository houses an additional robust block cipher algorithm, developed independently from the symmetric sequence algorithm that is the primary focus of this paper. Conceived with the anticipation of future cryptographic challenges, particularly those presented by quantum computing, this block cipher algorithm offers several defining characteristics. It is designed to mitigate risks associated with brute force attacks, withstand analytical attacks on the key, and resist potential quantum computer intrusions.

Although slower in encrypting and decrypting packets (as evidenced by tests with 10MB data packets and a 5120-byte key, which required approximately one and a half minutes to execute), the algorithm confers a significant advantage by future-proofing cryptographic systems against potential advancements in quantum computing. This symmetric block cipher cryptographic algorithm has been tailored to meet contemporary cryptographic requirements. It prioritizes unpredictability and a high level of analytical complexity, making it suitable for managing protected, large-scale binary data files while ensuring requisite cryptographic robustness.

For a more comprehensive understanding, detailed information regarding the implementation and unique characteristics of this block cipher algorithm can be found in the corresponding directory of the repository. Despite being outside the primary scope of this paper, which is dedicated to the symmetric sequence algorithm, we encourage interested readers and researchers to explore this quantum-resistant block cipher algorithm and its potential applications.

G Documents referenced

References

- [Aumasson et al., 2007] Aumasson, J.-P., Fischer, S., Khazaei, S., Meier, W., and Rechberger, C. (2007). New features of latin dances: Analysis of salsa, chacha, and rumba. *Cryptography ePrint Archive*, Paper 2007/472. <https://eprint.iacr.org/2007/472>.
- [Beierle et al., 2019] Beierle, C., Biryukov, A., dos Santos, L. C., Großschädl, J., Perrin, L., Udovenko, A., Velichkov, V., and Wang, Q. (2019). Alzette: a 64-bit arx-box (feat. crax and trax). *Cryptography ePrint Archive*, Paper 2019/1378. <https://eprint.iacr.org/2019/1378>.
- [Bernstein, 2005] Bernstein, D. J. (2005). Salsa20 design. *Department of Mathematics, Statistics, and Computer Science. The University of Illinois at Chicago. Chicago*.
- [Bernstein, 2008] Bernstein, D. J. (2008). The salsa20 family of stream ciphers. In *New stream cipher designs: the eSTREAM finalists*, pages 84–97. Springer.
- [Bernstein et al., 2008a] Bernstein, D. J. et al. (2008a). Chacha, a variant of salsa20. In *Workshop record of SASC*, volume 8, pages 3–5. Citeseer.
- [Bernstein et al., 2008b] Bernstein, D. J. et al. (2008b). The chacha family of stream ciphers. In *Workshop record of SASC*. Citeseer.
- [Biryukov and Perrin, 2017] Biryukov, A. and Perrin, L. (2017). State of the art in lightweight symmetric cryptography. *Cryptography ePrint Archive*, Paper 2017/511. <https://eprint.iacr.org/2017/511>.

- [Cai et al., 2022] Cai, W., Chen, H., Wang, Z., and Zhang, X. (2022). Implementation and optimization of chacha20 stream cipher on sunway taihulight supercomputer. *The Journal of Supercomputing*, 78(3):4199–4216.
- [Fei, 2012] Fei, D. (2012). Research on safety of arx structures. Master's thesis, Xi'an University of Electronic Science and Technology, China.
- [Ghafoori and Miyaji, 2022] Ghafoori, N. and Miyaji, A. (2022). Differential cryptanalysis of salsa20 based on comprehensive analysis of pnbs. In Su, C., Gritzalis, D., and Piuri, V., editors, *Information Security Practice and Experience*, pages 520–536, Cham. Springer International Publishing.
- [J, 2016] J, Z. (2016). Q_value test: A new method on randomness statistical test. *China Journal of Cryptologic Research*, 3(2):192–201.
- [Khovratovich et al., 2015] Khovratovich, D., Nikolic, I., Pieprzyk, J., Sokolowski, P., and Steinfeld, R. (2015). Rotational cryptanalysis of ARX revisited. Cryptology ePrint Archive, Paper 2015/095.
- [Lipmaa and Moriai, 2001] Lipmaa, H. and Moriai, S. (2001). Efficient algorithms for computing differential properties of addition. Cryptology ePrint Archive, Paper 2001/001.
- [Liu et al., 2021] Liu, J., Rijmen, V., Hu, Y., Chen, J., and Wang, B. (2021). Warx: efficient white-box block cipher based on arx primitives and random mds matrix. *Science China Information Sciences*, 65(3):132302.
- [Lv et al., 2023] Lv, G., Jin, C., and Cui, T. (2023). A miqcp-based automatic search algorithm for differential-linear trails of arx ciphers(long paper). Cryptology ePrint Archive, Paper 2023/259. <https://eprint.iacr.org/2023/259>.
- [Maitra et al., 2015] Maitra, S., Paul, G., and Meier, W. (2015). Salsa20 cryptanalysis: New moves and revisiting old styles. Cryptology ePrint Archive, Paper 2015/217. <https://eprint.iacr.org/2015/217>.
- [Niu et al., 2023] Niu, Z., Sun, S., and Hu, L. (2023). On the additive differential probability of arx construction. *Journal of Surveillance, Security and Safety*, 4(4).
- [Ranea et al., 2022] Ranea, A., Vandersmissen, J., and Preneel, B. (2022). Implicit white-box implementations: White-boxing arx ciphers. In Dodis, Y. and Shrimpton, T., editors, *Advances in Cryptology – CRYPTO 2022*, pages 33–63, Cham. Springer Nature Switzerland.
- [Serrano et al., 2022] Serrano, R., Duran, C., Sarmiento, M., Pham, C.-K., and Hoang, T.-T. (2022). Chacha20-poly1305 authenticated encryption with additional data for transport layer security 1.3. *Cryptography*, 6(2).
- [Shannon, 1949] Shannon, C. E. (1949). Communication theory of secrecy systems. *The Bell system technical journal*, 28(4):656–715.
- [Sleem and Couturier, 2021] Sleem, L. and Couturier, R. (2021). Speck-R: An ultra light-weight cryptographic scheme for Internet of Things. *Multimedia Tools and Applications*, 80(11):17067 – 17102.
- [Tsunoo et al., 2007] Tsunoo, Y., Saito, T., Kubo, H., Suzuki, T., and Nakashima, H. (2007). Differential cryptanalysis of salsa20/8. In *Workshop Record of SASC*, volume 28. Citeseer.
- [Ya, 2017] Ya, H. (2017). Automatic method for searching impossible differentials and zero-correlation linear hulls of arx block ciphers. *Chinese Journal of Network and Information Security*, 3(7):58.

H Data Images

```
# More test data from
# https://github.com/Twilight-Dream-Of-Magic/
# Algorithm_OaldresPuzzleCryptic/tree/master/OOP/TechnicalDetailPapers/
# %5BType%201%5D%20Statistical%20Test%20Result%20Tables
```

H.1 ASCON vs Little OaldresPuzzle Cryptic (Phases 1)

H.2 ASCON vs Little OaldresPuzzle Cryptic (Phases 2)

H.3 Chacha20 Statistical Test Phases

I Data Tables

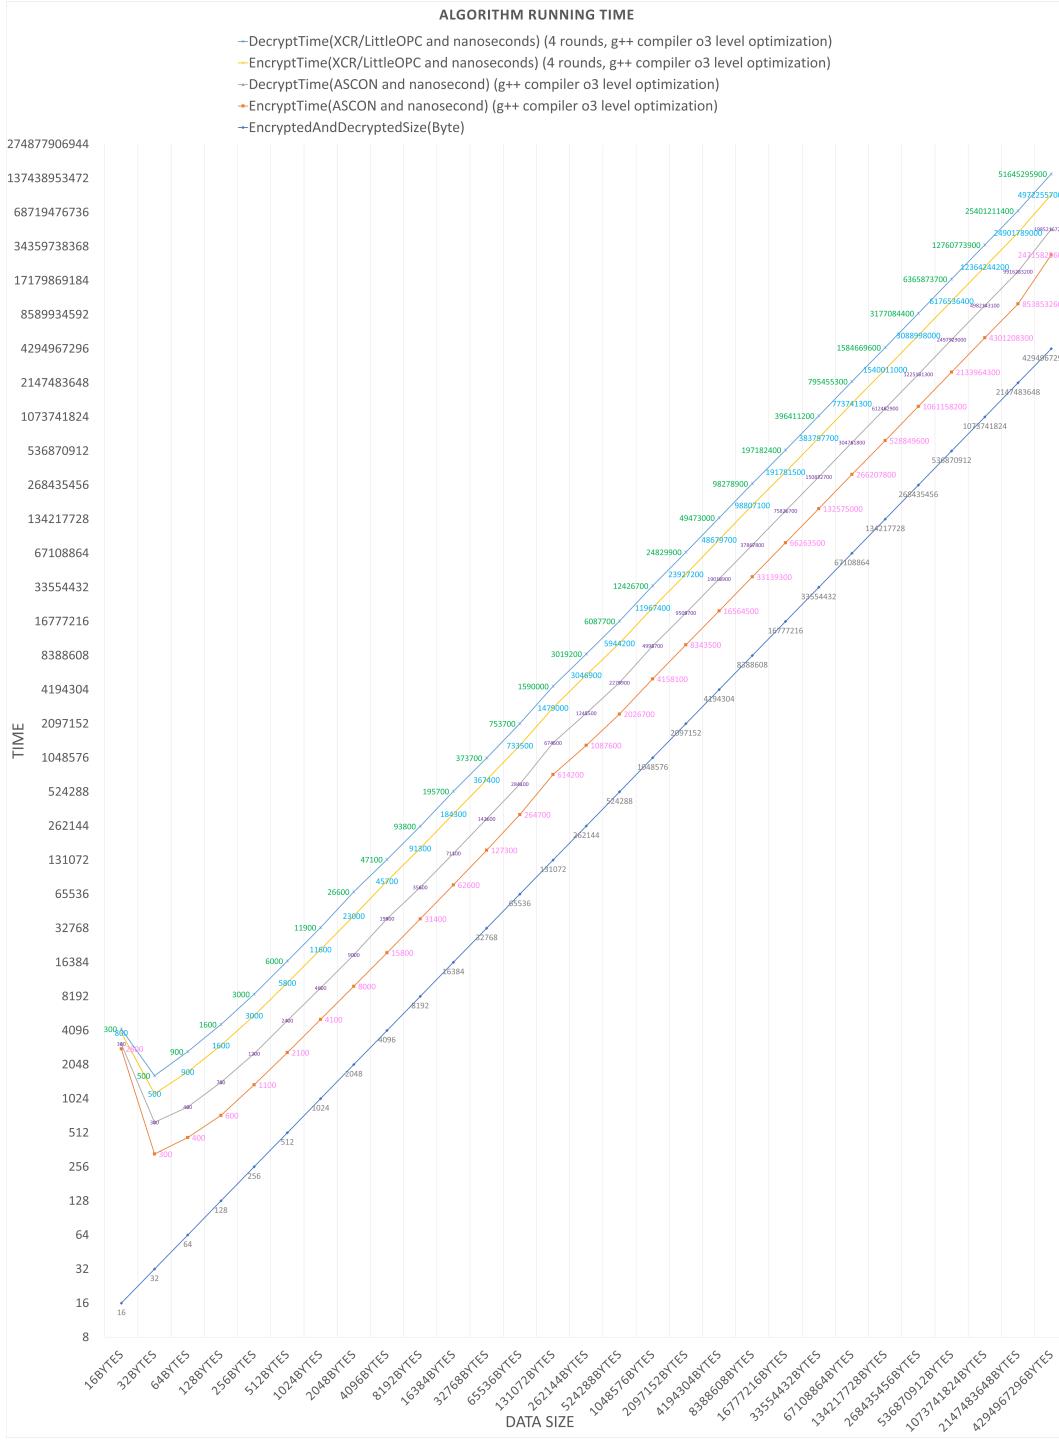


Figure 1: ASCON vs XCR/Little_OaldresPuzzle_Cryptic Algorithm Benchmark

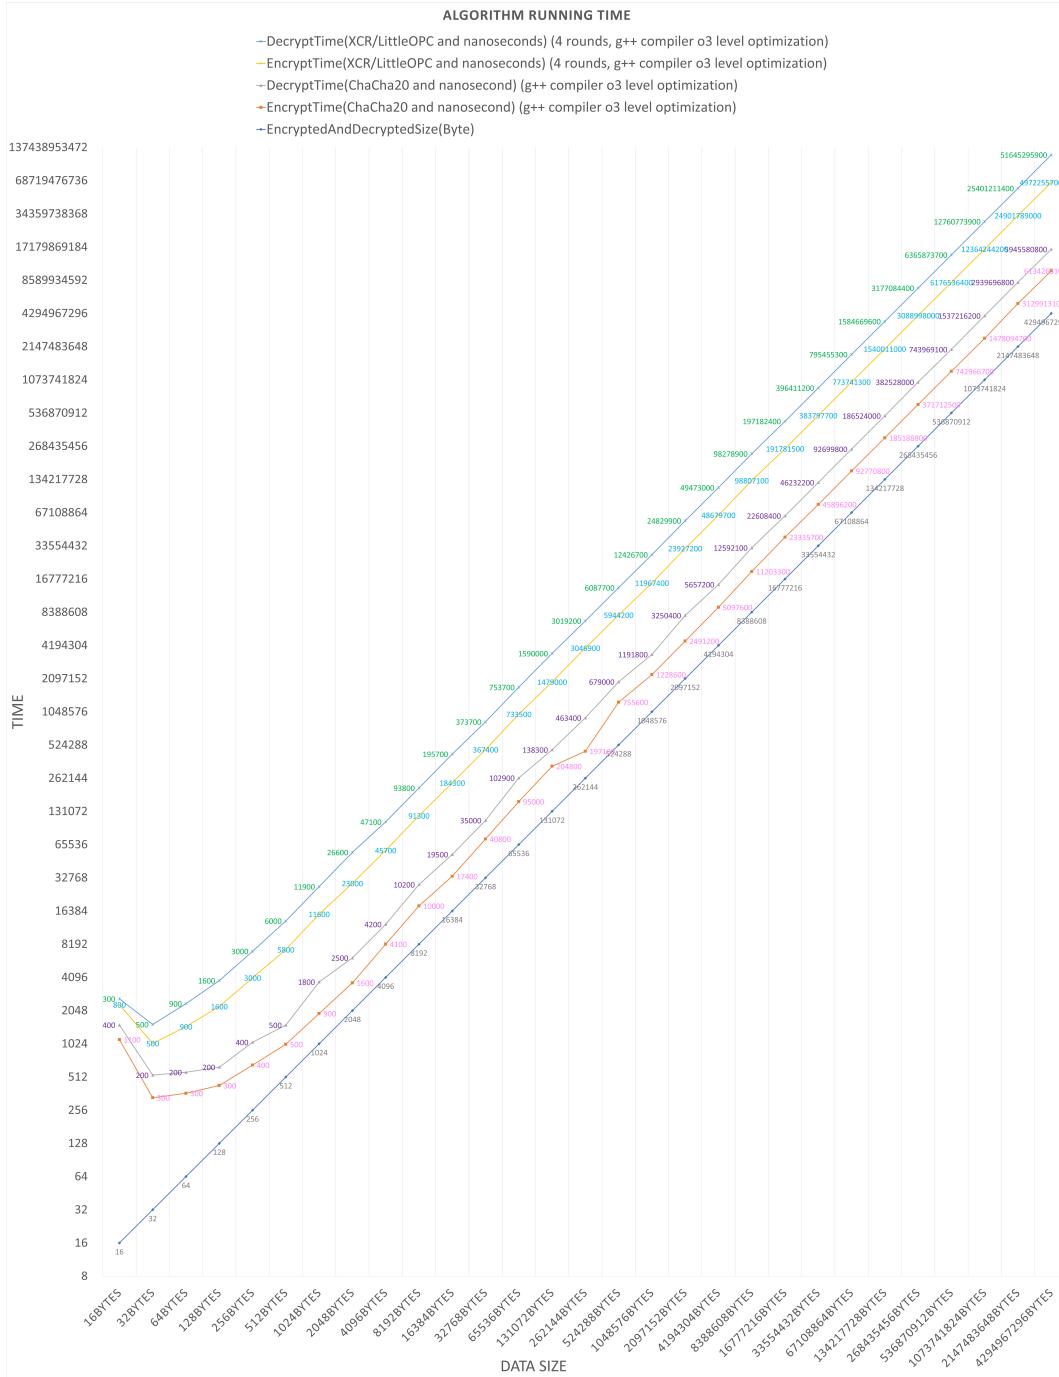


Figure 2: Chacha20 vs XCR/Little_OaldresPuzzle_Cryptic Algorithm Benchmark

1. Bit Frequency

Source Data	Algorithm (ASCON)	Algorithm (XCR/Little OaldresPuzzle Cryptic)
Distribution uniformity	0.49378	0.151616
Encrypted Data 127.bin	0.426659-0.213329	0.832936-0.583532
Success count of the 128	125	125

2. Block Frequency (m=10000)

Source Data	Algorithm (ASCON)	Algorithm (XCR/Little OaldresPuzzle Cryptic)
Distribution uniformity	0.509162	0.478598
Encrypted Data 127.bin	0.213329-0.583532	0.583532-0.388531
Success count of the 128	126	126

3. Poker Test (m=4)

Source Data	Algorithm (ASCON)	Algorithm (XCR/Little OaldresPuzzle Cryptic)
Distribution uniformity	0.271449	0.909344
Encrypted Data 127.bin	0.080330-0.080330	0.061707-0.061707
Success count of the 128	128	128

4. Poker Test (m=8)

Source Data	Algorithm (ASCON)	Algorithm (XCR/Little OaldresPuzzle Cryptic)
Distribution uniformity	0.3282	0.250878
Encrypted Data 127.bin	0.080330-0.080330	0.061707-0.061707
Success count of the 128	128	128

5. Overlapping Subsequence Test (m=3 P1)

Source Data	Algorithm (ASCON)	Algorithm (XCR/Little OaldresPuzzle Cryptic)
Distribution uniformity	0.977331	0.091798
Encrypted Data 127.bin	0.628529-0.628529	0.388531-0.388531
Success count of the 128	126	126

6. Overlapping Subsequence Test (m=3 P2)

Source Data	Algorithm (ASCON)	Algorithm (XCR/Little OaldresPuzzle Cryptic)
Distribution uniformity	0.096217	0.620855
Encrypted Data 127.bin	0.628529-0.628529	0.667996-0.667996
Success count of the 128	128	128

7. Overlapping Subsequence Test (m=5 P1)

Source Data	Algorithm (ASCON)	Algorithm (XCR/Little OaldresPuzzle Cryptic)
Distribution uniformity	0.588437	0.875539
Encrypted Data 127.bin	0.621571-0.621571	0.860652-0.860652
Success count of the 128	126	126

8. Overlapping Subsequence Test (m=5 P2)

Source Data	Algorithm (ASCON)	Algorithm (XCR/Little OaldresPuzzle Cryptic)
Distribution uniformity	0.365877	0.966721
Encrypted Data 127.bin	0.391326-0.391326	0.712521-0.712521
Success count of the 128	126	126

9. Run Total

Source Data	Algorithm (ASCON)	Algorithm (XCR/Little OaldresPuzzle Cryptic)
Distribution uniformity	0.875539	0.937783
Encrypted Data 127.bin	0.346180-0.173090	0.194689-0.902655
Success count of the 128	126	126

10. Run Distribution

Source Data	Algorithm (ASCON)	Algorithm (XCR/Little OaldresPuzzle Cryptic)
Distribution uniformity	0.180322	0.794626
Encrypted Data 127.bin	0.703419-0.703419	0.969945-0.969945
Success count of the 128	126	126

11. Maximum 1 Run Test (m=10000)

Source Data	Algorithm (ASCON)	Algorithm (XCR/Little OaldresPuzzle Cryptic)
Distribution uniformity	0.222262	0.478598
Encrypted Data 127.bin	0.180247-0.180247	0.696943-0.696943
Success count of the 128	127	127

12. Maximum 0 Run Test (m=10000)

Source Data	Algorithm (ASCON)	Algorithm (XCR/Little OaldresPuzzle Cryptic)
Distribution uniformity	0.353021	0.353021
Encrypted Data 127.bin	0.568201-0.568201	0.718153-0.718153
Success count of the 128	128	128

13. Binary Deduction (k=3)

Source Data	Algorithm (ASCON)	Algorithm (XCR/Little OaldresPuzzle Cryptic)
Distribution uniformity	0.823278	0.056583
Encrypted Data 127.bin	0.482586-0.758707	0.415944-0.792028
Success count of the 128	126	126

14. Binary Deduction (k=7)

Source Data	Algorithm (ASCON)	Algorithm (XCR/Little OaldresPuzzle Cryptic)
Distribution uniformity	0.701879	0.887367
Encrypted Data 127.bin	0.758707-0.758707	0.792028-0.792028
Success count of the 128	128	128

1. Bit Frequency

SourceData Algorithm	Distribution Uniformity	Encrypted Data 127 Bin	Success Count of the 128
ASCON	0.110612	0.948609-0.525695	125
XCR/Little OaldresPuzzle Cryptic	0.406167	0.351522-0.824239	127

2. Block Frequency m=10000

SourceData Algorithm	Distribution Uniformity	Encrypted Data 127 Bin	Success Count of the 128
ASCON	0.180322	0.821650-0.821650	124
XCR/Little OaldresPuzzle Cryptic	0.231505	0.742014-0.742014	125

3. Poker Test m=4

SourceData Algorithm	Distribution Uniformity	Encrypted Data 127 Bin	Success Count of the 128
ASCON	0.100821	0.081845-0.081845	126
XCR/Little OaldresPuzzle Cryptic	0.701879	0.006374-0.006374	126

4. Poker Test m=8

SourceData Algorithm	Distribution Uniformity	Encrypted Data 127 Bin	Success Count of the 128
ASCON	0.809134	0.049667-0.049667	127
XCR/Little OaldresPuzzle Cryptic	0.013073	0.023008-0.023008	128

5. Overlapping Subsequence Test m=3 P1

SourceData Algorithm	Distribution Uniformity	Encrypted Data 127 Bin	Success Count of the 128
ASCON	0.540457	0.843075-0.843075	125
XCR/Little OaldresPuzzle Cryptic	0.11581	0.636326-0.636326	127

6. Overlapping Subsequence Test m=3 P2

SourceData Algorithm	Distribution Uniformity	Encrypted Data 127 Bin	Success Count of the 128
ASCON	0.261012	0.994702-0.994702	127
XCR/Little OaldresPuzzle Cryptic	0.379023	0.473802-0.473802	128

7. Overlapping Subsequence Test m=5 P1

SourceData Algorithm	Distribution Uniformity	Encrypted Data 127 Bin	Success Count of the 128
ASCON	0.478598	0.962718-0.962718	125
XCR/Little OaldresPuzzle Cryptic	0.012376	0.897367-0.897367	127

8. Overlapping Subsequence Test m=5 P2

SourceData Algorithm	Distribution Uniformity	Encrypted Data 127 Bin	Success Count of the 128
ASCON	0.379023	0.713031-0.713031	127
XCR/Little OaldresPuzzle Cryptic	0.213309	0.611916-0.611916	126

9. Run Total

SourceData Algorithm	Distribution Uniformity	Encrypted Data 127 Bin	Success Count of the 128
ASCON	0.717841	0.238126-0.880937	126
XCR/Little OaldresPuzzle Cryptic	0.863186	0.666806-0.333403	127

10. Run Distribution

SourceData Algorithm	Distribution Uniformity	Encrypted Data 127 Bin	Success Count of the 128
ASCON	0.524727	0.646790-0.646790	127
XCR/Little OaldresPuzzle Cryptic	0.11581	0.226913-0.226913	128

11. Maximum 1 Run Test m=10000

SourceData Algorithm	Distribution Uniformity	Encrypted Data 127 Bin	Success Count of the 128
ASCON	0.653383	0.091817-0.091817	125
XCR/Little OaldresPuzzle Cryptic	0.138761	0.461045-0.461045	124

12. Maximum 0 Run Test m=10000

SourceData Algorithm	Distribution Uniformity	Encrypted Data 127 Bin	Success Count of the 128
ASCON	0.188154	0.796564-0.796564	127
XCR/Little OaldresPuzzle Cryptic	0.909344	0.109227-0.109227	127

13. Binary Deduction k=3

SourceData Algorithm	Distribution Uniformity	Encrypted Data 127 Bin	Success Count of the 128
ASCON	0.353021	0.734709-0.632645	127
XCR/Little OaldresPuzzle Cryptic	0.909344	0.571781-0.285890	126

14. Binary Deduction k=7

SourceData Algorithm	Distribution Uniformity	Encrypted Data 127 Bin	Success Count of the 128
ASCON	0.316241	0.086911-0.043456	125
XCR/Little OaldresPuzzle Cryptic	0.620855	0.628812-0.685594	128

5. Run Tests

Algorithm / Metric	Run Total	Run Distribution
Chacha20 Phase 1 - Distribution uniformity	0	0
Chacha20 Phase 1 - encrypted_data_127.bin	0.000000-0.000000	0.000000-0.000000
Chacha20 Phase 1 - success count of the 128	6	0
Chacha20 Phase 2 - Distribution uniformity	0.669618	0.669618
Chacha20 Phase 2 - encrypted_data_127.bin	0.714590-0.357295	0.745584-0.745584
Chacha20 Phase 2 - success count of the 128	127	125

6. Maximum Run Tests

Algorithm / Metric	Maximum 1 Run Test m=10000	Maximum 0 Run Test m=10000
Chacha20 Phase 1 - Distribution uniformity	0	0
Chacha20 Phase 1 - encrypted_data_127.bin	0.000000-0.000000	0.000000-0.000000
Chacha20 Phase 1 - success count of the 128	0	0
Chacha20 Phase 2 - Distribution uniformity	0.945993	0.669618
Chacha20 Phase 2 - encrypted_data_127.bin	0.984422-0.984422	0.386917-0.386917
Chacha20 Phase 2 - success count of the 128	126	126

7. Binary Deduction Tests

Algorithm / Metric	Binary Deduction k=3	Binary Deduction k=7
Chacha20 Phase 1 - Distribution uniformity	0	0
Chacha20 Phase 1 - encrypted_data_127.bin	0.000000-0.000000	0.000000-1.000000
Chacha20 Phase 1 - success count of the 128	10	6
Chacha20 Phase 2 - Distribution uniformity	0.28219	0.072289
Chacha20 Phase 2 - encrypted_data_127.bin	0.832174-0.583913	0.810905-0.405453
Chacha20 Phase 2 - success count of the 128	127	125

8. Autocorrelation Tests

Algorithm / Metric	Autocorrelation d=1	Autocorrelation d=2	Autocorrelation d=8	Autocorrelation d=16
Chacha20 Phase 1 - Distribution uniformity	0	0	0	0
Chacha20 Phase 1 - encrypted_data_127.bin	0.000000-0.000000	0.000000-0.000000	0.000000-1.000000	0.000000-0.000000
Chacha20 Phase 1 - success count of the 128	9	5	15	3
Chacha20 Phase 2 - Distribution uniformity	0.669618	0.875539	0.556333	0.809134
Chacha20 Phase 2 - encrypted_data_127.bin	0.715665-0.357833	0.663165-0.331583	0.492997-0.246499	0.990650-0.495325
Chacha20 Phase 2 - success count of the 128	127	128	127	127

1. Frequency Tests

Algorithm / Metric	Bit Frequency	Block Frequency m=10000
Chacha20 Phase 1 - Distribution uniformity	0	0
Chacha20 Phase 1 - encrypted_data_127.bin	0.000000-0.000000	0.000000-0.000000
Chacha20 Phase 1 - success count of the 128	2	23
Chacha20 Phase 2 - Distribution uniformity	0.749263	0.919445
Chacha20 Phase 2 - encrypted_data_127.bin	0.490538-0.245269	0.061062-0.061062
Chacha20 Phase 2 - success count of the 128	126	128

2. Poker Tests

Algorithm / Metric	Poker Test m=4	Poker Test m=8
Chacha20 Phase 1 - Distribution uniformity	0	0
Chacha20 Phase 1 - encrypted_data_127.bin	0.000000-0.000000	0.000000-0.000000
Chacha20 Phase 1 - success count of the 128	0	0
Chacha20 Phase 2 - Distribution uniformity	0.406167	0.110612
Chacha20 Phase 2 - encrypted_data_127.bin	0.817546-0.817546	0.496781-0.496781
Chacha20 Phase 2 - success count of the 128	128	124

3. Overlapping Subsequence m=3 Tests

Algorithm / Metric	Overlapping Subsequence Test m=3 P1	Overlapping Subsequence Test m=3 P2
Chacha20 Phase 1 - Distribution uniformity	0	0
Chacha20 Phase 1 - encrypted_data_127.bin	0.000000-0.000000	0.000000-0.000000
Chacha20 Phase 1 - success count of the 128	0	1
Chacha20 Phase 2 - Distribution uniformity	0.3282	0.985508
Chacha20 Phase 2 - encrypted_data_127.bin	0.810226-0.810226	0.611189-0.611189
Chacha20 Phase 2 - success count of the 128	126	126

4. Overlapping Subsequence m=5 Tests

Algorithm / Metric	Overlapping Subsequence Test m=5 P1	Overlapping Subsequence Test m=5 P2
Chacha20 Phase 1 - Distribution uniformity	0	0
Chacha20 Phase 1 - encrypted_data_127.bin	0.000000-0.000000	0.000000-0.000000
Chacha20 Phase 1 - success count of the 128	0	0
Chacha20 Phase 2 - Distribution uniformity	0.540457	0.041862
Chacha20 Phase 2 - encrypted_data_127.bin	0.513439-0.513439	0.213006-0.213006
Chacha20 Phase 2 - success count of the 128	125	127

9. Matrix Rank Detection

Algorithm / Metric	Matrix Rank Detection
Chacha20 Phase 1 - Distribution uniformity	0
Chacha20 Phase 1 - encrypted_data_127.bin	0.000000-1.000000
Chacha20 Phase 1 - success count of the 128	0
Chacha20 Phase 2 - Distribution uniformity	0.448892
Chacha20 Phase 2 - encrypted_data_127.bin	0.171130-0.171130
Chacha20 Phase 2 - success count of the 128	126

10. Cumulative Sum Tests

Algorithm / Metric	Cumulative Sum Forward Detection	Cumulative Sum Backward Detection
Chacha20 Phase 1 - Distribution uniformity	0	0
Chacha20 Phase 1 - encrypted_data_127.bin	0.000000-0.000000	0.000000-0.000000
Chacha20 Phase 1 - success count of the 128	2	2
Chacha20 Phase 2 - Distribution uniformity	0.733647	0.637119
Chacha20 Phase 2 - encrypted_data_127.bin	0.315216-0.315216	0.879297-0.879297
Chacha20 Phase 2 - success count of the 128	127	126

11. Approximate Entropy Tests

Algorithm / Metric	Approximate Entropy m=2	Approximate Entropy m=5
Chacha20 Phase 1 - Distribution uniformity	0	0
Chacha20 Phase 1 - encrypted_data_127.bin	0.000000-0.000000	0.000000-0.000000
Chacha20 Phase 1 - success count of the 128	0	0
Chacha20 Phase 2 - Distribution uniformity	0.3282	0.340461
Chacha20 Phase 2 - encrypted_data_127.bin	0.810323-0.810323	0.213800-0.213800
Chacha20 Phase 2 - success count of the 128	126	127

12. Linear Complexity Tests

Algorithm / Metric	Linear Complexity m=500	Linear Complexity m=1000
Chacha20 Phase 1 - Distribution uniformity	0	0
Chacha20 Phase 1 - encrypted_data_127.bin	0.000000-0.000000	0.000000-0.000000
Chacha20 Phase 1 - success count of the 128	2	0
Chacha20 Phase 2 - Distribution uniformity	0.059452	0.406167
Chacha20 Phase 2 - encrypted_data_127.bin	0.598511-0.598511	0.681561-0.681561
Chacha20 Phase 2 - success count of the 128	125	128

13. Maurer Universal Statistical Test

Algorithm / Metric	Maurer Universal Statistical Test L=7 Q=1280
Chacha20 Phase 1 - Distribution uniformity	0.000012
Chacha20 Phase 1 - encrypted_data_127.bin	0.000012-0.000012
Chacha20 Phase 1 - success count of the 128	14
Chacha20 Phase 2 - Distribution uniformity	0.572333
Chacha20 Phase 2 - encrypted_data_127.bin	0.946864-0.473432
Chacha20 Phase 2 - success count of the 128	127

14. Discrete Fourier

Algorithm / Metric	Discrete Fourier
Chacha20 Phase 1 - Distribution uniformity	0
Chacha20 Phase 1 - encrypted_data_127.bin	0.000000-1.000000
Chacha20 Phase 1 - success count of the 128	0
Chacha20 Phase 2 - Distribution uniformity	0.271449
Chacha20 Phase 2 - encrypted_data_127.bin	0.403700-0.798150
Chacha20 Phase 2 - success count of the 128	126

Algorithm	Type	Key size (bits)	Block size (bits)
Ascon	Block cipher	80, 128	64
LED	Block cipher	64, 128	64
PHOTON	Hash function	N/A	N/A
SPONGENT	Hash function	N/A	N/A
PRESENT	Block cipher	80 , 128	64
CLEFIA	Block cipher	128, 192, 256	128
LEA	Block cipher	128, 192, 256	128
Grain-128a	Stream cipher	128	N/A
Enocoro-128v2	Stream cipher	128	N/A
Lesamnta-LW	Hash function	N/A	N/A
ChaCha	Stream cipher	128, 256	N/A
LBlock	Block cipher	80	64
SIMECK	Block cipher	64, 128	32, 48, 64
SIMON	Block cipher	64, 72, 96, 128, 144, 192	32, 48, 64, 96, 128
PRIDE	Block cipher	128	64
TWINE	Block cipher	80, 128	64
ESF	Block cipher	128	128

Table 3: Cryptography Algorithm Specifications (No Internal State or Initial Vector) - Part 1

AEAD mode
Yes
No
Yes
No

Table 4: AEAD Mode for Cryptography Algorithms

Algorithm	Rounds of use round function
Ascon	12 or 8
LED	48 or 32
PRESENT	31
CLEFIA	18, 22, 26
LEA	24, 28, or 32
SIMECK	32, 36, or 44
SIMON	32, 36, 42, 44, 52, 54, 68, 69, 72, 84
PRIDE	20
TWINE	36

Table 5: Rounds of Use Round Function for Cryptography Algorithms

Algorithm	Internal state size (bits)	Initial vector size (bits)
Ascon	320	128
Grain-128a	256	96
Enocoro-128v2	128	128
Lesamanta-LW	256	N/A
ChaCha	512	64 or 96

Table 6: Cryptography Algorithm Specifications (With Internal State and Initial Vector)

Key size (bits)	Block size (bits)	AEAD mode
80 , 128	64	Yes
128, 256	32	Yes
128 or Custom	128 or Mutable	Yes (eg. Poly1305)

Table 7: Comparison of Key Size, Block Size, and AEAD Mode for Cryptography Algorithms: Ascon, ChaCha20, and XCR/Little_OaldresPuzzle_Cryptic (Excluding Internal State and Initialization Vector Sizes, and Rounds of Use Round Function)

Internal State Size (bits)	Initialization Vector Size (bits)
64	128
512	64 or 96
192	Not specified

Table 8: Internal State and Initialization Vector Sizes for The opposite algorithm we chose to compare with Ascon, Chacha20, XCR/Little_OaldresPuzzle_Cryptic

Rounds of use round function
12 or 8
8, 12, or 20
Custom

Table 9: Rounds of Use Round Function for The opposite algorithm we chose to compare with Ascon, Chacha20, XCR/Little_OaldresPuzzle_Cryptic